



Data Scientist

STUDENT GUIDE



© Copyright 2021
Pegasystems Inc., Cambridge, MA
All rights reserved.

This document describes products and services of Pegasystems Inc. It may contain trade secrets and proprietary information. The document and product are protected by copyright and distributed under licenses restricting their use, copying, distribution, or transmittal in any form without prior written authorization of Pegasystems Inc.

This document is current as of the date of publication only. Changes in the document may be made from time to time at the discretion of Pegasystems. This document remains the property of Pegasystems and must be returned to it upon request. This document does not imply any commitment to offer or deliver the products or services provided.

This document may include references to Pegasystems product features that have not been licensed by your company. If you have questions about whether a particular capability is included in your installation, please consult your Pegasystems service consultant.

PegaRULES, Process Commander, SmartBPM® and the Pegasystems logo are trademarks or registered trademarks of Pegasystems Inc. All other product names, logos and symbols may be registered trademarks of their respective owners.

Although Pegasystems Inc. strives for accuracy in its publications, any publication may contain inaccuracies or typographical errors. This document or Help System could contain technical inaccuracies or typographical errors. Changes are periodically added to the information herein. Pegasystems Inc. may make improvements and/or changes in the information described herein at any time.

This document is the property of:
Pegasystems Inc.
1 Rogers Street
Cambridge, MA 02142
Phone: (617) 374-9600
Fax: (617) 374-9620
www.pega.com

Mission: Data Scientist

Product: Pega Customer Decision Hub™ 8.6

URL: <https://academy.pega.com/mission/data-scientist/v3>

Date: 14 May 2021

Contents

Pega AI overview.....	5
Next-Best-Action paradigm.....	6
One-to-one customer engagement paradigm.....	11
Action arbitration.....	17
Next-Best-Action Designer.....	23
Prediction Studio.....	30
Enabling AI in the NBA framework.....	36
Adaptive analytics.....	37
Predictors and outcomes of an adaptive model.....	39
Advanced settings of an adaptive model.....	43
Adaptive model outputs.....	45
Configuring an adaptive model.....	46
Monitoring adaptive models.....	53
Measuring lift using a control group.....	54
Regular monitoring of adaptive models.....	60
Inspecting adaptive models.....	62
Exporting historical data.....	69
Creating and understanding decision strategies.....	74
Decision strategies.....	75
Decision strategy canvas.....	80
Creating a decision strategy.....	88
Decision strategy execution.....	94
Creating predictions.....	105
Predictive models drive predictions.....	106
Creating a prediction.....	111
Using predictions in engagement strategies.....	113
Creating predictive models.....	119
Predictive models.....	120

Building models with Pega machine learning.....	124
Importing predictive models.....	132
Using machine learning services.....	137
Model transparency.....	144
Upgrading predictions.....	146
Model operations process.....	147
Placing a predictive model in shadow mode.....	153
Text analytics for email routing.....	156
Text analytics.....	157
Email routing.....	162
Training a topic model to improve email routing.....	168

Pega AI overview

Description

Familiarize yourself with the one-to-one customer engagement paradigm and discover how Pega's omni-channel AI delivers the right action during every customer interaction.

Prediction Studio is the dedicated workspace for data scientists to control the life cycles of predictions and the predictive models that drive them. Prediction Studio offers prediction and model reports that allow the user to monitor and spot predictions and models that underperform.

Learning objectives

- Explain the basics of the Next-Best-Action approach
- Describe the purpose of Next-Best-Action Designer and the user interface
- Explain the types of predictions that are available in Prediction Studio
- Describe the purpose of the control group
- Describe the bubble chart that visualizes the adaptive model performance
- Recognize the transparency settings for predictive models

Next-Best-Action paradigm

Introduction

The value of big data and analytics is fully realized when every customer conversation delivers exactly the right message, the right offer, or the right level of service to provide a great experience while maximizing the customer's value to the organization. With Pega Next-Best-Action, business experts develop decision strategies that combine predictive and adaptive analytics with traditional business rules to maximize this value.

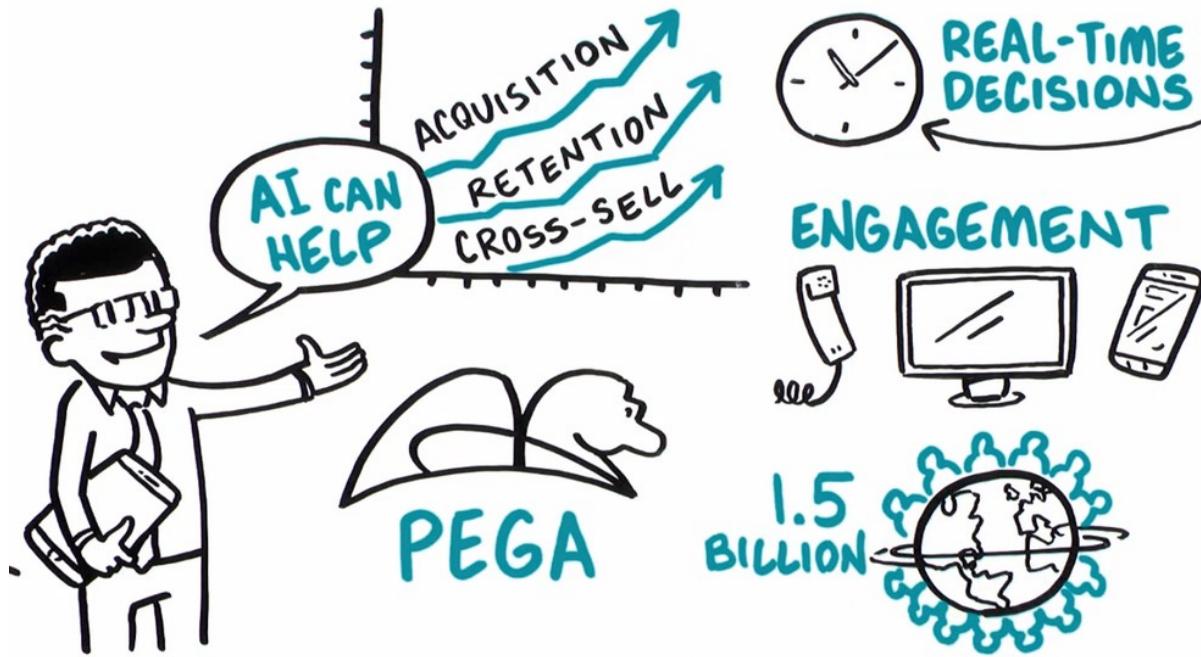
Transcript

This is your customer. You want him to buy your products, use your services and have a great experience. And your competitors want the same thing. To compete, you have to take the right action at every customer touch, ensuring that each conversation delivers exactly the right message, offer and level of service. You want to provide a great experience, while maximizing the customer's value to your organization.



Artificial Intelligence, or AI, can help—if you can get past the hype. Pega has been using AI to create real business value for years, driving real-time decisions that deliver awesome

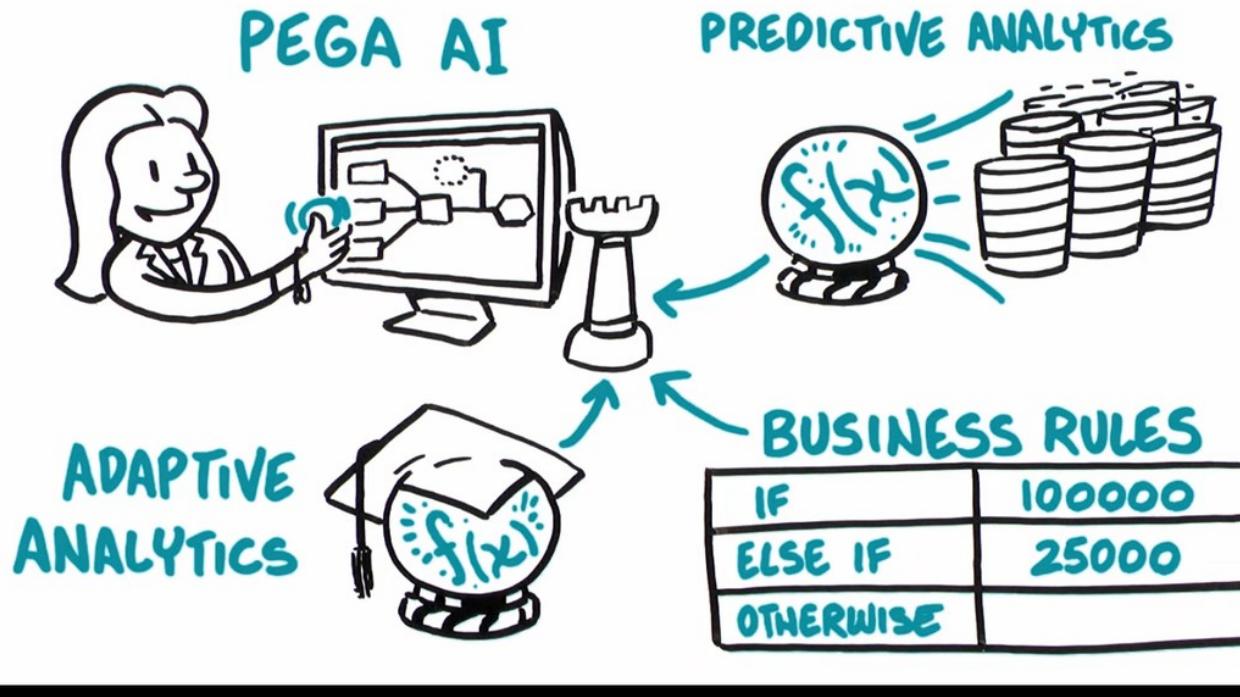
engagement on any channel and improving experiences for over 1.5 billion customers across the globe.



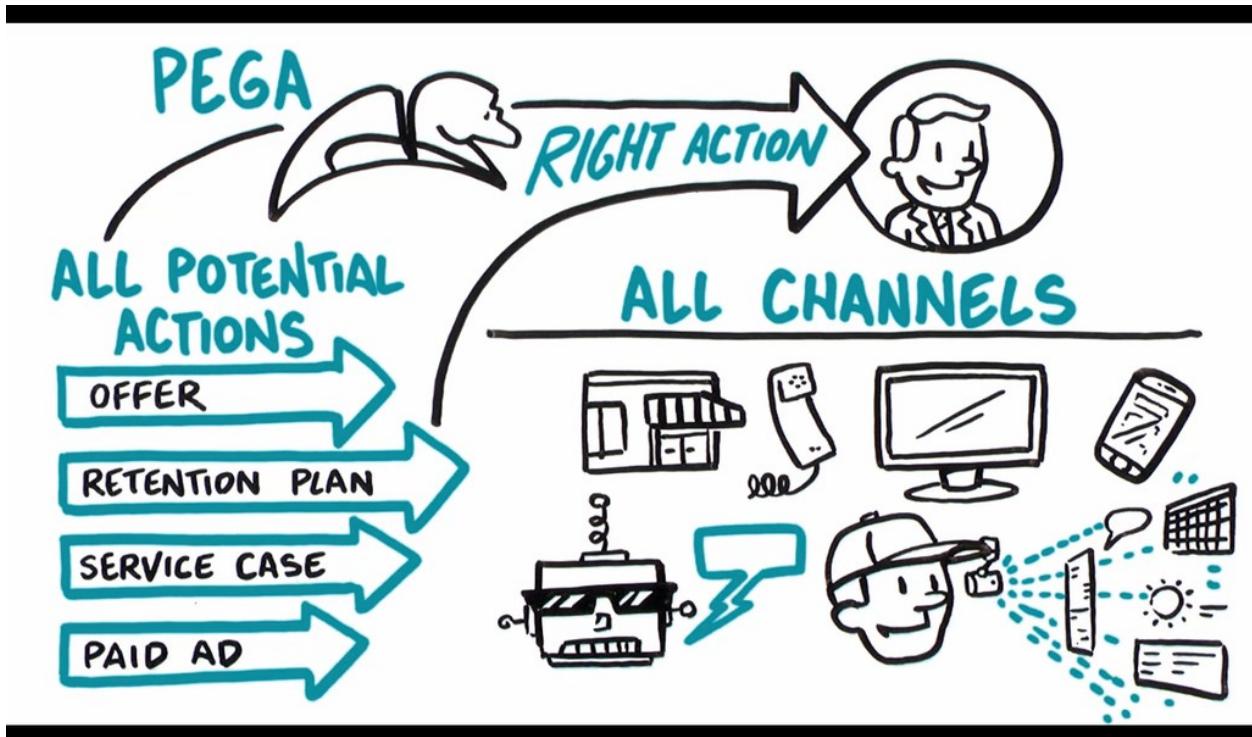
Pega's omni-channel AI delivers the right action at every customer touch by crunching millions of data points in real-time. Make an offer, initiate a retention plan, predict a problem before it happens. Every decision generates the next-best-action for your customer, and your business.



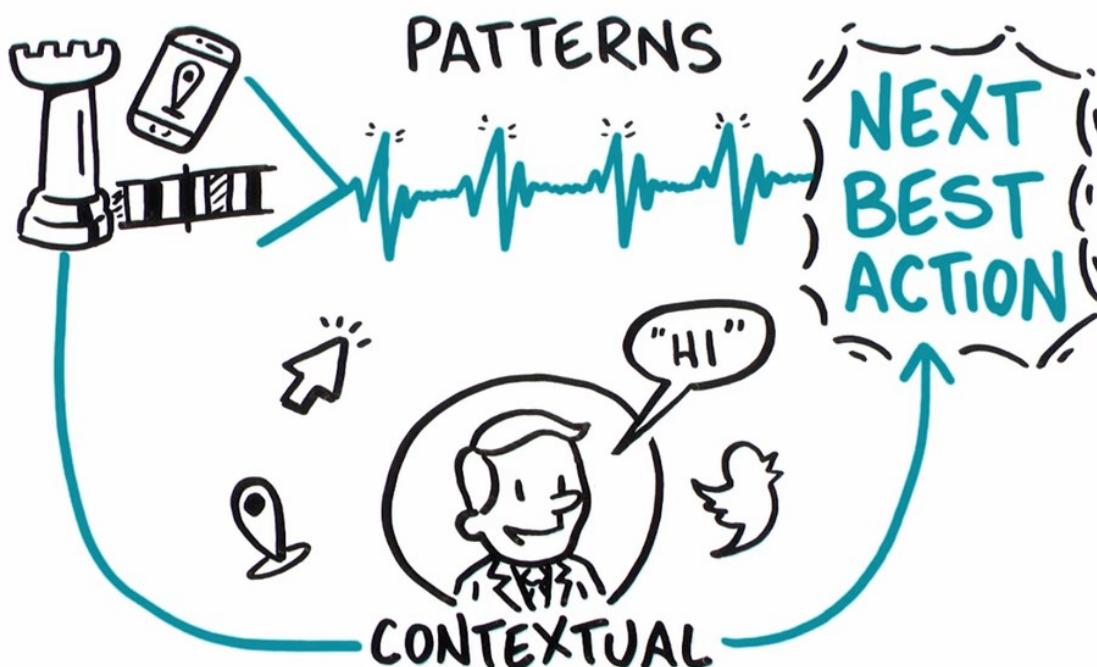
Pega's AI is built for business people, not scientists or developers. They design visual decision strategies that combine predictive analytics, algorithms developed through mining large sets of data, adaptive analytics, machine-learning algorithms that improve with each interaction, and traditional business rules that allow users to prioritize and arbitrate between decisions.



Pega uses the strategy to look across all the potential actions you may take with a customer, make an offer, initiate a retention plan, open a service case, place an ad, and ensure exactly the right action is taken at every interaction and it works across all channels to provide a consistent experience in a store, on the phone, on the web, mobile, with the chat bot, or just some crazy tech that hasn't even been invented yet.



And Pega connects to streams like mobile locations or network events to detect patterns and drive the Next Best Action proactively. And strategies are completely contextual. Any change in the customer's context — a click, a reply, a location change, a Tweet — will trigger the Next Best Action. So, you can really listen to your customers and act accordingly.



Pega's real-time, omni-channel AI puts the power in your hands, so you can optimize every customer interaction for experience, and value.



One-to-one customer engagement paradigm

Introduction

The optimal outcome of every customer interaction is to provide a great experience while maximizing the customer's value to the company. To achieve this, you have to be able to perform the right action in the right channel at the right moment for each customer. We call this capability, "1-to-1 Customer Engagement".

Transcript

In this video, learn about the 1-to-1 Customer Engagement paradigm and how the principles of Next-Best-Action are implemented using the Pega Customer Decision Hub™.

Customers are more empowered than ever before. As a result, they have very high expectations of the experiences they receive from their service providers. Their experiences must make sense within the context of their lives. This means they must be meaningful, consistent, and personalized across every channel they interact with.



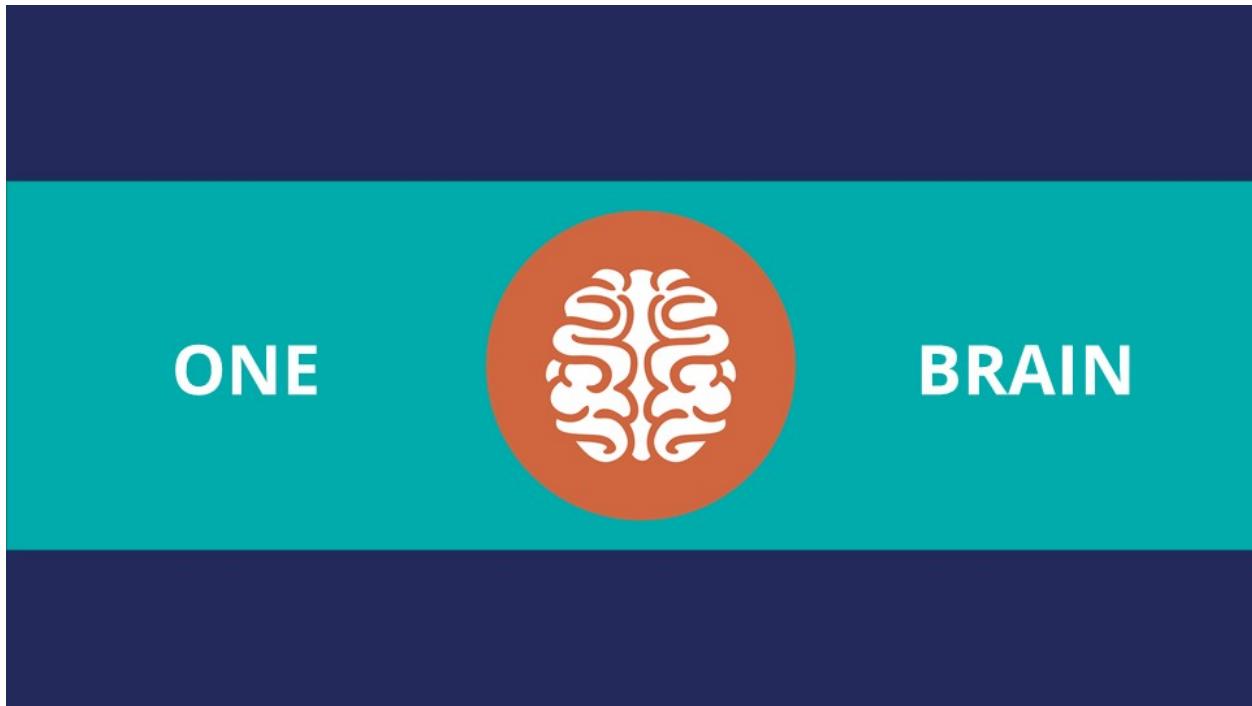
In business, the optimal outcome of every customer interaction is to provide a great experience while maximizing the customer's value to the company. To achieve this, you have to be able to perform the right action in the right channel at the right moment for each customer.

We call this capability, "1-to-1 Customer Engagement".

1-to-1 Customer Engagement

1-to-1 Customer Engagement enables companies to transition their marketing away from a traditional one-to-many campaign-driven approach. A one-to-one approach allows companies to have consistent, contextual and relevant conversations with individual customers across any channel or touch point.

The key to achieving 1-to-1 Customer Engagement is an idea that's simple to conceive, but very difficult to execute: one centralized brain.



In other words, one piece of intelligence that acts as a single decision authority across your application ecosystem.

Each channel or system profits from this single source of customer intelligence and can leverage it to gain insights or perform relevant actions.

In Pega Marketing™, this centralized brain is called the Pega Customer Decision Hub, and it leverages AI to enable 1-to-1 Customer Engagement.

In Pega Infinity™, the Pega Customer Decision Hub forms the core of the customer engagement platform, which sits at the center of existing systems and channels in an enterprise.

The Customer Decision Hub

a Single Decision Authority



Data from every customer engagement across the enterprise is collected by the Brain and used to make predictions and decisions about every interaction in every channel.

Continuous learning and decision-making are the foundation of a 1-to-1 Customer Engagement solution.

The Customer Decision Hub combines analytics, business rules, customer data, and data collected during each customer interaction to create a set of actionable insights that it uses to make intelligent decisions. These decisions are known as the Next-Best-Action.

Every Next-Best-Action weighs customer needs against business objectives to optimize decisions based on priorities set by the business manager.

In the milliseconds before interacting with a customer, the Customer Decision Hub processes thousands of predictive and adaptive models to determine customer needs, considering the customer's immediate context to ensure the Next-Best-Action is relevant, timely, and contextual. These models can be propensity, risk, or churn models.

Next, the decision strategy considers business rules and matches those with the customer's context and higher-level business goals.



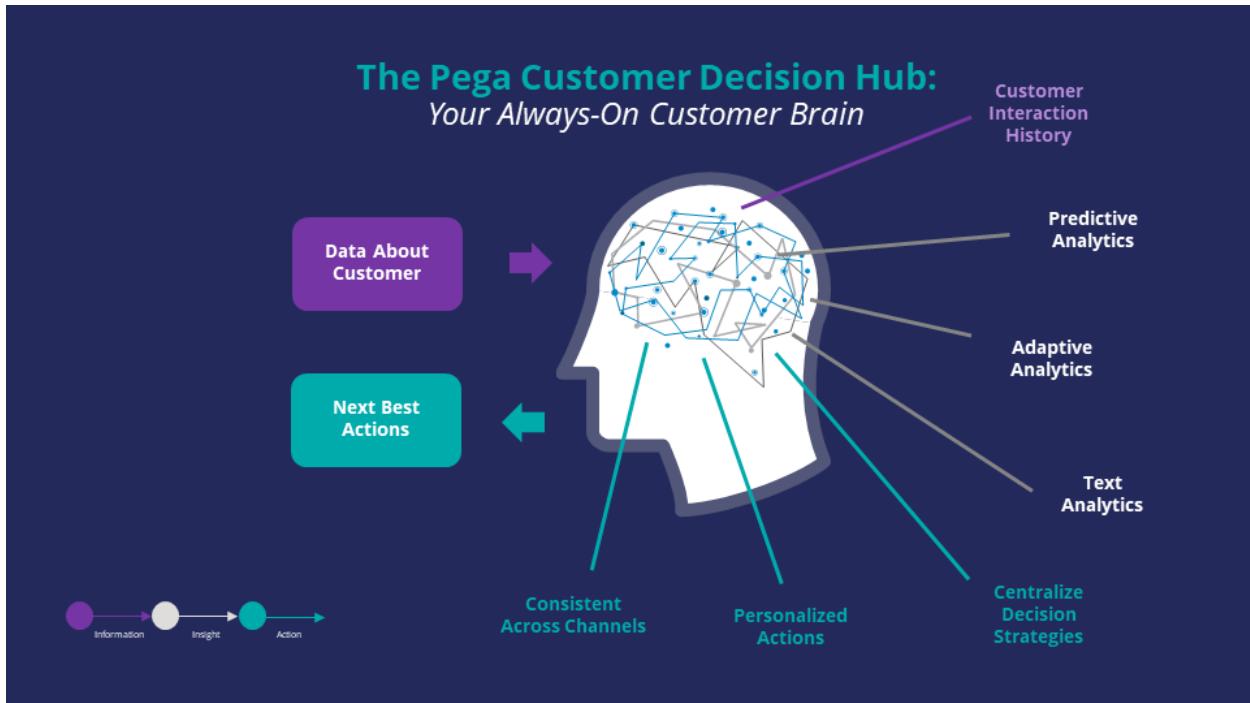
All of this information is used by the Next-Best-Action decision strategy to evaluate every potential action that could be taken with a particular customer in a given situation. The decision strategy then recommends the best way to interact with the customer to achieve the optimal result.

Using the Next-Best-Action approach, the Customer Decision Hub is able to identify the best moments for making a sale, providing a service, making a retention offer, or doing nothing at all (e.g. if nothing is relevant enough to warrant the customer's attention). Next-Best-Action is even able to select which offers are most likely to be accepted by the customer in a sales or retention situation. Next-Best-Action decisions are distributed, in real-time, to each of your real-time owned channels, such as web, mobile, and contact center. Through Pega Marketing, Next-Best-Actions can also be distributed to real-time paid channels such as Google, YouTube, Facebook, LinkedIn and Instagram. Pega Marketing also integrates with non-real time outbound channels such as data management platforms (DMPs) and email.

Once the Next-Best-Actions are distributed and customer responses have been received by the Brain, the whole process begins again, and new Next-Best-

Actions are distributed within milliseconds. Every outbound channel, including a data management platform, is dynamically updated with the Next-Best-Action to ensure consistency and an optimized customer experience no matter which channel the customer interacts with.

In summary, the Pega Customer Decision Hub is the Always-On Brain that acts as a single, centralized decision authority.



It uses data about the customer, including past interactions, as input.

It leverages advanced AI techniques to make predictions.

And it uses decision strategies (which combine traditional business rules with predictive, adaptive and text analytics), to deliver consistent and personalized Next-Best-Actions across all channels.

Action arbitration

Introduction

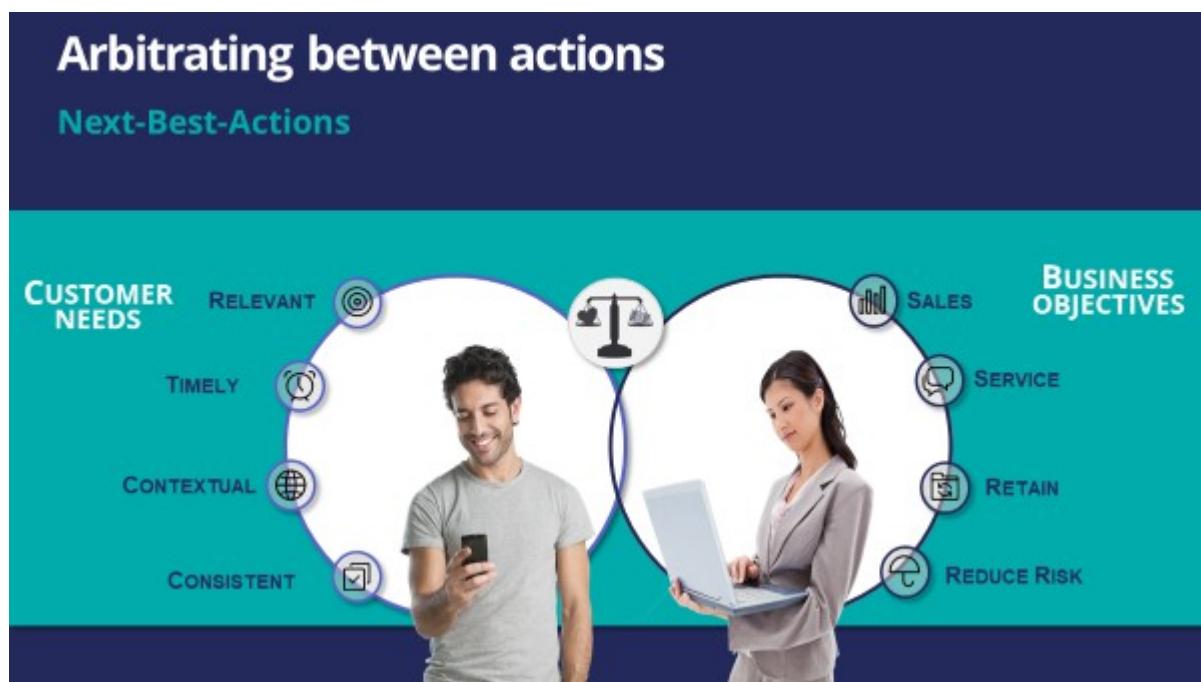
Pega Customer Decision Hub combines analytics, business rules, customer data, and data collected during each customer interaction to create a set of actionable insights that it uses to make intelligent decisions. Arbitration aims to balance customer relevance with business priorities by weighing numerical values for the following factors: propensity, context weighting, action value, and business levers. Learn to create a simple formula for arriving at a prioritization value, which is used to select the top actions.

Transcript

This video explains the concept of action arbitration.

Pega Customer Decision Hub™ combines analytics, business rules, customer data, and data collected during each customer interaction to create a set of actionable insights that it uses to make intelligent decisions. These decisions are known as Next-Best-Action.

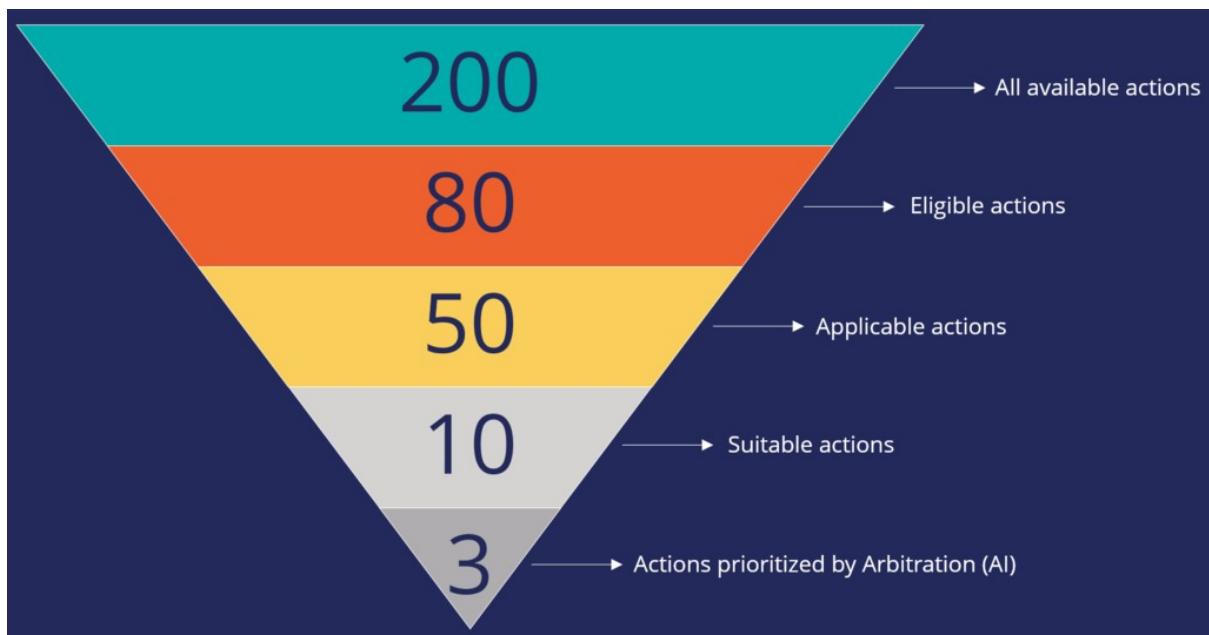
Every Next-Best-Action weighs customer needs against business objectives to optimize decisions based on priorities set by the business manager.



U+ Bank, a retail bank, has several actions for its customers and has configured engagement policies to suit both customer needs and business objectives.

In this scenario, a marketer for U+ has designed 200 actions that can be presented to customers. To select the Next-Best-Actions from these, Pega Customer Decision Hub first checks the eligibility conditions and filters the actions. Then, the applicability conditions are run to filter it further. Next, Customer Decision Hub checks the suitability conditions to derive the final set of available actions.

These actions move through one final stage before being presented to customers: the arbitration stage. Arbitration is used to prioritize and choose the best actions based on what is relevant for the customer right now.



Arbitration aims to balance customer relevance with business priorities. The factors weighed are **Propensity**, **Context Weighting**, **Action Value**, and **Business Levers**, each represented by numerical values. A simple formula is used to arrive at a prioritization value, which is used to select the top actions. The number of top actions selected depends on the channel of interaction. For example, the top three actions, plus two tiles and one hero treatment, can be selected for display on a bank's website.



Propensity is the likelihood of a customer responding positively to an action; this is calculated by AI. For example, the higher the likelihood of a customer accepting an offer, the higher the Propensity value for that offer.

Context Weighting allows Pega Customer Decision Hub to consider the situational context for each action. For example, if a customer contacts the bank to close their account, the highest-priority action is to ensure that the customer is retained. The priority of an action is increased by a specified value when the context is detected.

Action Value enables you to assign a financial value to an action and prioritize high-value actions over low-value ones. This value is typically normalized across Issues and Groups. For example, an unlimited data plan is more profitable than a limited data plan. So, in a situation where a customer is eligible for both plans, the unlimited plan has higher priority.

Business Levers allow the business to assert some level of control over the prioritization of actions defined within the system. Levers are used to manually nudge Customer Decision Hub toward Next-Best-Actions based on external factors. For example, the recommended Next-Best-Action might be to offer a credit card to a customer when they visit the home page. But to meet a business goal, the Mortgage Line of Business favors a mortgage offer even if that offer is ranked a little lower on the list of possible actions.

Consider an example where three actions are selected for arbitration. At the moment, only the Propensity is used for prioritization.

Rank	Issues	Groups	Actions	Propensity	Context weighting	Action value	Business levers	Priority
	Sales	Credit cards	Gold Card	0.5	1	1	1	0.5
	Retention	Proactive	10% discount	0.55	1	1	1	0.55
	Service	Administrative	Address change	0.4	1	1	1	0.4

Action arbitration with propensity before prioritization

The result of the arbitration is that the top action is the one with the highest Propensity.

Rank	Issues	Groups	Actions	Propensity	Context weighting	Action value	Business levers	Priority
1	Retention	Proactive	10% discount	0.55	1	1	1	0.55
2	Sales	Credit cards	Gold Card	0.5	1	1	1	0.5
3	Service	Administrative	Address change	0.4	1	1	1	0.4

Action arbitration with propensity after prioritization

Examine what happens when Context Weighting together with Propensity are considered for arbitration. For example, if the intent of a customer calling customer service is to change their address, the Context Weight of a Service action increases.

Rank	Issues	Groups	Actions	Propensity	Context weighting	Action value	Business levers	Priority
1	Retention	Proactive	10% discount	0.55	1	1	1	
2	Sales	Credit cards	Gold Card	0.5	1	1	1	
3	Service	Administrative	Address change	0.4	2	1	1	

Action arbitration with context weight before prioritization

As a result, the Arbitration caters to the current need of the customer and presents a Service action as the top action for the customer. Thus, the Arbitration caters to the current need of the customer and presents a Service action as the top action for the customer.

Rank	Issues	Groups	Actions	Propensity	Context weighting	Action value	Business levers	Priority
1	Service	Administrative	Address change	0.4	2	1	1	0.8
2	Retention	Proactive	10% discount	0.55	1	1	1	0.55
3	Sales	Credit cards	Gold Card	0.5	1	1	1	0.5

Action arbitration with context weight after prioritization

Consider another scenario in which a customer is eligible for two credit cards and two other actions. Now, consider that the Action Value is also used in arbitration when prioritizing. In this case, the Platinum Card is assigned a higher value by the business than the Gold Card.

Rank	Issues	Groups	Actions	Propensity	Context weighting	Action value	Business levers	Priority
1	Sales	Credit cards	Gold Card	0.6	1	1	1	
2	Sales	Credit cards	Platinum Card	0.55	1	2	1	
3	Retention	Proactive	10% discount	0.2	1	1	1	
4	Service	Administrative	Address change	0.1	1	1	1	

Action arbitration with action value before prioritization

Thus, the arbitration selects the Platinum Card as the top action.

Rank	Issues	Groups	Actions	Propensity	Context weighting	Action value	Business levers	Priority
1	Sales	Credit cards	Platinum Card	0.55	1	2	1	1.1
2	Sales	Credit cards	Gold Card	0.6	1	1	1	0.6
3	Retention	Proactive	10% discount	0.2	1	1	1	0.2
4	Service	Administrative	Address change	0.1	1	1	1	0.1

Action arbitration with action value after prioritization

Finally, consider an example in which all four parameters are used for arbitration. In this case, U+ Bank wants to promote two new checking account offers under the Sales issue. The bank sets a higher Business Lever value for the Checking Accounts actions.

Rank	Issues	Groups	Actions	Propensity	Context weighting	Action value	Business levers	Priority
1	Sales	Credit cards	Gold Card	0.6	1	1	1	
2	Sales	Credit cards	Platinum Card	0.55	1	1	1	
3	Sales	Checking Accounts	Premium Checking	0.55	1	1	2	
4	Sales	Checking Accounts	Student Checking	0.5	1	1	2	
5	Retention	Proactive	10% discount	0.2	1	1	1	
6	Service	Administrative	Address change	0.1	1	1	1	

Action arbitration with business levers before prioritization

Although the Propensity of the Checking Accounts actions is low, they are selected as the top actions due to their high Lever values.

Rank	Issues	Groups	Actions	Propensity	Context weighting	Action value	Business levers	Priority
1	Sales	Checking Accounts	Premium Checking	0.55	1	1	2	1.1
2	Sales	Checking Accounts	Student Checking	0.5	1	1	2	1
3	Sales	Credit cards	Gold Card	0.6	1	1	1	0.6
4	Sales	Credit cards	Platinum Card	0.55	1	1	1	0.55
5	Retention	Proactive	10% discount	0.2	1	1	1	0.2
6	Service	Administrative	Address change	0.1	1	1	1	0.1

Action arbitration with business levers after prioritization

Next-Best-Action Designer

Introduction

Next-Best-Action Designer guides you through the creation of a Next-Best-Action strategy for your business. Its intuitive interface, proven best practices and sophisticated underlying decisioning technology enable you to automatically deliver personalized customer experiences across inbound, outbound and paid channels. Next-Best-Action Designer is organized according to the high-level sequence of steps needed to configure the Next-Best-Action strategy for your organization.

Transcript

Next-Best-Action Designer guides you through the creation of a Next-Best-Action strategy for your business. Its intuitive interface, proven best practices and sophisticated underlying decisioning technology enable you to automatically deliver personalized customer experiences across inbound, outbound and paid channels.

The Next-Best-Action Designer user interface allows you to easily define, manage and monitor Next-Best-Actions.

The tabs across the top of the user interface represent the steps that need to be completed to define Next-Best-Actions.

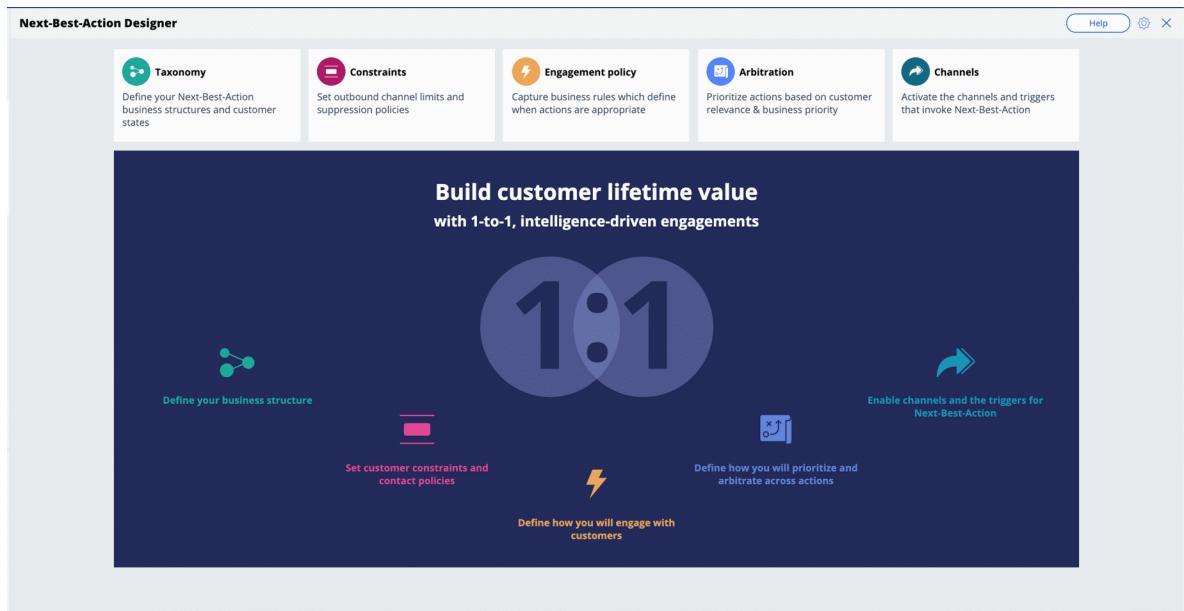
Use the **Taxonomy** component to define the business structure for your organization.

Use the **Constraints** component to implement channel limits and constraints.

Use the **Engagement policy** component to define the rules that control which actions are offered to which customers.

Use the **Arbitration** component to configure how actions are prioritized.

Use the **Channels** component to configure when and where Next-Best-Action is triggered.



The system uses these definitions to create an underlying Next-Best-Action Strategy framework. This framework leverages best practices to generate Next-Best-Action decision strategies at the enterprise level. These decision strategies are a combination of the business rules and AI models that form the core of the Pega Centralized Decision Hub, which determines the personalized set of Next-Best-Actions for each customer.

Use the **Taxonomy** component to define the hierarchy of Business Issues and Groups to which an action belongs.

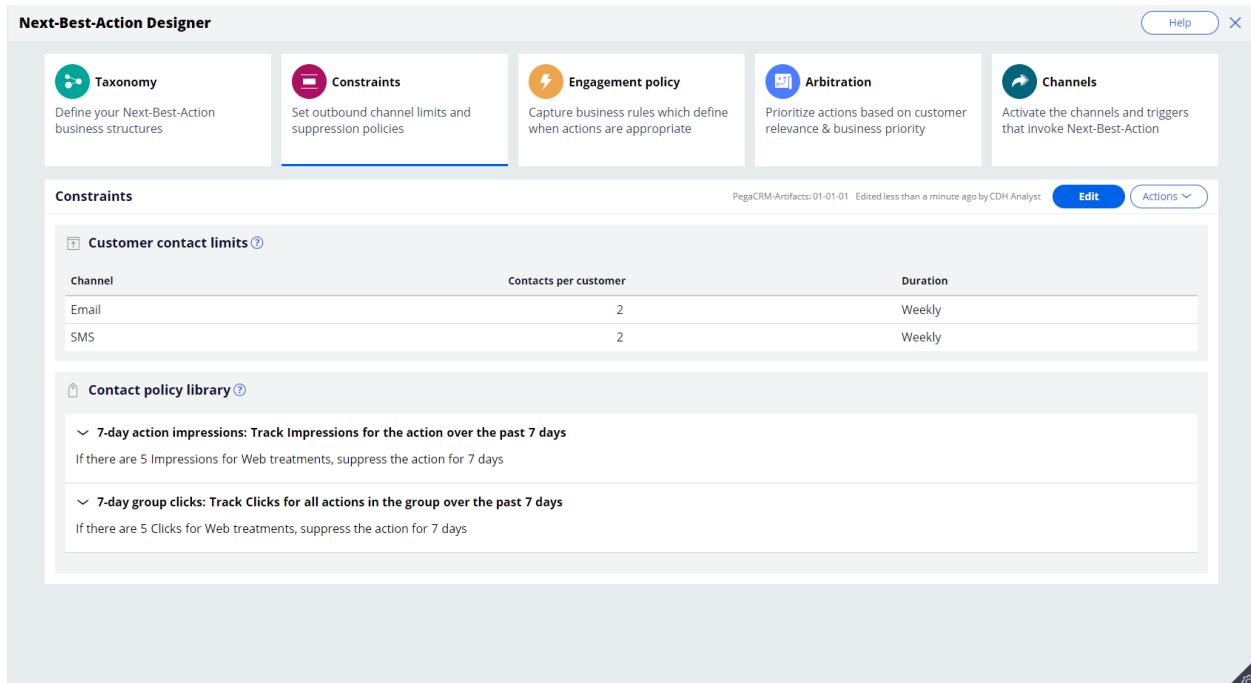
The screenshot shows the 'Taxonomy' component within the 'Next-Best-Action Designer'. It displays a table with columns for 'Issues / Groups', 'Description', and 'Action naming'. The table contains the following data:

Issues / Groups	Description	Action naming
Acquire	Customer acquisition	
↳ Mortgage	Home mortgage offerings for acquisition	
↳ Cards	Credit card offerings for acquisition	Promotion
Retain	Customer retention	
↳ Mortgage	Home mortgage offerings for retention	
↳ Cards	Credit card offerings for retention	

A Business Issue is the purpose behind the actions you offer to customers. For example, actions with the purpose of retaining existing customers should be grouped under the business Issue of Retention. Actions with the purpose of acquiring new customers belong to the business Issue of Acquisition.

Business Groups are used to organize customer actions into categories. For example, as part of the business Issue of Acquisition, you can create Groups for products like Credit Cards, Mortgages, or Personal Loans, with the intention of offering these to potential customers.

Use **Constraints** to specify outbound contact limits as well as to limit overexposure to a specific action or group of actions.



The screenshot shows the 'Next-Best-Action Designer' interface. At the top, there are five tabs: 'Taxonomy', 'Constraints' (which is selected), 'Engagement policy', 'Arbitration', and 'Channels'. Below the tabs, there's a section titled 'Constraints' with two expandable sections: 'Customer contact limits' and 'Contact policy library'. Under 'Customer contact limits', there's a table with two rows: one for Email (2 contacts per customer, weekly duration) and one for SMS (2 contacts per customer, weekly duration). Under 'Contact policy library', there are two collapsed sections: '7-day action impressions' and '7-day group clicks', each with a note about tracking and suppressing actions for 7 days if there are 5 impressions or clicks respectively.

Customer contact limits allow you to limit the number of interactions that a customer can receive over a given period of time on a specific channel. For example, you can decide that you do not want your customers to receive more than two emails per week or two SMS messages per week.

On the Constraints tab of Next-Best-Action Designer, you can define more extensive suppression rules by creating Contact Policy rules in the library. Contact Policy rules are reusable across all Business Issues and Groups.

In the Contact Policy library, you define suppression rules that automatically put an action on hold after a specific number of outcomes are recorded for some or all channels. For example, an action can be suppressed for a customer for seven days after the customer has seen an ad for that action five times. Suppressing or pausing an action prevents over-exposure by limiting the number of times a customer is exposed to the same action.

Use **Engagement policies** to define when specific actions or groups of actions are appropriate for customers.

The screenshot shows the 'Next-Best-Action Designer' interface. At the top, there are five tabs: 'Taxonomy', 'Constraints', 'Engagement policy' (which is selected and highlighted in orange), 'Arbitration', and 'Channels'. Below the tabs, on the left, is a sidebar titled 'Business structure' with categories like 'Acquire', 'Mortgage', 'Cards', 'Retain', 'Mortgage', 'Cards', 'Sales', 'CreditCards' (which is selected and highlighted in blue), and 'Mortgage'. The main panel shows a 'Sales / CreditCards' entry. Under 'Engagement policy', there are four sections: 'Eligibility' (with conditions '(isCustomer is true) and (Age is greater than 18)'), 'Applicability' (with condition '(Has Cards is equal to N)'), 'Suitability' (with condition '(CreditScore is greater than 500)'), and 'Contact policy' (with two collapsed sections: '7-day action impressions: Track Impressions for the action over the past 7 days' and '7-day group clicks: Track Clicks for all actions in the group over the past 7 days').

There are four types of engagement policies:

Eligibility determines whether or not a customer qualifies for an action or group of actions. For example, an action may only be available for customers over a specific age or living in a specific geographic location.

Applicability determines if an action or group of actions is relevant for a customer at a particular point in time. For example, a discount on a specific credit card may not be relevant for a customer who already owns a card.

Suitability determines if an action or group of actions is appropriate for a customer for ethical or empathetic reasons. For example, a new loan offer may not be appropriate for a customer whose credit score is low, even though it might be profitable for the bank.

Contact Policies determine when an action or group of actions should be suppressed and for how long. For example, you can suppress an action after a specific number of promotional messages has been sent to customers. To activate Contact Policy rules created in the library on the Constraints tab, add them to the Engagement Policy tab.

Arbitration determines how the Customer Decision Hub prioritizes the list of eligible and appropriate actions that come out of each group.

The screenshot shows the 'Next-Best-Action Designer' interface. At the top, there are five tabs: 'Taxonomy', 'Constraints', 'Engagement policy', 'Arbitration' (which is selected and highlighted in blue), and 'Channels'. Below the tabs, there's a section titled 'Arbitration' with two main categories: 'Customer relevance' and 'Business priority'. Each category has two circular icons with arrows between them, labeled 'Propensity' and 'Context weighting' respectively. Under 'Customer relevance', there are also 'Action value' and 'Business levers'. On the left side, there are three sections: 'Propensity' (with a toggle switch off), 'Context weighting' (with a table showing a row for 'CallReason' with value 'Enquire credit cards' and issue/group 'Sales / CreditCards' with weight '20%'), and 'Action value' (with a toggle switch off). A status bar at the bottom indicates 'PegaCRM-Artifacts:01-01-01 Edited 21 hours ago by CRM Decisioning Analyst'.

The factors weighed in arbitration are: Propensity, Context weighting, Action value, and Business levers, each represented by numerical values. A simple formula is used to arrive at a prioritization value, which is used to select the top actions.

Propensity is the likelihood of a customer responding positively to an action. This is calculated by Artificial Intelligence (AI). For example, a click on an offer banner or an accept of an offer in the contact center are considered positive behaviors.

Real-time contextual data is an important part of making highly relevant recommendations. **Context weighting** allows you to assign weighting to a specific context value for all actions within an Issue or Group. For example, if a customer contacts the bank to change their address, the weight of the Service context will increase, and the highest priority action will be to ensure that the relevant service is delivered to the customer.

Action value enables you to assign a financial value to an action and prioritize high-value actions over low-value ones. For example, promoting an unlimited data plan might be more

profitable for the company than a limited data plan. Action values are typically normalized across Issues and Groups.

Business levers enable you to accommodate ad hoc business priorities by specifying a weight for an action or Group of actions and/or its associated Business Issue.

Next-Best-Action Designer enables Next-Best-Actions to be delivered via inbound, outbound and paid channels.

The screenshot shows the 'Next-Best-Action Designer' interface. At the top, there are five tabs: 'Taxonomy', 'Constraints', 'Engagement policy', 'Arbitration', and 'Channels'. The 'Channels' tab is active, indicated by a blue border and the word 'EDITING' at the top right. Below the tabs, there are two main sections: 'Channels' and 'Triggers'.

Channels:

- Call center: On
- Email: On
- Mobile: On
- Other: Off
- Paid: On
- Push notification: On
- Retail: On
- SMS: On
- Web: On

Triggers:

Real-time containers

Status	Name	Description	Business structure level
ACTIVE	TopOffers	Top Offers	All Issues / All Groups
ACTIVE	MortgagesLandingPage	Mortgages Landing Page	Retain / Mortgage
ACTIVE	AccountCreditCards	Account Credit Cards	Sales / CreditCards

Events

These channels can be toggled on or off. If a channel is toggled off, the Next-Best-Actions will not be delivered to that channel.

An external real-time channel is any channel that presents actions selected by the Customer Decision Hub to a customer. These channels can include a website, or a call-center or mobile application. A real-time container is a placeholder for content in an external real-time channel.

A trigger is a mechanism whereby an external channel invokes the execution of a Next-Best-Action decisioning process for specific Issues and Groups. The result will be delivered back to the invoking channel. For example, when a real-time container called "Mortgages Landing Page" is configured, the website invokes this real-time container before loading the mortgage page.

As you have seen in this video, Next-Best-Action Designer is organized according to the high-level sequence of steps needed to configure the Next-Best-Action strategy for your organization. These steps involve:

- Defining the business structure for your organization
- Implementing the channel limits and constraints
- Defining the rules that control which actions are offered to which customers
- Configuring how actions are prioritized
- Configuring when and where Next-Best-Action is triggered

Prediction Studio

Introduction

Prediction Studio is the dedicated workspace for data scientists to control the life cycles of predictions created in Pega Customer Decision Hub™ to optimize 1:1 customer interactions. From Prediction Studio, you can also manage additional resources, such as data sets, summaries and predictive models.

Transcript

This video shows how you can use Prediction Studio to control the life cycles of predictions and the predictive models that drive them.



Manage predictions

Predictions combine predictive analytics and best practices in data science.

In Pega Customer Decision Hub™, many predictions are available out-of-the-box. They are listed on the **Predictions** landing page in Prediction Studio.

For example, **Predict web propensity** aims at predicting the likelihood that a customer clicks a web banner.

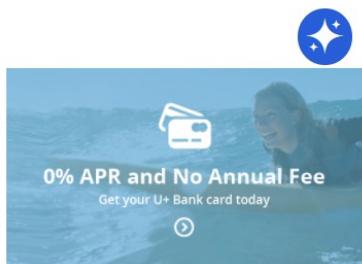
Predict Web Propensity

No performance available yet

Decision Request Jul 28, 2020

Open prediction :

Consider as an example the cross-sell on the web scenario of U+ Bank. The bank shows a personalized credit card offer to eligible customers when they log in to the banks' website.



Standard card

0% APR and no annual fee

[Learn more](#)

When a customer is eligible for multiple credit cards, the prediction calculates the propensity to receive a positive response from the customer for each card.

Customer Decision Hub decides which credit card to offer based on business rules, interaction context, and predictions.

A predictive model drives every prediction. The predictive model that is active in the **Predict Web Propensity** prediction is the **Web Click Through Rate** model.

Supporting models

Name	Component name	Type	Performance	Status
Web_Click_Through_Rate	Web_Click_Through_Rate_Customers	Adaptive model	73.87 AUC	ACTIVE

This is an adaptive model that continuously learns from customer interactions. It adapts to changes in customer behavior over time.

You can configure several aspects of a prediction.

A control group is a small group of customers who receive random offers, as opposed to the test group.

Customers in the test group receive the offers that they are most likely to accept, based on the propensity that the prediction calculates for each customer.

The purpose of the control group is to calculate lift by comparing the success rate in the control group with the success rate in the test group.

The random offers also allow for continuous exploration of the predictive models.

Based on the lift, you can determine the effectiveness of your prediction, for example, in increasing conversion rates.

The control group is typically defined as 2% of all customers, but this can be changed.

The response labels represent the possible outcomes of a prediction.

The propensity is computed based on the number of outcomes registered under the target label versus the alternative label.

For example, in **Predict Web Propensity**, as you want to predict the likelihood of a customer clicking a banner, the Target label (which in this case represents the positive outcome) is **Clicked**.

The alternative label that represents a negative outcome in this case, is **NoResponse**.

The NoResponse response can be captured on request or automatically depending on the response time-out setting.

The response time-out defines how long to wait for the customer to respond to your offer.

After the specified amount of time elapses, the system automatically records the alternative outcome for the interaction.

Response timeout

You can choose how long you want to wait for a response.
If this period elapses, the alternative label will be recorded.

Propensity to Click

Indefinitely Fixed time frame

Amount Time unit

10 minutes ▾

This setting depends on the use case.

When predicting a click on a web banner, you typically set it at 30 minutes or less, but in an outbound email offer, a waiting time of several days is more appropriate.

The **Web Click Through Rate** model that drives the prediction is the template for a model for each of the credit card offers.

You can monitor the performance of the models in a diagram that shows the success rate versus model performance.

Colored circles represent the models. The size of the circles indicates the number of responses captured by the model.

Suppose that U+ Bank wants to introduce proactive retention by showing a retention offer instead of a credit card offer to customers that are about to leave the bank.

That configuration requires a prediction that aims to calculate the churn risk for each customer, ideally based on the banks' historical data.

Pega supports creating three types of predictions: predictions that optimize the engagement with customers, predictions that improve business processes, and predictions that analyze the text of incoming messages.

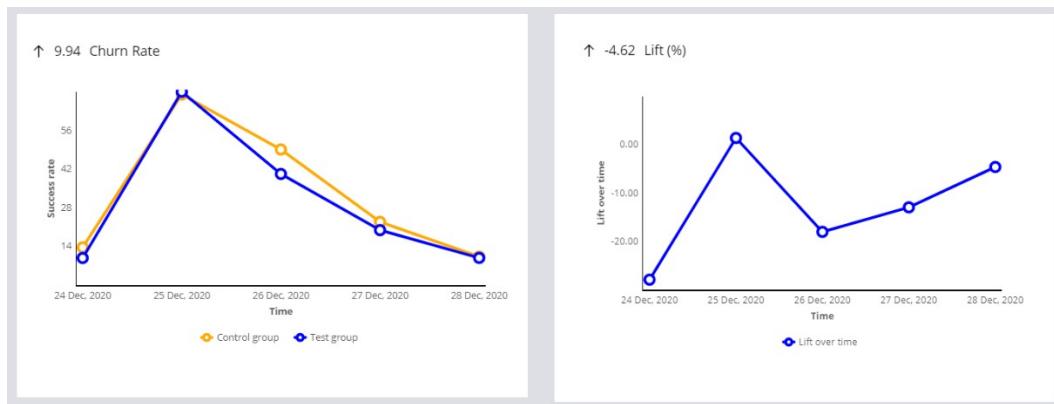
Create a prediction

Where will you be using the prediction?

- Customer Decision Hub
Optimize the engagement with your customers
- Case management
Use predictions to improve the automation in cases
- Text analytics
Analyze the text that comes through your channels

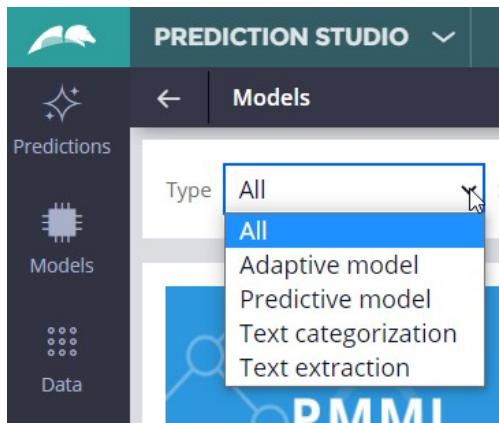
To analyze how successful the prediction is, monitor it over time.

Available metrics are the success rate, the lift calculated using the control group, the prediction performance, and the total number of cases.



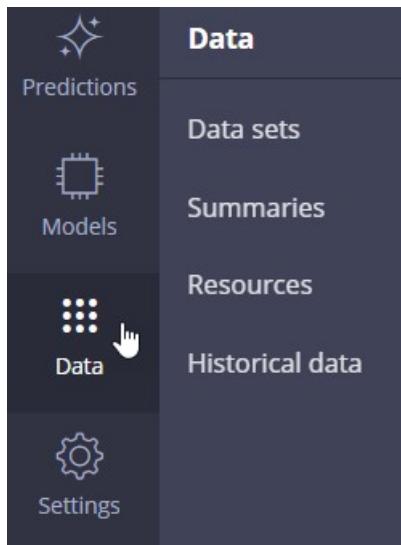
Prediction Studio generates actionable insights and notifies the user when predictions and predictive models show unexpected behavior (for example, a significant drop in success rate).

On the **Models** landing page, all available models are listed. These are adaptive models, predictive models, and text categorization and extraction models.



Resources such as data sets and summaries are available from the **Data** tab.

When enabled, historical data that is used for the training of adaptive models and monitoring of predictive models is recorded for offline analysis.



Under **Settings**, you can review the company policies regarding model transparency.

Model transparency is becoming an important requirement for many businesses.

In risk management, decisions must be explainable. In marketing, more accurate models may be allowed at the expense of transparency.

The lead data scientist sets the transparency thresholds for different business issues.

Each model type is assigned a transparency score ranging from 1 to 5, where 1 means that the model is opaque, and 5 means that the model is transparent.

Adaptive model Pega Compliance All business issues	Bivariate model Pega Compliance All business issues	Genetic algorithm Pega Compliance All business issues
Clustering model Compliance All business issues	Ensemble model Compliance All business issues	General regression Compliance All business issues
Neural network Compliance All business issues	Random forest Compliance All business issues	Scorecard Compliance All business issues

Depending on these settings, some types of models can be non-compliant for a specific business issue.

You have reached the end of this video. What did it show you?

- How you can manage the lifecycles of predictions and the predictive models that drive them in Prediction Studio
- How to examine the model transparency setting of the business

Enabling AI in the NBA framework

Description

Learn how every Next-Best-Action weighs customer needs against business objectives to optimize decisions based on priorities set by the business. Configure the AI model for the web treatment with additional customer data as potential predictors. Confirm that the model learns from customer behavior.

Learning objectives

- Describe how AI uses customer behavior to calculate propensity
- Explain the $(P*C*V*L)$ prioritization formula
- Configure additional potential predictors for an adaptive model

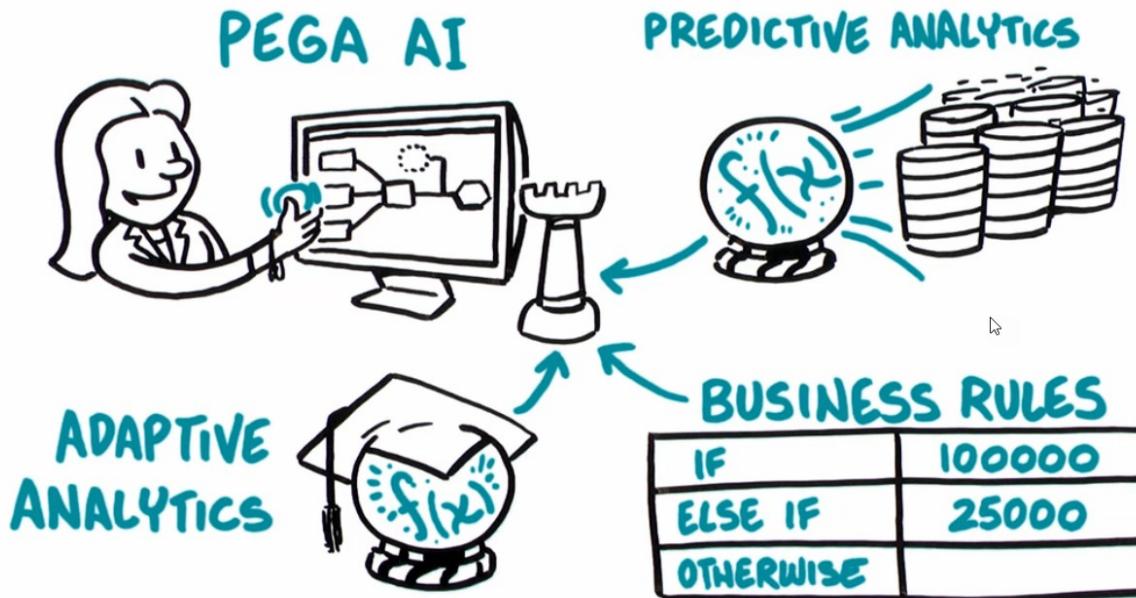
Adaptive analytics

The effectiveness of adaptive models

Applying simple business rules to a sales strategy enables you to identify eligible actions for a customer. However, business rules alone will not enable you to select the best action for a customer, or the action the customer is most likely to accept. As a result, action acceptance rates can be low when only business rules are used to make sales decisions.

Please be aware that actions can be renamed and are sometimes referred to as propositions or offers.

To improve acceptance rates, augment the business rules in a decision strategy with analytics.



Applying adaptive analytics to your decision strategies enables the strategies to detect changes in customer behavior in real time so that you can act on the changes immediately.

Pega Adaptive Decision Manager

Pega Adaptive Decision Manager (ADM) is a component that allows you to build self-learning adaptive models that continuously improve predictions for a customer. ADM can automatically detect changes in customer behavior and act on the changes in real time, which enables business processes and customer interactions to adapt instantly to the changing interests and needs of customers.

Adaptive decisioning continuously increases the accuracy of its decisions by learning from each response to an action. For example, if a customer is offered and then accepts a product, the likelihood that customers with a similar profile also accept that offer increases slightly. The mathematical expressions of these probabilities in the model are regularly updated.

ADM is a closed-loop system that automates the model creation, deployment, and monitoring process. The component can manage a large number of models without human intervention.

In contrast to predictive analytics, which requires historical data and human resources to develop a reliable predictive model, adaptive decisioning can start to calculate who is likely to accept or reject an offer without using any historical information, learning on the fly. Adaptive decisioning captures and analyzes response data in real time, which is useful in situations where the behavior itself is volatile. A typical use case is the real-time detection of complex fraud patterns or predicting customer behavior following the introduction of a new offering.

You can use predictive models as an alternative to, or in conjunction with, cases where data is available for offline modeling.

Adaptive decisioning creates binary models and uses these models for predictions. The full adaptive modeling cycle consists of the following steps:

1. Capture response data in real time from every customer interaction.
2. Regularly:
 - a. Use sophisticated auto-grouping to create coarse-grained, statistically reliable numeric intervals or sets of symbols.
 - b. Use predictor grouping to assess inter-correlations in the data.
 - c. Use predictor selection to establish an uncorrelated view that contains all relevant aspects of the action.
 - d. Use the resulting, statistically robust adaptive binary model for scoring customers.
3. Whenever new data is available, update the data model.

Adaptive decisioning can also build channel-specific models that account for differences in customer responses to outbound versus real-time inbound offers.

Predictors and outcomes of an adaptive model

Predictors

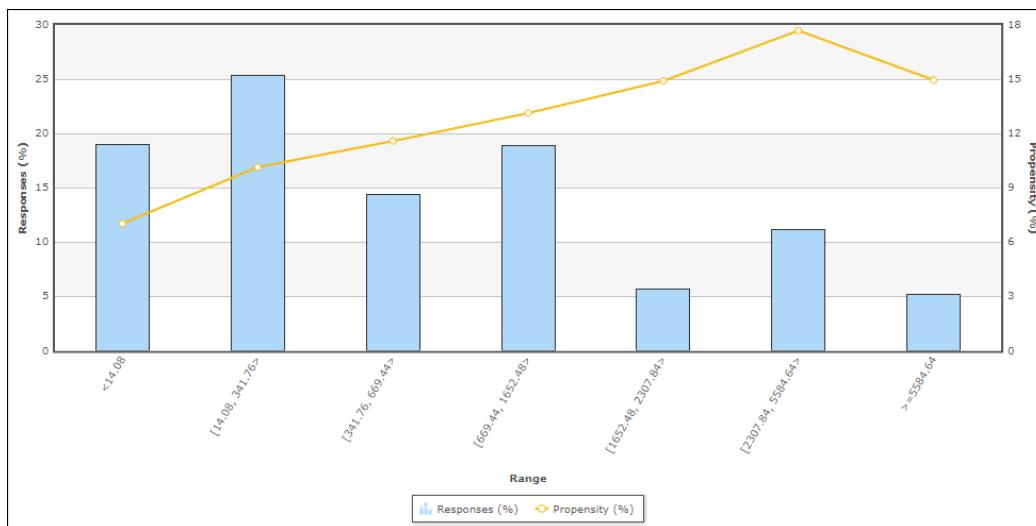
Configuring adaptive models involves selecting potential predictors and setting outcomes that identify positive and negative customer behavior. Unless you are a highly experienced data scientist, it is strongly recommended to leave the advanced settings at their default.

The input fields you select as predictor data for an adaptive model play a crucial role in the predictive performance of that model. A model's predictive power is at its highest when you include as much relevant, yet uncorrelated, information as possible. In Pega, it is possible to make a wide set of candidate predictors available, as many as several hundred or more.

Adaptive Decision Manager (ADM) automatically selects the best subset of predictors. ADM groups predictors into sets of correlated predictors and then selects the best predictor from each group, that is, the predictor that has the strongest relationship to the outcome. In adaptive decisioning, this predictor selection process repeats periodically.

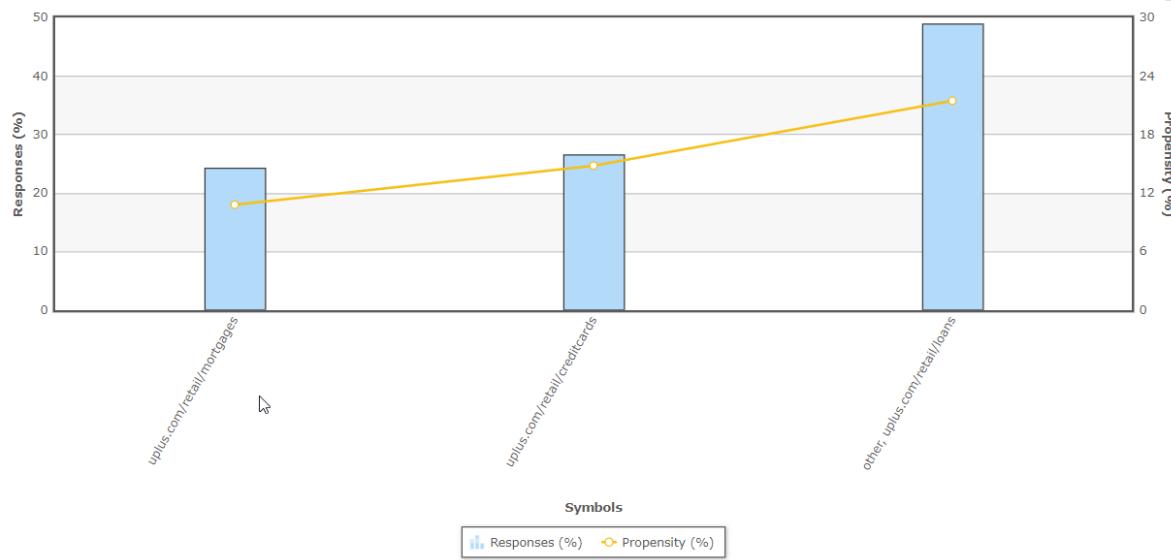
You can use several data types in adaptive analytics, including:

Numeric data - Basic numeric data such as age, income, and customer lifetime value can be used without any preprocessing. Your model automatically divides that data into relevant value ranges by dynamically defining the bin boundaries.



Symbolic data - You can feed predictors with up to 200 distinct string values without any preprocessing. Such data is automatically categorized into relevant value groups, such as

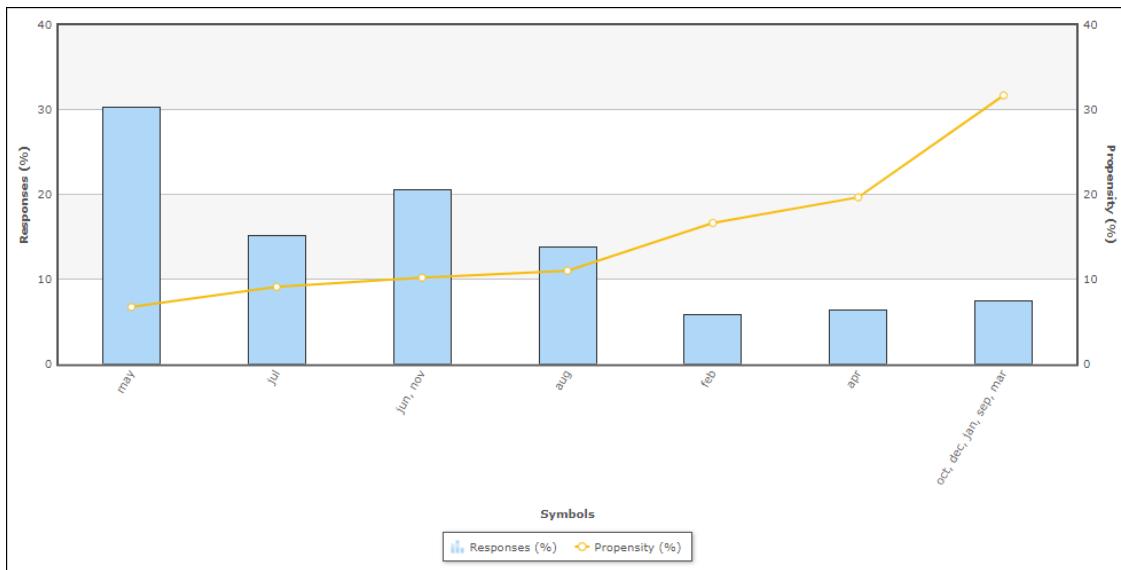
the **PreviousWebpage** predictor in the following example. For predictors with more than 200 distinct values, group the data into fewer categories for better model performance.



Customer identifiers - Customer identifiers are symbolic or numeric variables that have a unique value for each customer. Typically, they are not useful as predictors, although they might be predictive in special cases. For example, customer identifiers that are handed out sequentially might be predictive in a churn model, as they correlate to tenure.

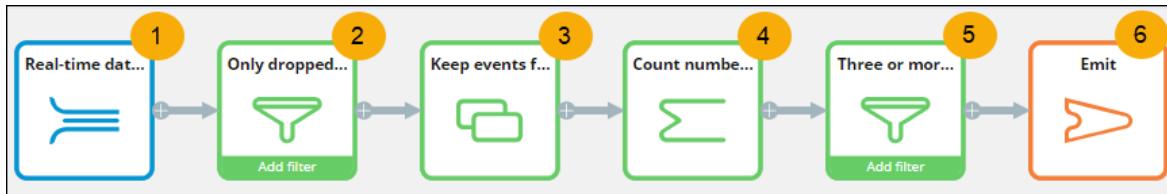
Codes - For meaningful numeric fields, feed code fragments to the model as separate predictors. Simple values require only basic transformation. For example, you can shorten postal codes to the first 2 or 3 characters which, in most countries, denote geographical location.

Dates - Avoid using absolute date/time values as predictors. Instead, take the time span until now (for example, derive age from the DateOfBirth field), or the time difference between various pairs of dates in your data fields (such as the DurationLastSubscription field). Additionally, you can improve predictor performance by extracting fields that denote a specific time of day, week, or month.



Text - Do not use plain text to create predictors without any preprocessing; it contains too many unique values. Instead, extract values such as intent, topic and sentiment to use as predictors. Pega features a Text Analyzer rule for this purpose.

Event streams - Do not use event streams as predictors without preprocessing, aggregate the data instead. Pega features event strategies for this purpose. As an example, this event strategy detects dropped calls.



First, (1) it listens to a real-time dataset; then (2) it filters out dropped customer calls; next (3) it stores the terminated calls for one day; (4) it counts the number of terminated calls within the one-day timeframe; and (5) it creates an event if three calls are terminated within the one-day timeframe; lastly, (6) it emits the event. The aggregates can be stored and used like any other symbolic or numeric field.

Interaction History - Past interactions are usually very predictive. You can use the Interaction History (IH) to extract fields such as the number of recent purchases, the time since last purchase, and so on. To summarize and preprocess IH data for predictions, use IH summaries. Several predictors based on IH summaries are enabled by default (and require no additional setup) for all new adaptive models. These are the group that was referenced in the last interaction, the number of days since the last interaction, and the total number of interactions.

Multidimensional data - For models that inform the initial customer decision, things such as lists of products, activities, and transaction outcomes are useful sources of information for predictors. Use your intuition and data science insight to determine the possibly relevant derivatives, for example, number-of-products, average-sentiment-last-30-days, and so on.

Interaction context - To increase the efficiency and performance of your models, do not limit the data to customer data alone. By supplementing decision process data with the interaction context, you can adjust the predictions for a customer and provide different outcomes depending on their context. Contextual data might include the reason for a call, or the way the customer uses the website or mobile app to interact with the company, etc.

Customer behavior and usage - Customer behavior and interactions, such as financial transactions, claims, calls, and complaints, are typically transactional in nature. From an adaptive analytics perspective, you can use that data to create derived fields that summarize or aggregate this data for better predictions. Examples of this type of data include average length of a call, average gigabyte usage last month, and the increase or decrease in usage over the last month compared to previous months.

Model scores - Scores from predictive models for different but related outcomes as well as other data science output might be predictive as well. If you decide to use scores as predictors in your models, evaluate whether the models that include such a score perform better at the model level by verifying the area under the curve (AUC) and success rate metrics.

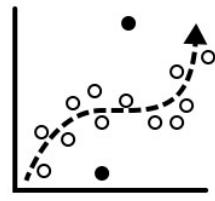
In summary, to achieve the best results, use predictors that provide data from many different sources, including:

Customer profile data such as age, income, gender, and current product subscriptions. This information is usually part of the Customer Analytic Record (CAR) and is refreshed regularly.

Interaction context data such as recent web browsing information, call reasons, or input that is gathered during a conversation with the customer. This information can be highly relevant and, therefore, very predictive.

Customer behavior data such as product usage or transaction history. The strongest predictors of future behavior typically contain data about past behavior.

Model scores, which are scores derived from the off-line execution of external models.



Outcomes

The responses that indicate positive or negative behavior must be identified. When predicting the click-through rate for a web banner, the default value for positive behavior is **Clicked** and the default value for negative behavior is **NoResponse**.

Applications may use different words to identify positive or negative behavior, for example, **Accepted** may be identified as positive behavior and **Rejected** may be identified as negative behavior. You can add these values when needed.

Positive outcome ⓘ

Add outcome

Clicked

Negative outcome ⓘ

Add outcome

NoResponse

Accepted

Rejected

Advanced settings of an adaptive model

Default values

The default values for the adaptive model advanced settings are based on best practices and should only be changed by a highly experienced data scientist.

Update frequency and scope

When a model is updated, Prediction Studio re-trains the model with a specified number of responses. You can set the number of responses that will trigger the update.

Model update frequency

Update model after every
500 responses

You can also set the scope of the update. By default, all responses received during each update cycle are used. If you want to assign more weight to recent responses when updating a model, use a subset of the responses.

When updating a model
 Use all responses
 Use subset of responses

0 weighted last responses

By default, all historical data is used to monitor the performance of the model. If required, model performance can be monitored for the most recent responses.

Monitor performance for the last
500 weighted last responses

Grouping

The default values for **Grouping granularity** (the granularity of predictor binning) and **Grouping minimum cases** (the minimum percentage of cases per interval) are based on best practices and should not be changed casually.

Data analysis binning

Grouping granularity

0.25

Grouping minimum cases

0.05

The higher the value for **Grouping granularity**, the more bins are created. This value represents a statistical threshold that indicates when predictor bins with similar behavior are merged.

The **Grouping minimum cases** setting controls how predictor grouping is established. Higher values result in a decreasing number of groups, which can be used to increase the robustness of the model. Lower values result in an increasing number of groups, which can be used to increase the performance of the model.

The selection of the *active* predictors is guided by thresholds for predictor performance and the correlation between predictors.

Predictor selection

Activate predictors with a performance above

0.52 AUC

Group predictors with a correlation above

0.8

The performance of a predictor is measured as the area under the curve (AUC). A higher value results in fewer predictors in the final model. The minimum AUC value is 0.5, therefore the value of the performance threshold should always be set to at least 0.5.

The value for the correlation between predictors determines when predictors are considered similar, and only the best of those predictors are used for adaptive learning. The measure is the correlation between the probabilities of positive behavior within pairs of predictors.

Adaptive model outputs

Model outputs

Adaptive models produce four outputs: Propensity, Evidence, Performance, and Positives.

Propensity is the predicted likelihood of positive behavior, for example, the likelihood of a customer accepting an offer. The propensity for every action starts at 0.5 or 50% (the same as a flip of a coin) because in the beginning, the model has no response behavior on which to base its predictions.

Evidence is the number of responses used in the calculation of the Propensity.

Performance is how well the model can differentiate between positive and negative behavior. Again, the initial value is 50%, with 100% being perfect performance. As a result, the performance value is somewhere between 50 and 100.

Positives is the number of positive outcomes that has been received by the model.

Mapping

In strategies, model propensity is automatically mapped to the strategy property called `.pyPropensity`. There is no automatic mapping for the Evidence, Performance or Positives outputs, but a strategy designer can manually map the outputs to any of the strategy properties under the **Output mapping** tab.

The screenshot shows a user interface for mapping adaptive model outputs. At the top, there are three tabs: 'Source components', 'Adaptive model', and 'Output mapping'. The 'Output mapping' tab is currently selected. Below the tabs, there is a section titled 'Default mapping' with the sub-instruction: 'Component sets .pyPropensity equal to the propensity of the adaptive model.' Underneath this, there is a checked checkbox labeled 'Enable additional mapping'. This checkbox is followed by three mapping entries, each consisting of a 'Set' field and a 'equal to' field. The first entry is 'Set .ModelEvidence' and 'equal to Evidence'. The second entry is 'and .ModelPerformance' and 'equal to Performance'. The third entry is 'and .ModelPositives' and 'equal to Positives'. The entire interface is contained within a light blue rectangular box.

Configuring an adaptive model

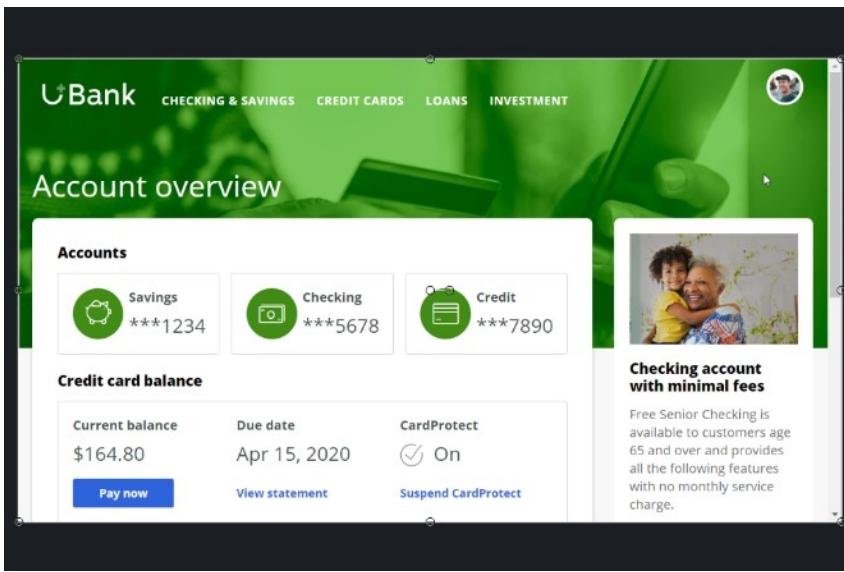
Introduction

Explore how AI-based arbitration works. Use adaptive models to leverage a website as a marketing channel by recommending more relevant banner ads to customers when they visit their personal portal. Use the Pega Prediction Studio to configure the out-of-the-box adaptive model for the web treatment.

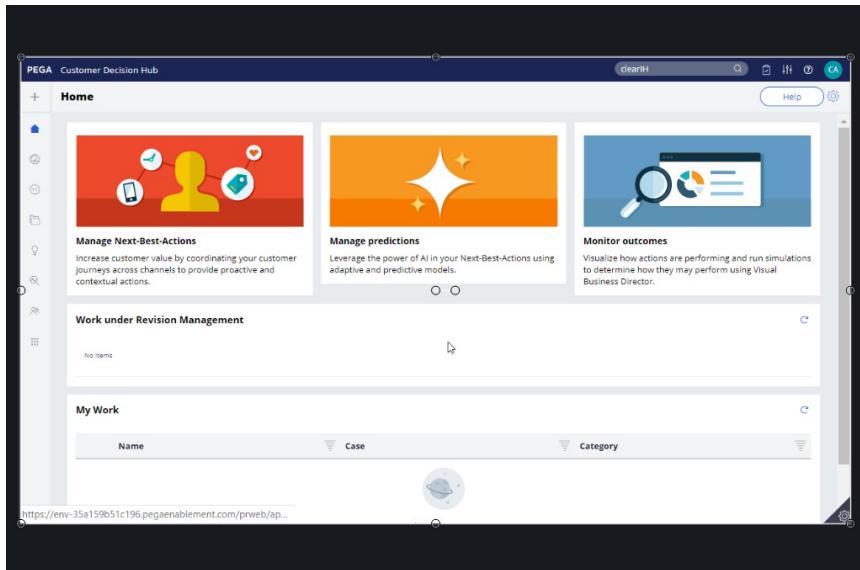
Transcript

This demo will explore how AI-based arbitration works and show you how to configure an adaptive model.

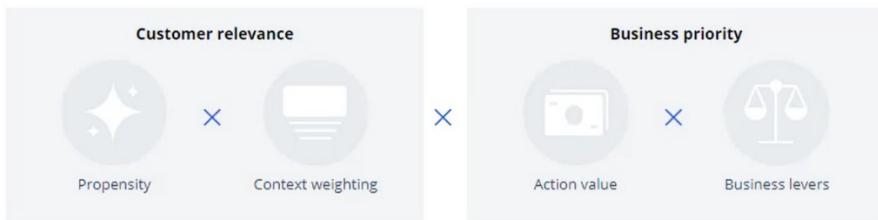
Currently, the U+Bank website shows a banner containing the running marketing offer that all customers see when they log in to the U+Bank website.



The bank would like to move forward and display more relevant offers to customers based on their behavior. To that end, the bank will use Pega Customer Decision Hub™ to display personalized marketing offers on the website.



These are the arbitration settings defined in Pega Customer Decision Hub's Next-Best-Action Designer.



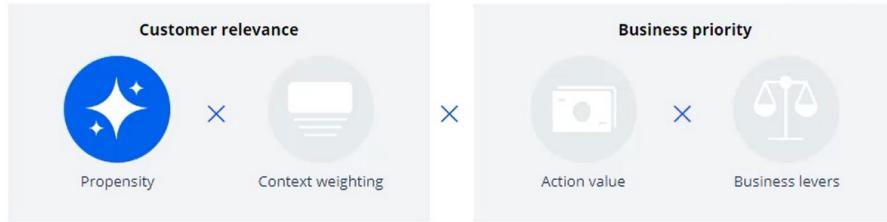
Arbitration aims to balance customer relevance with business priorities to decide which offer to show to the customer.

To achieve this balance, propensity (P), context weighting (C), action value (V), and business levers (L) are represented by numerical values and plugged into a simple formula, $P * C * V * L$.

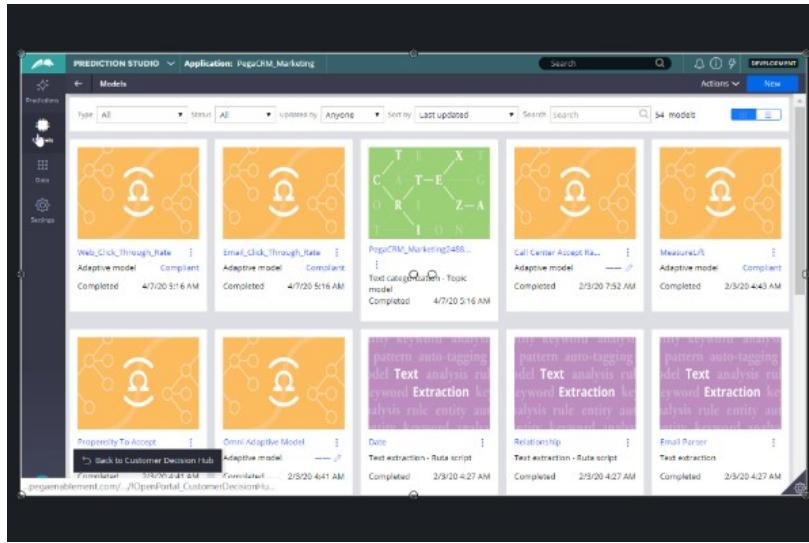
This formula is used to arrive at a prioritization value, which is used to select the top actions.

Propensity is the predicted likelihood of positive behavior, such as the likelihood of a customer accepting an offer. The propensity value is calculated using AI.

How is this AI configured? First, let's activate propensity in the arbitration formula to allow the AI to weigh in.



The AI is driven by an adaptive model. All predictive models in the system are available on the Model landing page of Prediction Studio.



An adaptive model is a self-learning predictive model that uses machine learning to calculate propensity scores. It automatically determines the factors that help in predicting customer behavior.

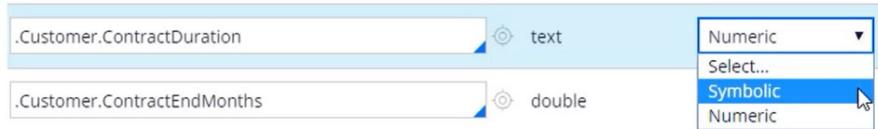
The Customer Decision Hub is configured to calculate the propensity for each treatment. For the web treatment, **Web_Click_Through_Rate** is the out-of-the-box adaptive model. It predicts the likelihood of positive customer behavior, such as clicking on a banner.

The model is used by the NBA Designer-generated decision strategy to determine a propensity for each treatment. After the initial application setup by a decisioning architect, the **Web_Click_Through_Rate** model is automatically generated.

All customer fields are set up as potential predictors by default. The model can be enhanced by adding additional fields.

Predictors can be one of two types: numeric or symbolic. The system will default the predictor type to the property type during the initial set-up, but you can change this.

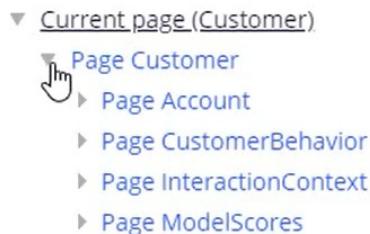
When you know a numeric predictor has a small number of distinct values, for example, when the contract duration is either 12 or 24 months, change the predictor type from numeric to symbolic.



Bear in mind that changing the predictor type effectively means removing and adding a predictor. Therefore, best practice is to make these changes early in the process, as there is no way to retain previous responses. Pega only stores binned aggregates.

However, this usually has little impact, as models typically learn quickly given the response rates in most applications.

It is highly recommended to add many uncorrelated predictors, as the models will figure out which ones to use. Additional predictors may include customer behavior, contextual information, past interactions with the bank, and even scores from external models.



Customer responses are recorded in the Interaction History dataset. Aggregated fields are automatically provided to the models as predictors. An example of such a predictor is the group to which the most recently accepted web offer belongs.

[Fields \(58\)](#) [Parameters \(0\)](#) [IH Summaries \(Enabled\)](#)

Aggregated fields from interaction history summaries are automatically provided to the models as predictors.
IH summaries leverage historical customer interactions to improve the predictions. [?](#)

Predictors based on interaction history summaries are [Enabled](#) ▾

IH - All time for each Subject ID, Subject Type, Channel, Direction, Outcome

Predictor	Aggregate	Field from interaction history
IH.{Channel}.{Direction}.{Outcome}.pxLastGroupID	Last	pyGroup
IH.{Channel}.{Direction}.{Outcome}.pxLastOutcomeTime.DaysSince Last		pxOutcomeTime
IH.{Channel}.{Direction}.{Outcome}.pyHistoricalOutcomeCount	Count	

By default, the **action group**, the **time since the last accept**, and the **number of accepted actions** from that group are available as predictor information.

Next, check the positive and negative outcomes of the adaptive model.

When a customer clicks the web banner, it is recorded as **Clicked** and considered positive behavior, because the customer has shown interest in the offer.

If the banner is shown, but the customer does not click it within a given time period, it is recorded as **NoResponse**, which is considered negative behavior.

The screenshot shows two sections: 'Positive outcome' and 'Negative outcome'. Each section has an 'Add outcome' button and a list box containing one item. The 'Positive outcome' list box contains 'Clicked' and the 'Negative outcome' list box contains 'NoResponse'. Both list boxes have a trash can icon at the bottom right.

The default values for the advanced settings in an adaptive model are based on best practices and should only be changed by a highly experienced data scientist.

The **model update frequency** represents the number of responses an adaptive model will accumulate before it starts updating the model.

The screenshot shows a configuration screen for 'Model update frequency'. It has a text input field labeled 'Update model after every' with the value '500' and the unit 'responses'.

There are additional parameters used for binning responses.

The screenshot shows a configuration screen for 'Data analysis binning'. It has two text input fields: 'Grouping granularity' with the value '0.25' and 'Grouping minimum cases' with the value '0.05'.

You can now save the model.

Now that the AI is configured, when Troy logs into the U+Bank website the Standard Card credit card offer is presented instead of the generic offer.

The screenshot shows a summary card for the AI model. It includes the following information:

Treatment	Standard card tile
0% APR and no annual fee	
Category	
Propensity	0.5
Priority (P*C*V*L)	1

Note the propensity and priority values for the **Standard Card**. The propensity for every action starts at 0.5 or 50%, the same as the flip of a coin. This is because in the beginning, the AI has no past customer behavior on which to base its predictions.

Notice also that although only propensity is currently enabled for arbitration, the priority value does not match the propensity value.

The AI uses Thompson sampling to add some noise to the propensity values. If the models have received very few responses, a lot of noise is added.

When the number of responses received increases, less noise is added and therefore the propensities returned are very close to the raw model propensity.

If Troy doesn't click on the current offer this time, a different offer will be shown the next time he visits the website. The next offer Troy is eligible for, the **Rewards Card**, is then selected for display.

If Troy ignores this offer as well (by not clicking on it) then the next time he logs in, the **Standard Card** offer will be displayed again.

Why this behavior? First, Troy only qualifies for two credit card offers. Second, the AI model is configured to record an impression when a customer is presented with an offer but doesn't click on it. This is considered negative behavior.

As a result, the propensity, and therefore the priority, of the not-clicked-on offer decreases. Notice that the propensity value of the **Standard Card** offer dropped from 0.5 to 0.25.

Treatment	Standard card tile 0% APR and no annual fee
Category	
Propensity	0.25
Priority (P*C*V*L)	0.625

Now, if Troy clicks on the **Learn more** link for the **Standard Card** offer, a positive response is recorded, and thus the propensity value of the **Standard Card** increases.

Treatment	Standard card tile 0% APR and no annual fee
Category	
Propensity	0.375
Priority (P*C*V*L)	0.53125

This demo has concluded. What did it show you?

- How AI uses customer behavior to calculate propensity.
- How propensity smoothing is used to jump-start AI learning.

- How the prioritization value is calculated using the (P*C*V*L) formula.
- How to enable AI on the web channel.
- How to configure additional potential predictors for an adaptive model.

Monitoring adaptive models

Description

It is a regular data scientist task to inspect the health of the adaptive models and share the findings with the business. The predictive performance and success rate of individual adaptive models provide information that can help business users and decisioning consultants to refine the next best actions of the company. Learn how to monitor the performance of the adaptive models and measure the lift that the system produces.

Learn how to export the raw data that adaptive models have processed to inspect and validate the predictors.

Learning objectives

- Describe how predictions add best practices to predictive models
- Explain how the use of a model control group allows the measurement of lift
- Name the key metrics of adaptive models visualized in the bubble chart
- Inspect individual active and inactive predictors
- Explain how predictors with similar predictive performance are grouped
- Examine the propensity distribution and the trend for the whole model
- Export the raw data that is used by adaptive models

Measuring lift using a control group

Introduction

The boost in success rate achieved with adaptive models can be measured using a control group in Predictions. Predictions are strategies that add best practices to predictive models. Customers in the control group will receive a random offer instead of the one recommended by the AI. This allows a comparison between the control group and the rest of the audience. It also enables the models to explore alternative outcomes.

Transcript

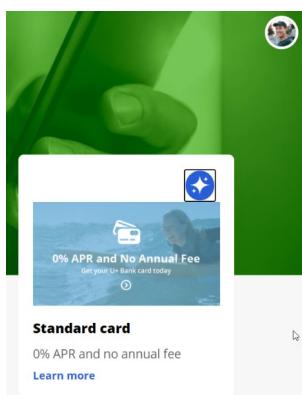
This demo shows you how to use a control group to measure the boost in success rate the adaptive models achieve.

U+ Bank is a small retail bank. When customers log in to the U+ Bank website, they see the credit card offers for which they qualify based on the engagement policy defined by the business.

When customers qualify for multiple credit card offers, adaptive models decide which is the best offer to show.

Adaptive models are self-learning and will automatically learn from customer interactions.

The models interpret a click on the offer banner as positive behavior. When a customer ignores the banner, this is interpreted as negative behavior. Models are updated frequently.



An important business metric is the lift in success rate that the models generate.

To measure and report on this metric, a data scientist can use Predictions. Predictions add best data science practices to predictive models.

One of these is to use a control group as a benchmark for measuring the lift that the models achieve.

All out-of-the-box predictions available are listed on the Predictions landing page in Prediction Studio.

The screenshot shows the Pega Prediction Studio interface with the application set to "PegaCRM_Marketing". On the left sidebar, there are navigation links for "Predictions", "Models", "Data", and "Settings". The main area displays three prediction cards:

- Predict Web Propensity**: A line chart showing performance over time. The Y-axis is labeled "Performance" with values from 0 to 100. The X-axis shows dates from Nov 12, 2020. The AUC is listed as 82.61. Below the chart, it says "Decision Request" and "Open prediction".
- Predict Outbound SMS Propensity**: Shows a blue starburst icon. It says "No performance available yet" and "Decision Request". Below it, "Open prediction" is shown.
- Predict Outbound Retail Propensi...**: Shows a blue starburst icon. It says "No performance available yet" and "Decision Request". Below it, "Open prediction" is shown.

Below these cards, there is a section titled "All predictions" with a table:

Name	Subject	Outcome	Performance (AUC)
Predict Web Propensity	Decision Request	Clicked	↓ 82.61
Predict Outbound SMS Propensity	Decision Request	Clicked	—

The **Predict Web Propensity** prediction monitors the responses from the U+ Bank website.

The outcome of this prediction is the propensity to click on a web banner.



A percentage of the total number of customers is randomly reserved for the control group.

The customers in the control group receive a random action instead of the action that the AI recommends.

A small percentage is sufficient to measure lift. By default, the control group is set to reflect 2% of the population, but you can change that value if desired.

Control group

The control group is used to measure lift by comparing the success rate in the target group with the control group. Customers in the control group will receive an action determined by a random propensity.

Percentage Field

Percentage
2.0 %

Alternatively, customers can be appointed to the control group based on a customer attribute.

Control group

The control group is used to measure lift by comparing the success rate in the target group with the control group. Customers in the control group will receive an action determined by a random propensity.

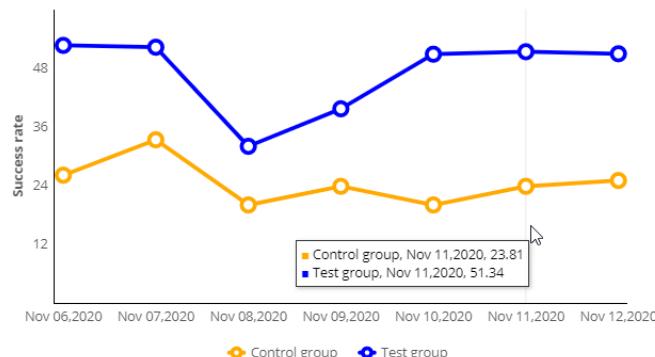
Percentage Field

Field
.ControlGroup equal to Yes

In the first chart, the yellow line shows the success rate, such as the accept rate, or, in the case of the **Predict Web Propensity** prediction, the click rate for the control group.

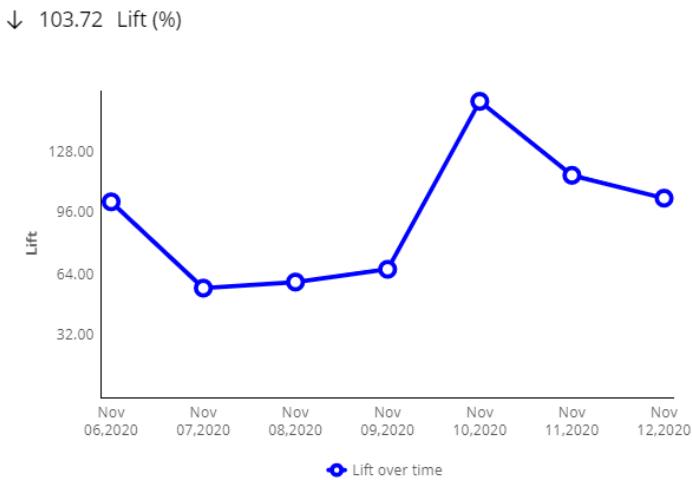
The blue line shows the success rate for all other customers, referred to as the test group.

↓ 50.93 Propensity to Click Rate



As customers in the control group are presented with a random offer, the success rate is expected to be lower than for the test group, for whom the offer is based on the propensity to click.

The second chart shows the lift over time, which is the difference between the blue and yellow lines, expressed as a percentage.



The use of a control group allows the measurement and monitoring over time of the efficiency of the machine learning process.

The control group also enables another data science best practice: mixing some exploration into the Prediction exploitation.

As the random offers made to the customers in the control group are not based on a particular customer profile, the models have a degree of freedom to explore.

In the notification section, two actionable insights related to lift are provided when applicable: absence of lift ...

... and a lift that is significantly lower than in the previous week.

❖ Predict Web Propensity

Outcome

Propensity to Click

Subject

Data_Decision_Request_Customer

⚙ Configure

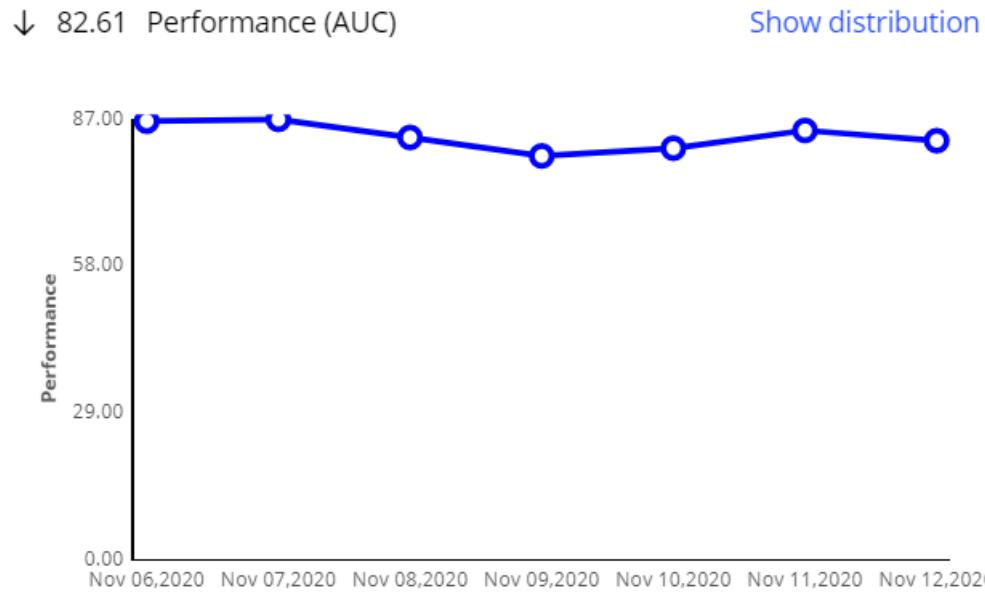
Insights

Lift has dropped by more than
10% compared to the previous
week

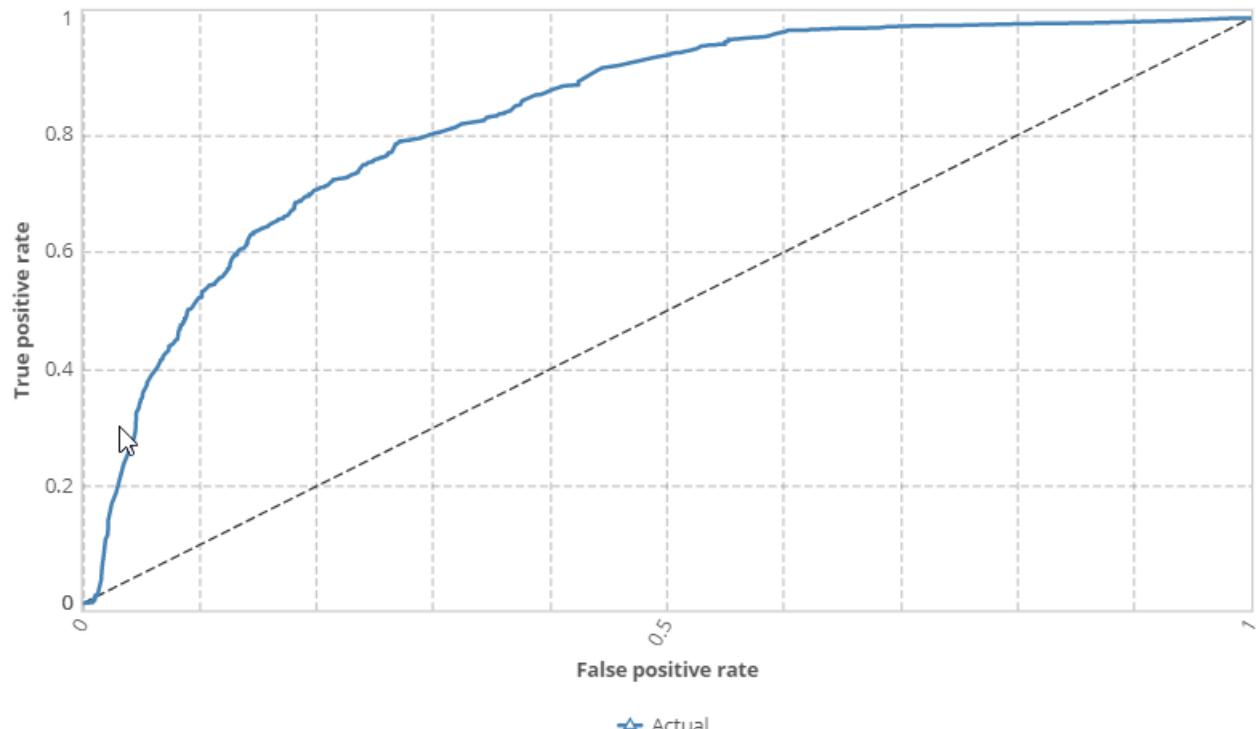
41 minutes ago

Both notifications prompt for an inspection of the model that drives the prediction as well as its predictors.

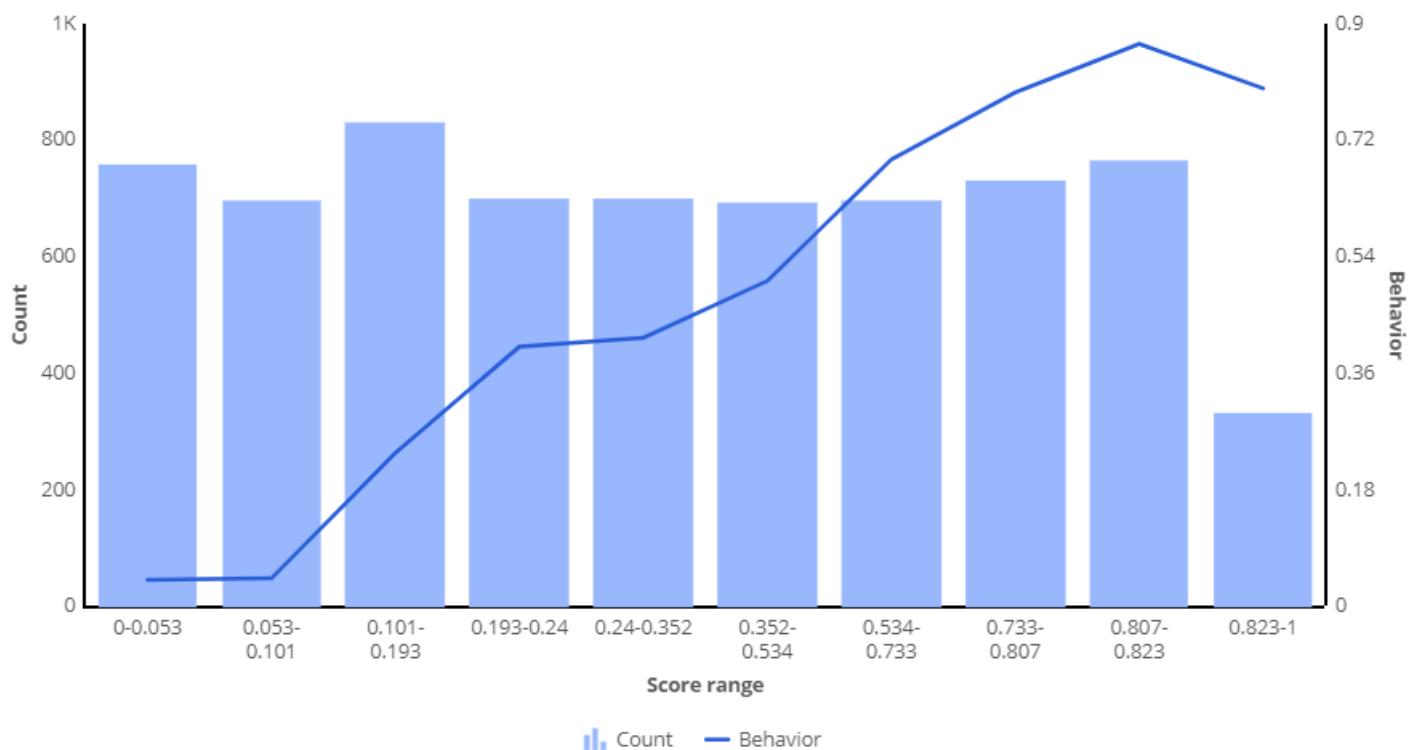
The third graph shows the performance of the model over time.



The ROC curve shows the true positive rate, or sensitivity, versus the false positive rate, or 1 minus the specificity, of the prediction.



Also, a decile distribution graph is provided for further inspection.



You have reached the end of this demo. What did it show you?

- How predictions add best practices to predictive models.
- How the use of a model control group allows the measurement of lift.
- How the use of a control group adds exploration to the exploitation of the models.

Regular monitoring of adaptive models

Regular monitoring of adaptive models

Adaptive models will learn from all customer interactions, adjusting to changing behavior over time. To confirm the continuing accuracy of your adaptive models, perform the following tasks regularly:

- Check the performance and success rate of your models every two weeks.
- Inspect predictors every two or three months.

The purpose of regular inspection is to detect factors that negatively influence the performance of the adaptive models and the success rate of the actions.

Identifying technical problems

Look for adaptive models with a success rate of zero. This means that the actions for these models do not have any positive responses.

Identifying actions for which the model is not predictive

Look for adaptive models with low performance. Consider adding additional data as predictors.

Identifying actions that have a low number of responses

Look for adaptive models with a low number of responses. Discuss the eligibility criteria set in the Next-Best-Action Designer with the business. Changing the exclusion settings may increase the number of responses.

Identifying actions that are offered so often that they dominate other actions

Look for adaptive models with a high number of responses. A high number of responses might be fine from the business point of view. However, if necessary, prioritization can be adjusted in the Next-Best-Action Designer.

Identifying actions with a low success rate

Look for adaptive models with a low success rate. If the model performance is high, the relevance to the customers is high, but the action is unattractive and should be discussed with the business.

Inspecting an adaptive model

Inspect your model after introducing a new action, adding or removing a predictor, or changing prioritization. Take note of the active and inactive predictors.

Inspecting predictors

Check the details of a predictor with a low performance score. A possible cause can be too many missing values for the predictor. Look at the top predictors and in the bins that have a particularly high or low success rate.

Identifying predictors that are never used

Because unused predictors have only a minor effect on model performance, you do not need to remove them from an adaptive model configuration; however, you can conduct an occasional cleanup as part of your maintenance activities. An unused predictor might still become relevant for a future action.

Inspecting adaptive models

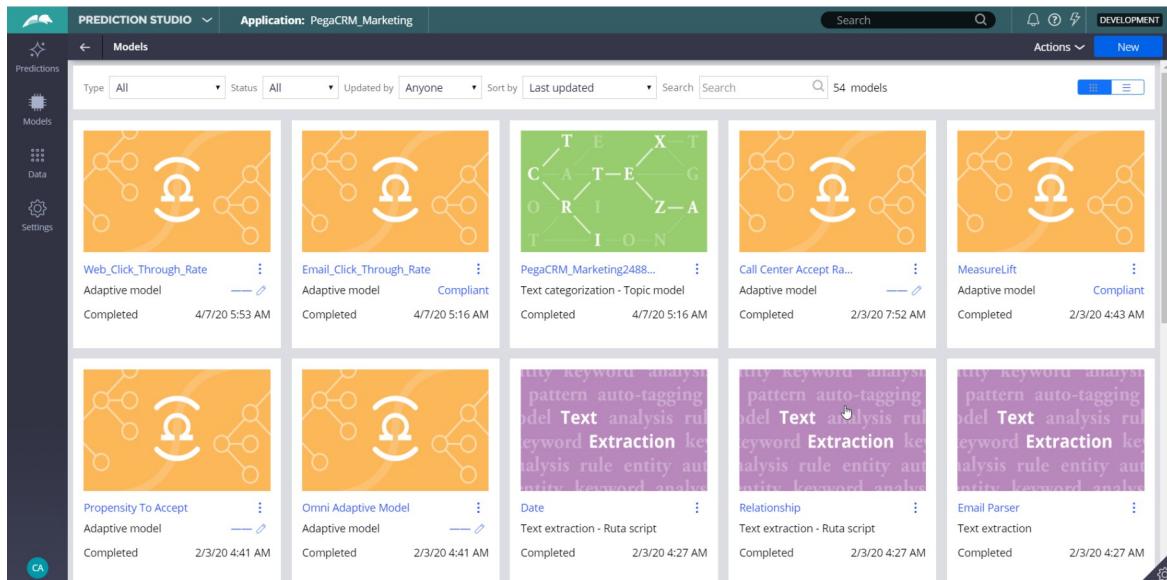
Introduction

The predictive performance and success rate of individual adaptive models provide information that can help business users and decisioning consultants to refine the Next-Best-Actions of the company. Monitoring of the health of adaptive models and their predictors is a regular data scientist task that can be performed in Prediction Studio.

Transcript

This demo will show how to inspect the health of your adaptive models and their predictors. This is a regular data scientist task.

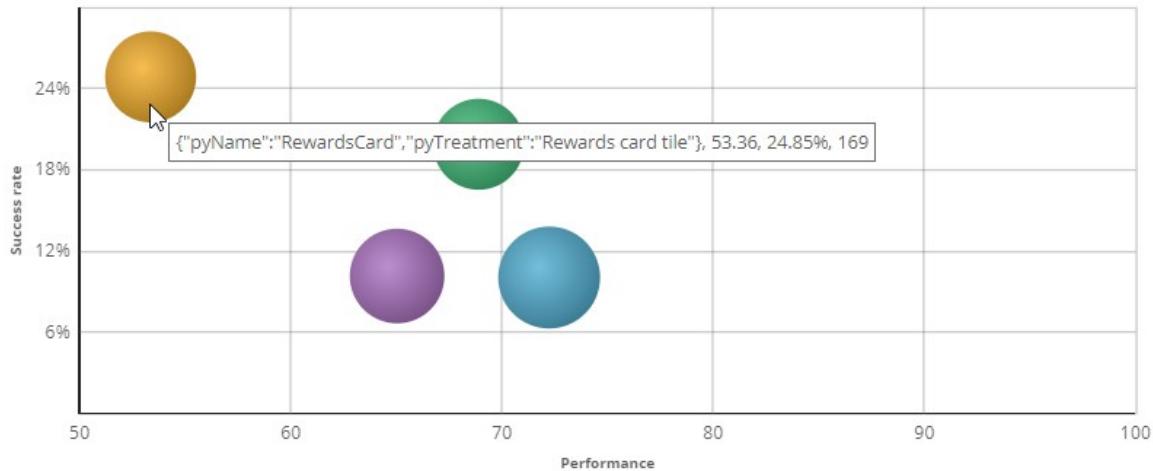
The predictive performance and success rate of individual adaptive models provide information that can help business users and decisioning consultants to refine the Next-Best-Actions of the company.



The screenshot shows the 'Models' tab in the Pega Prediction Studio interface. The top navigation bar includes 'PREDICTION STUDIO', 'Application: PegaCRM_Marketing', 'Search', 'Actions', and 'Development'. On the left, there are navigation icons for 'Predictions', 'Models', 'Data', and 'Settings'. The main area displays a grid of 10 adaptive models. Each model card includes a preview icon, the model name, type, status, and last updated time. The 'Web_Click_Through_Rate' model is highlighted with a purple background and has a tooltip showing its configuration details: 'Propensity To Accept', 'Adaptive model', 'Completed', '2/3/20 4:41 AM'. Other models include 'Email_Click_Through_Rate', 'PegaCRM_Marketing2488...', 'Call Center Accept Ra...', 'MeasureLift', 'Propensity To Accept', 'Omni Adaptive Model', 'Date', 'Relationship', and 'Email Parser'.

We will inspect the Web_Click_Through_Rate model, that calculates the propensity that a customer will respond positively to an offer made on the web channel.

The Monitor tab of an adaptive model configuration shows a bubble chart that visualizes the key metrics of all models generated.



Each bubble represents the model for a specific action.

The size of a bubble indicates the number of responses (positive and negative) to that action that have been used in the adaptive learning process.

In this example, there is a model for every action belonging to the Credit Card group.

When you hover the cursor over a bubble, you can view the name of the action, the performance, the success rate, and the number of responses.

In the legend, display of models can be toggled on and off.

- {"pyName": "PremierRewardsCard", "pyTreatment": "Premier Rewards card tile"}
- {"pyName": "RewardsCard", "pyTreatment": "Rewards card tile"}
- {"pyName": "RewardsPlusCard", "pyTreatment": "Rewards Plus card tile"}
- {"pyName": "StandardCard", "pyTreatment": "Standard card tile"}

The Performance axis indicates the accuracy of the outcome prediction.

The model performance is expressed in the Area Under the Curve (AUC) unit of measurement, which has a range between 50 and 100.

The higher the AUC, the better a model is at predicting the outcome.

The Success rate axis indicates the success rate expressed in percentages.

In this example, the success rate represents how often a web banner is clicked.

The system calculates this rate by dividing the number of times a banner is clicked by the total number of times the banner was shown on the website.

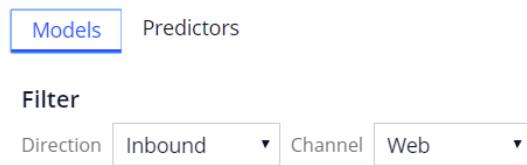
The information you see here is extracted from the Adaptive Data Mart, which is a reporting view of the Adaptive Decision Manager (ADM) server.

The Adaptive Data Mart is built automatically by a process running in the background. This process creates snapshots at regular time intervals.

This means that the data as reported may not be the same as the data you see when you view it in real-time. You can refresh the view to synchronize the data.

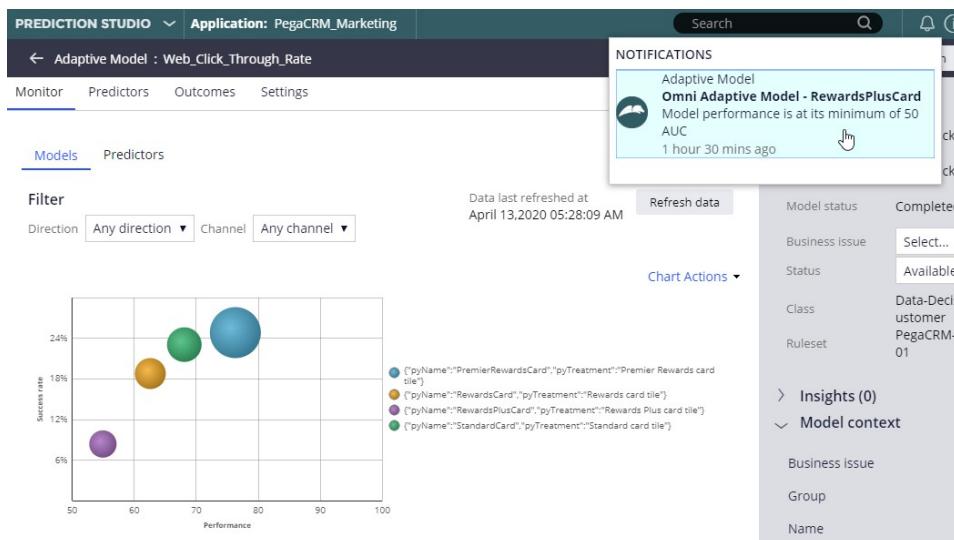
The model context includes the channel and direction, so you have different models for the Call Center, Email and Web channels, as well as for the inbound and outbound directions.

You can apply filtering to focus on models for a particular direction or channel, or a combination of the two.



A screenshot of the 'Predictors' tab in the 'Prediction Studio' interface. At the top, there are tabs for 'Models' (which is selected) and 'Predictors'. Below the tabs, there is a 'Filter' section with dropdown menus for 'Direction' (set to 'Inbound') and 'Channel' (set to 'Web').

Actionable insights are generated for individual models when the number of responses, model performance or success rate significantly changes over time.



On the Predictors tab, the number of models in which a predictor is active, and the performance of the predictor is displayed.

Predictor name	# Models active	# Models inactive
Customer.Age	4	0
Customer.AverageBalance	4	0
Customer.AverageSpent	4	0
Customer.CLV_VALUE	4	0
Customer.CreditScore	4	0
Customer.DebtToIncomeRatio	4	0
Customer.Gender	4	0
Customer.InteractionContext.PreviousWebpage	4	0
Customer.MonthlyPremium	4	0
Customer.NetPromoterScore	4	0
Customer.PrincipalLoan	4	0
Customer.RiskScore	4	0
IH.Web.Inbound.Clicked.pxLastOutcomeTime.DaysSince	4	0
Customer.HasMortgage	3	1

In this case, the Age predictor is used in all four models.

The HasMorgage predictor is active in three models and inactive in one model, where its predictive power is below the threshold.

The default value for this threshold is 52 percent.

Predictor selection

Activate predictors with a performance above

0.52 AUC

The system continuously monitors the predicting power of every predictor. If the predicting power of a predictor drops below the threshold value that predictor is deactivated.

The data that is used to visualize the models in the bubble chart is displayed in a table below the chart.

For each model number of responses, success rate and performance are shown.

From the adaptive model table, you can drill down into a model report for a specific adaptive model.

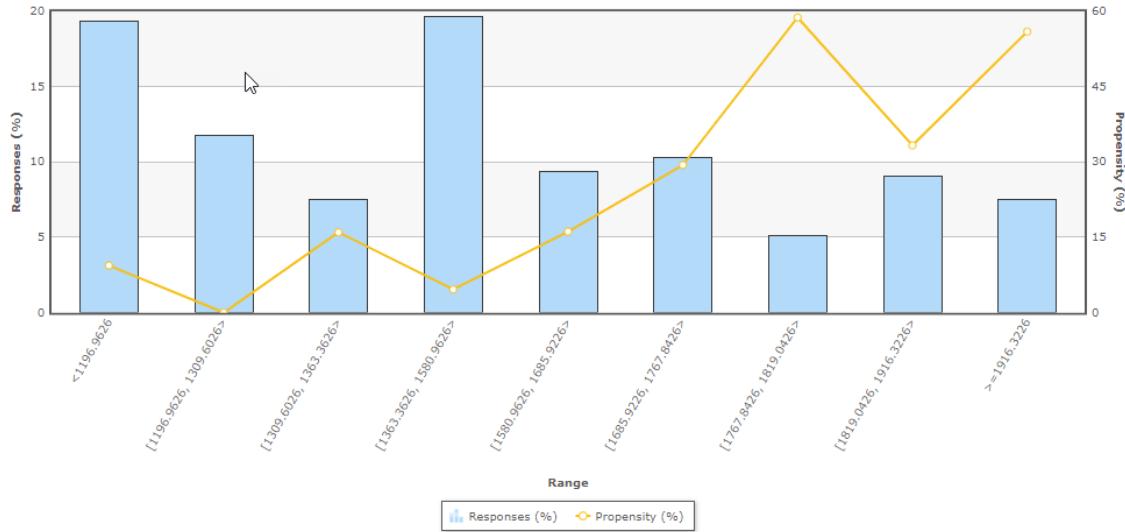
Predictors	Score distribution	Trend			
Correlated predictors are grouped and the best performing predictors become active in the model.					
Predictors	Status	Type	Performance (AUC)	Range/Symbols(#)	Bins(#)
Customer.AverageSpent	Active	Numeric	80.86	[1001.33; 1997.33]	9
Customer.Age	Active	Numeric	75.54	[19.0; 80.0]	9
Customer.InteractionContext.PreviousWebpage	Active	Symbolic	74.22	4.00	4
Customer.AverageBalance	Active	Numeric	73.51	[506.21; 1996.78]	9

There are three tabs, reporting on predictors, the model score distribution and the trend. In the predictors report, you can examine the performance of individual predictors.

Let's examine a couple of them.

In this case, the best performing predictor is AverageSpent. This a predictor of type numeric.

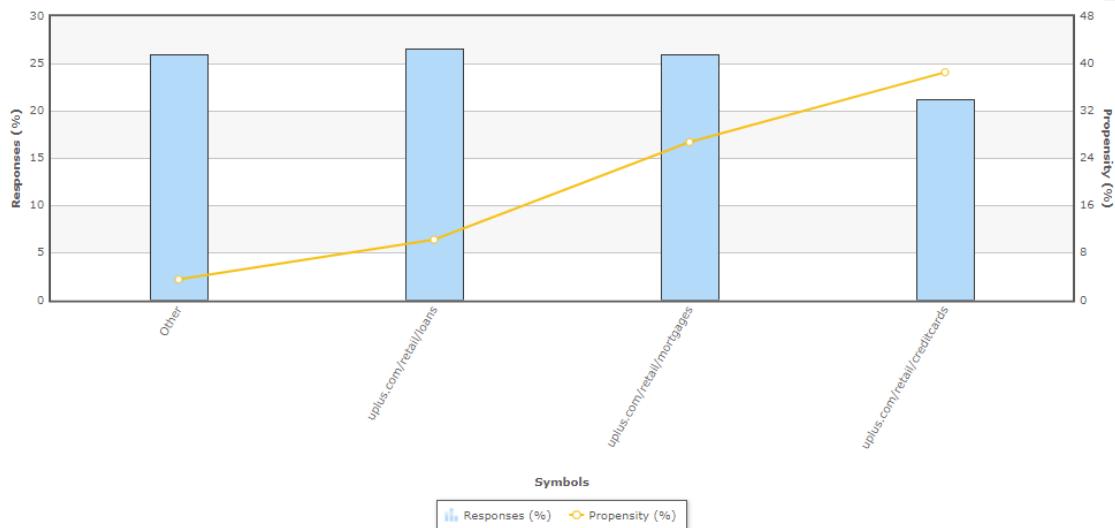
The system split the AverageSpent predictor into 9 bins. Each bin has its own offer propensity.



Propensity is the likelihood of positive customer behavior, which in this example is clicking on a web banner.

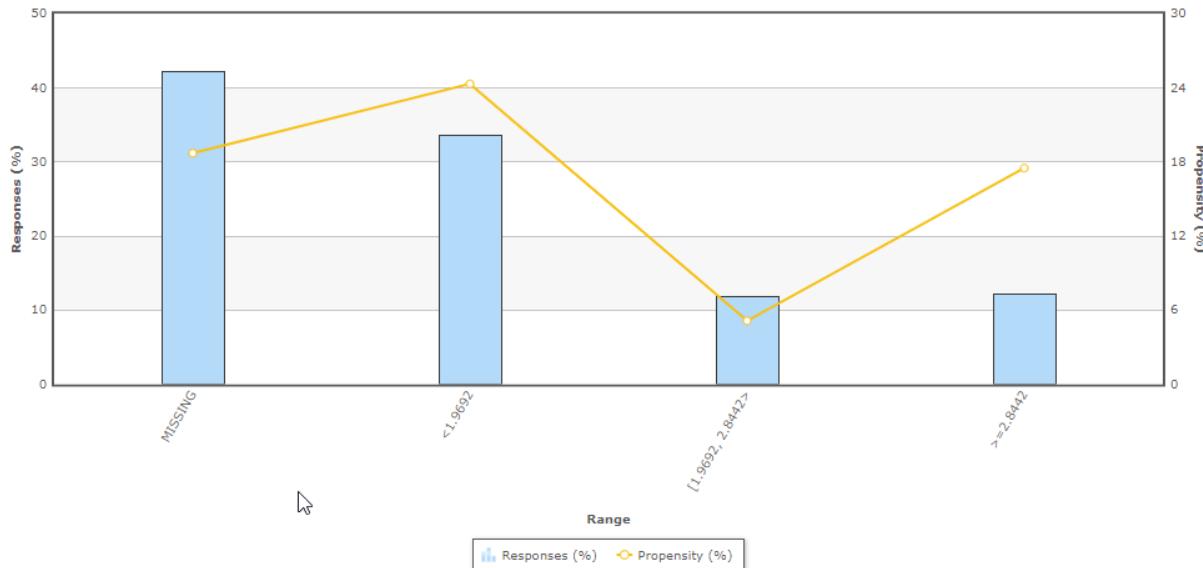
Now, let's examine the symbolic PreviousWebpage predictor.

The system split this predictor into 4 bins. The context of an interaction, in this case the previous web page visited by the customer, can be highly predictive.



To further improve the predictive power of the models the system uses Interaction History summaries.

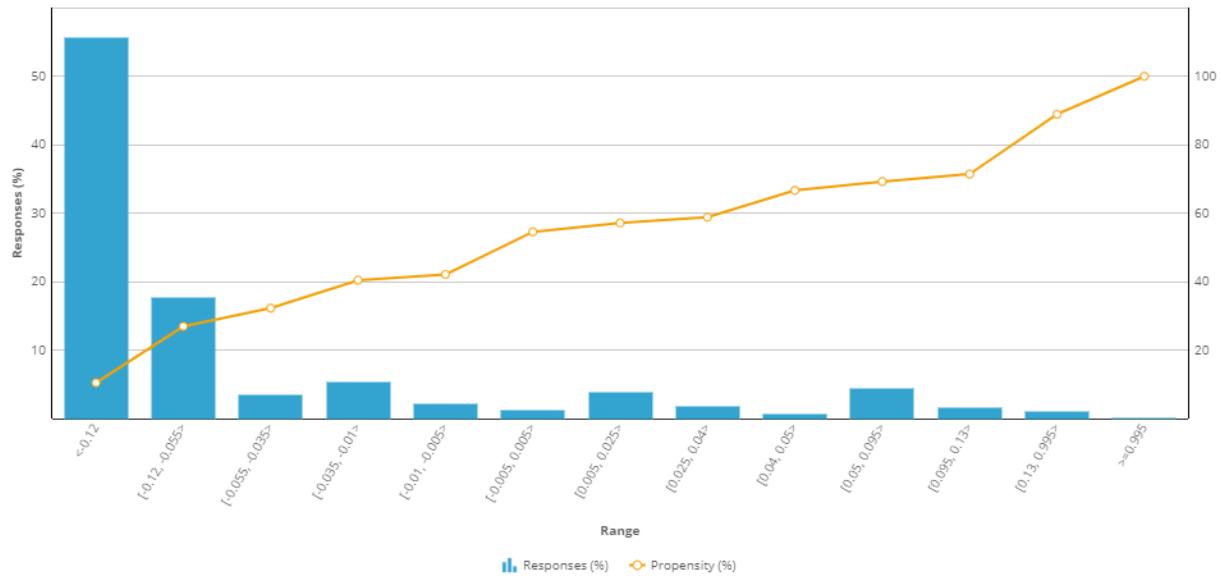
In this example, the adaptive system established that the number of days since the offer was accepted is a well-performing predictor.



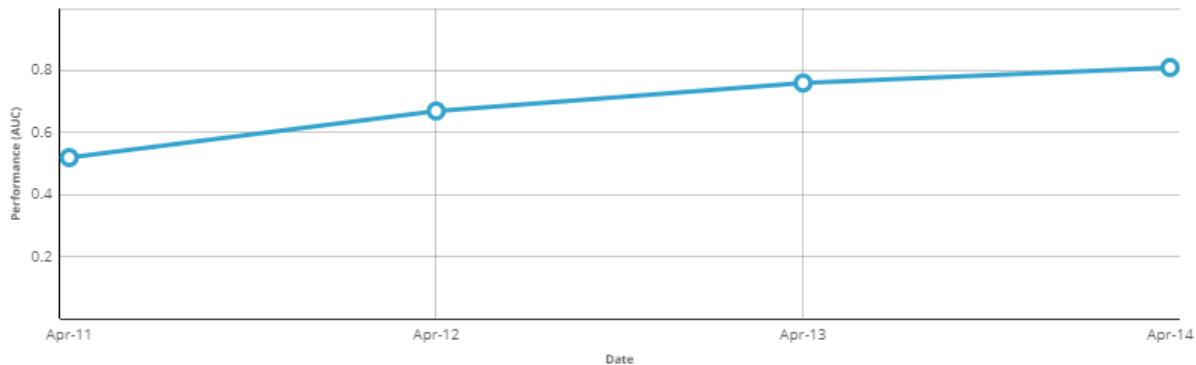
The system grouped three predictors that are correlated. It then marked two of them as inactive. Inactive predictors are not used in the propensity calculation.

▼ IH.Web.Inbound.Impression.pxLastOutcomeTime.DaysSince	Active	Numeric	56.52	[1.96; 3.9]	5
IH.Web.Inbound.Impression.pxLastGroupID	Inactive	Symbolic	54.48	2.00	2
IH.Web.Inbound.Impression.pyHistoricalOutcomeCount	Inactive	Numeric	54.48	[1.0; 1.0]	2

The Score distribution report enables you to examine the propensity distribution for the whole model.



And in the trend report you can see the performance of the model over time.



This demo has concluded. What did it show you?

- How the key metrics of adaptive models are visualized in a bubble chart.
- How you can customize the bubble chart by filtering.
- How to inspect active and inactive predictors.
- How to inspect individual predictors.
- How predictors with similar predictive performance are grouped.
- How to examine the propensity distribution for the whole model.
- How to examine the trend for the whole model.

Exporting historical data

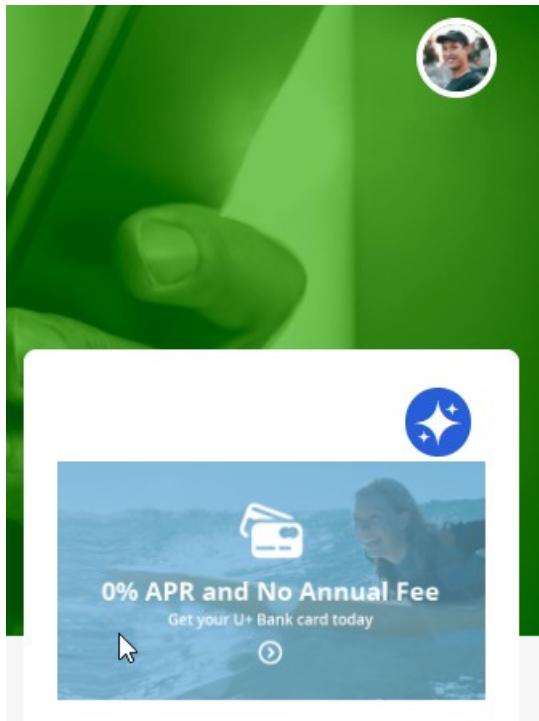
Introduction

Learn how to extract historical data (predictors and outcomes) from adaptive models in your application to perform offline analysis or use the data to build models using the machine learning service of your choice.

Transcript

This demo shows you how to export the customer interaction data that is used by adaptive models to make predictions, including all predictor data and associated outcomes, for offline analysis.

U+ Bank has implemented Pega Customer Decision Hub™ to show a personalized banner on their website that advertises credit card offers.



When customers are eligible for multiple credit cards, adaptive models decide which card to show them.

When the customer ignores the banner, the adaptive model that drives the decision regards this as negative behavior.

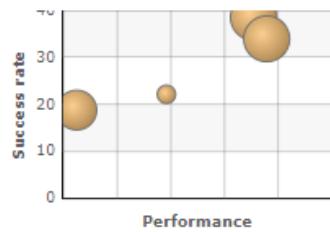
When the customer clicks on the banner, the model regards this as positive behavior.

As a data scientist, you may want to inspect the raw predictor data used by an adaptive model and the customer interaction outcome to validate data assumptions and check for concept drift.

You can also use the data to build various predictive models externally.

All models are managed in Prediction Studio.

The adaptive model that drives the decision about which banner to display is the **Web Click Through Rate** model.



[Web_Click_Through_Rate](#) :

To extract the data, enable the recording of historical data for a selected adaptive model.

A web banner typically has a low click-through rate and a significantly lower number of positive responses than negative responses.

In such cases, you can sample all positive outcomes and just one percent of the negative outcomes to limit the storage space needed.

Recording historical data

Save historical data in a repository to use for offline analysis.

You can find an overview of the historical data in

[Historical data overview](#).

Record historical data

	Sample percentage
Clicked	100.00%
NoResponse, Impression	1.00%

The sample percentages determine the likelihood that a customer response will be recorded.

The system stores the predictor data and outcome as a JSON file in the repository of your choice.

By default, the data is stored for 30 days in the **defaultstore** repository.

However, this repository points to a temporary directory, and a system architect should switch to a resilient repository to avoid data loss.

Supported repository types include Microsoft Azure and Amazon S3.

Select repository type

X



Artifactory



S3



File system



Azure

For this demo, we use the default store repository and create a data set to export the data.

New data set

X

Name *

ADMPayload

Type *

File

Apply to *

PegaCRM-Data-Customer

Development branch

[No branch] ▾

Add to ruleset

PegaCRM-Artifacts ▾

Ruleset version

01-01-99 ▾

Cancel

Create

The data set is mapped to the file that contains the recorded historical data.

File **Mapping**

Data source

Files on repository
 Embedded file

Connection

Repository configuration* **defaultstore**

File configuration

Use a file path
 Use a manifest file

File path* **ADM/Rule-Decision-AdaptiveModel/Data-Decision-Request**

With these settings in place, the input data used for the prediction and associated outcome are stored in the configured data set when customers see an offer and click on an offer.

The system architect can download the data set in DEV Studio.

Data set: ADMPayload [Available]
CL PegaCRM-Data-Customer **ID ADMPayload** **RS PegaCRM-Artifacts:01-01-99**

Save as Delete

File	Mapping	Specifications	History	Run
				Refresh
				Delegate
				Add to favorites
				Export
				Import

Data source

Files on repository
 Embedded file

Every record contains the predictor values used for the prediction, as well as the context and decision properties, including the outcome of the interaction.

All property names are automatically converted to comply with the JSON format.

The screenshot shows a software interface with a navigation bar at the top: Monitor, Predictors (highlighted in blue), Outcomes, and Settings. Below the navigation bar is a table titled 'Fields (71)'. The table has two columns: 'Name' and 'Value'. The 'Name' column lists various customer-related fields such as Customer.Age, Customer.AverageBalance, Customer.AverageSpent, Customer.BranchCode, Customer.ChurnScore, and Customer.City. The 'Value' column displays their corresponding JSON values. A red box highlights the 'Customer.Age' row, and another red box highlights the 'Customer Age' value within the JSON object.

Name	Value
Customer.Age	<pre>"Decision_SubjectID": "16", "Context_Treatment": "Premier Rewards card tile", "Decision_Rank": "1.0", "Context_Group": "CreditCards", "Customer_Date_of_Birth": "8817.0", "Customer_NetPromoterScore": "7.0", "Customer_pyFirstName": "Joanna", "Customer_pyID": "16", "negativeSampling": "100.0", "Customer_AverageBalance": "1100.23", "Customer_InteractionContext_VisitDuration": "60", "Customer_HasInsurance": "Y", "Decision_InteractionID": "-3320806451547993547", "Customer_Age": "25.0", "Customer_CLV_VALUE": "400.0", "Customer_CreditScore": "550.0", "Customer_MonthlyPremium": "250.0", "Decision_Outcome": "Clicked",</pre>

To use the JSON file for further analysis, import the file into a third-party analytics tool.

Keep in mind that when many customers visit the website, the file size becomes very large in a short time. To limit the storage space needed, you can lower the sample percentages.

You have reached the end of this demo. What did it show you?

- How to export the raw data that is used by adaptive models.
- What data is captured during a customer interaction.

Creating and understanding decision strategies

Description

Next-Best-Action Designer provides a guided and intuitive UI to bootstrap your application development with proven best practices that generate the underlying strategies for you. These strategies can be customized using designated extension points or by building decision strategies from scratch, depending on the business requirement.

Learning objectives

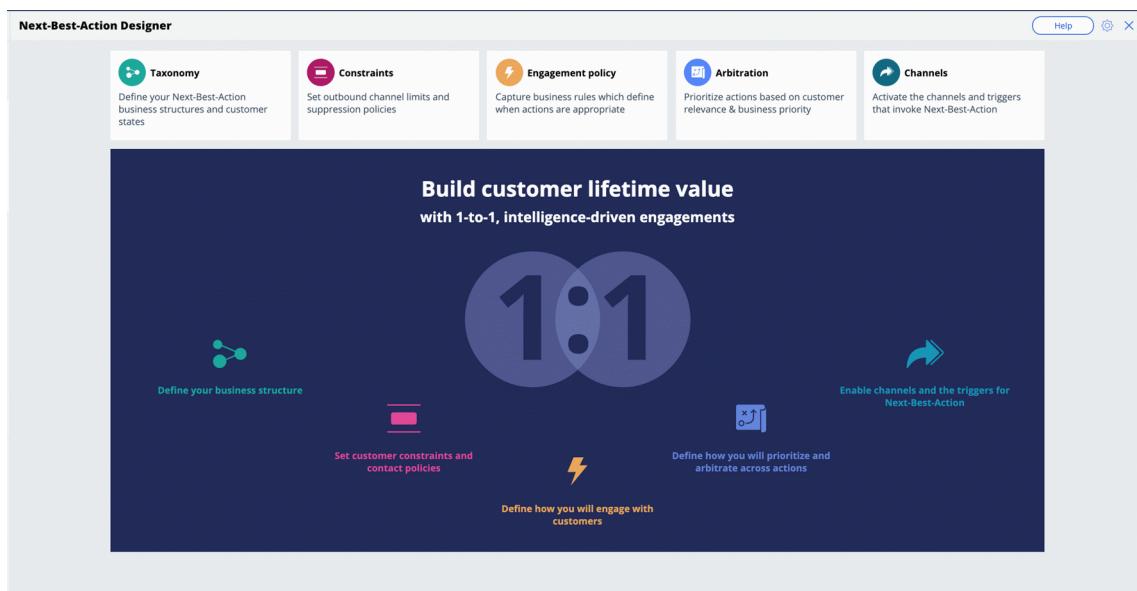
- Describe how decision strategies are used in the Next-Best-Action strategy framework
- Explain the decision strategy canvas and its building blocks
- Create decision strategies from scratch
- Explain what's going on inside each component when a decision strategy is executed

Decision strategies

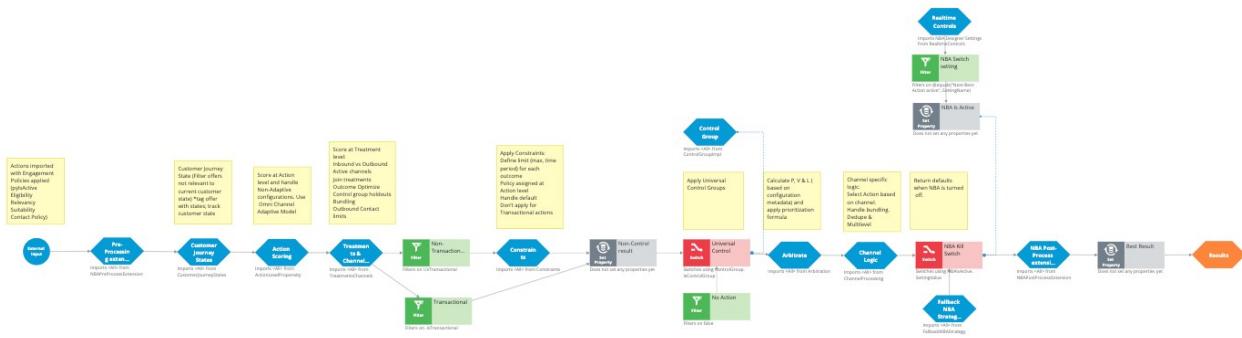
Transcript

Next-Best-Action Designer guides you through the creation of a Next-Best-Action strategy for your business. Its intuitive interface, proven best practices and sophisticated underlying decisioning technology enable you to automatically deliver personalized customer experiences across inbound, outbound and paid channels.

The Next-Best-Action Designer user interface allows you to easily define, manage and monitor Next-Best-Actions.



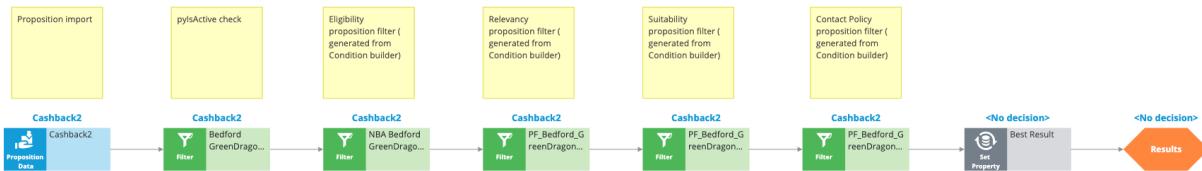
As you use the Next-Best-Action Designer user interface to define strategy criteria, the system uses these criteria to create the Next-Best-Action Strategy framework. This framework leverages best practices to generate Next-Best-Action decision strategies at the enterprise level. These decision strategies are a combination of the business rules and AI models that form the core of the Pega Centralized Decision Hub, which determines the personalized set of Next-Best-Actions for each customer.



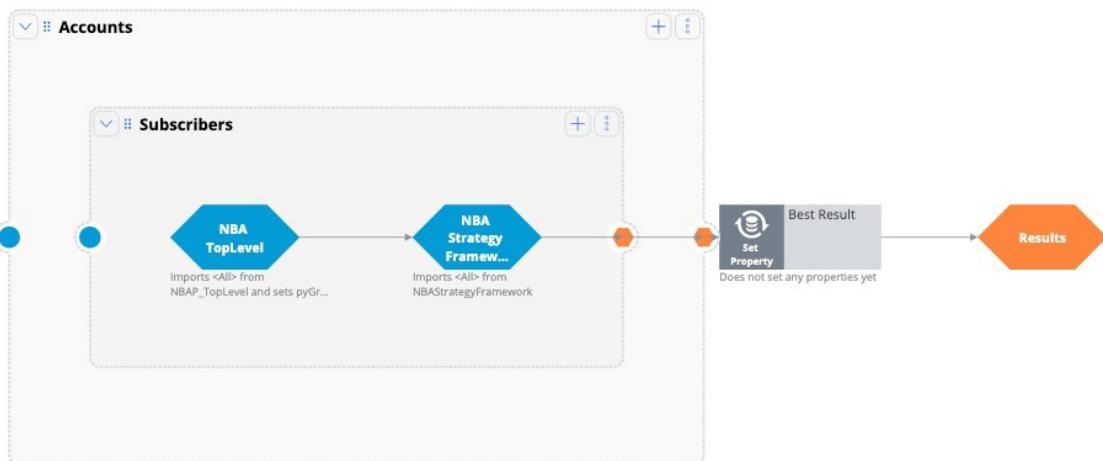
If you want to modify the strategy later, you can do that from Next-Best-Action Designer's simple and transparent interface.

The strategy framework is applied to all relevant Actions and Treatments after you define a Trigger in the Next-Best-Action Designer **Channels** tab.

Each Trigger generates a strategy that first imports the Actions from the appropriate level of the business structure and then applies the Eligibility, Relevance and Suitability rules.

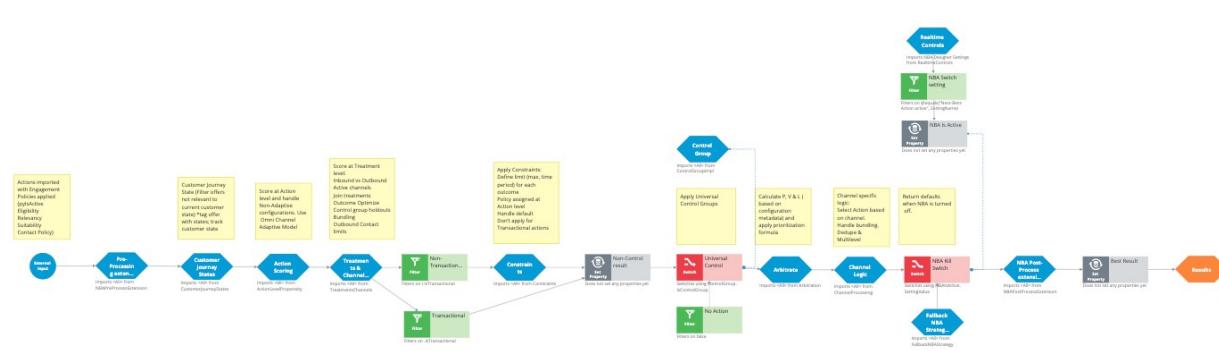


The strategy then passes these results to the strategy framework for processing.

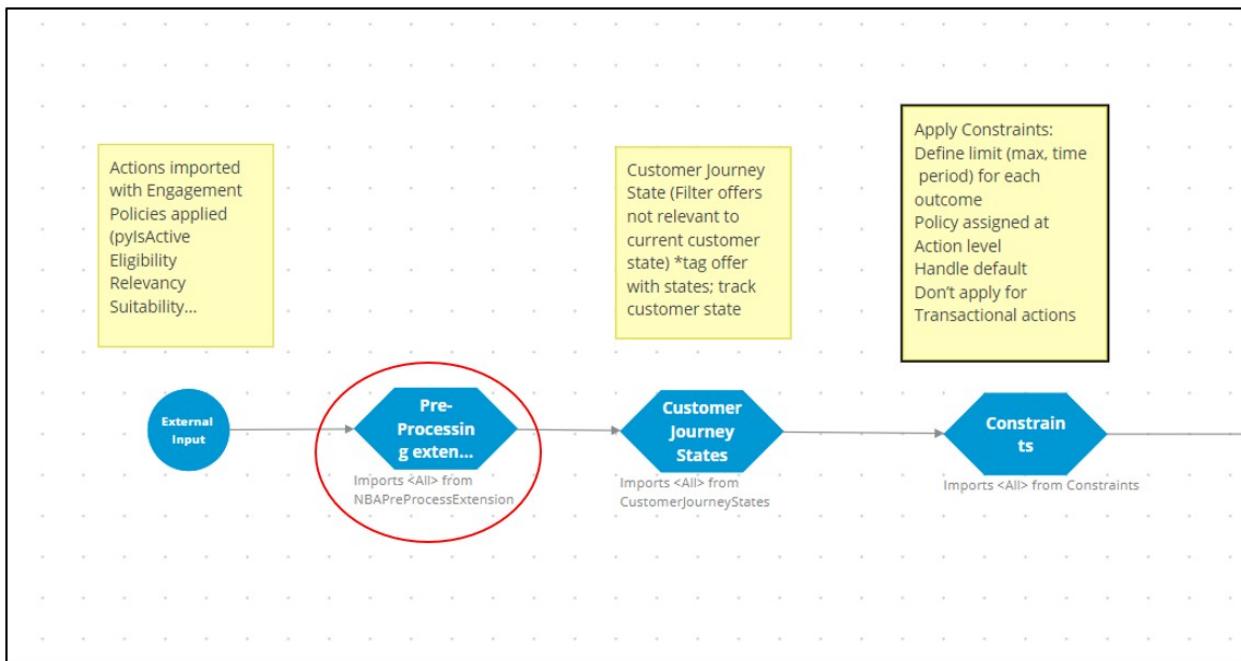


There are several extension points within the framework. An extension point is an empty rule or activity that is intended to be overridden to meet the specific needs of the application. When building an implementation of the current framework, the decision strategy designers must override the empty activity with a functioning interface to their customer master file.

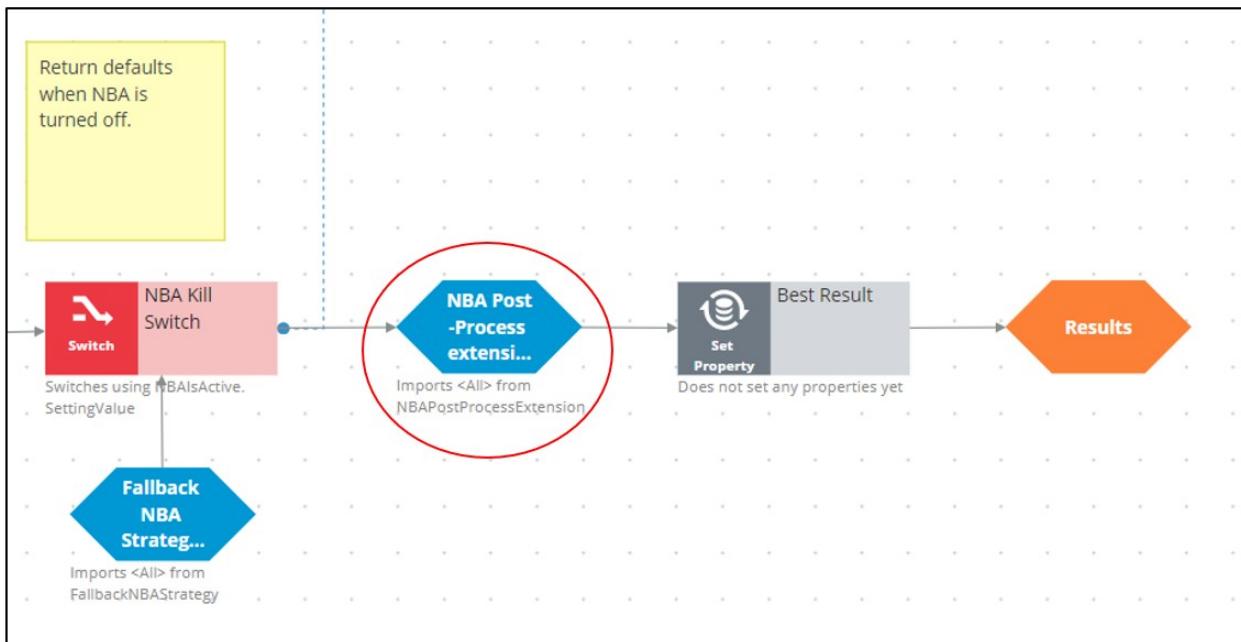
This is the NBA framework strategy when applied to each of the Actions.



The first component within the strategy framework is an extension point for any Action pre-processing you might need to perform.



The last functional component within the strategy framework is another extension point for any post-processing that must be performed.



Similarly, there are many other extension points such as the outbound limits extension points and business value extension points.

To ensure upgradeability, avoid overriding any part of the framework that is not a designated extension point.

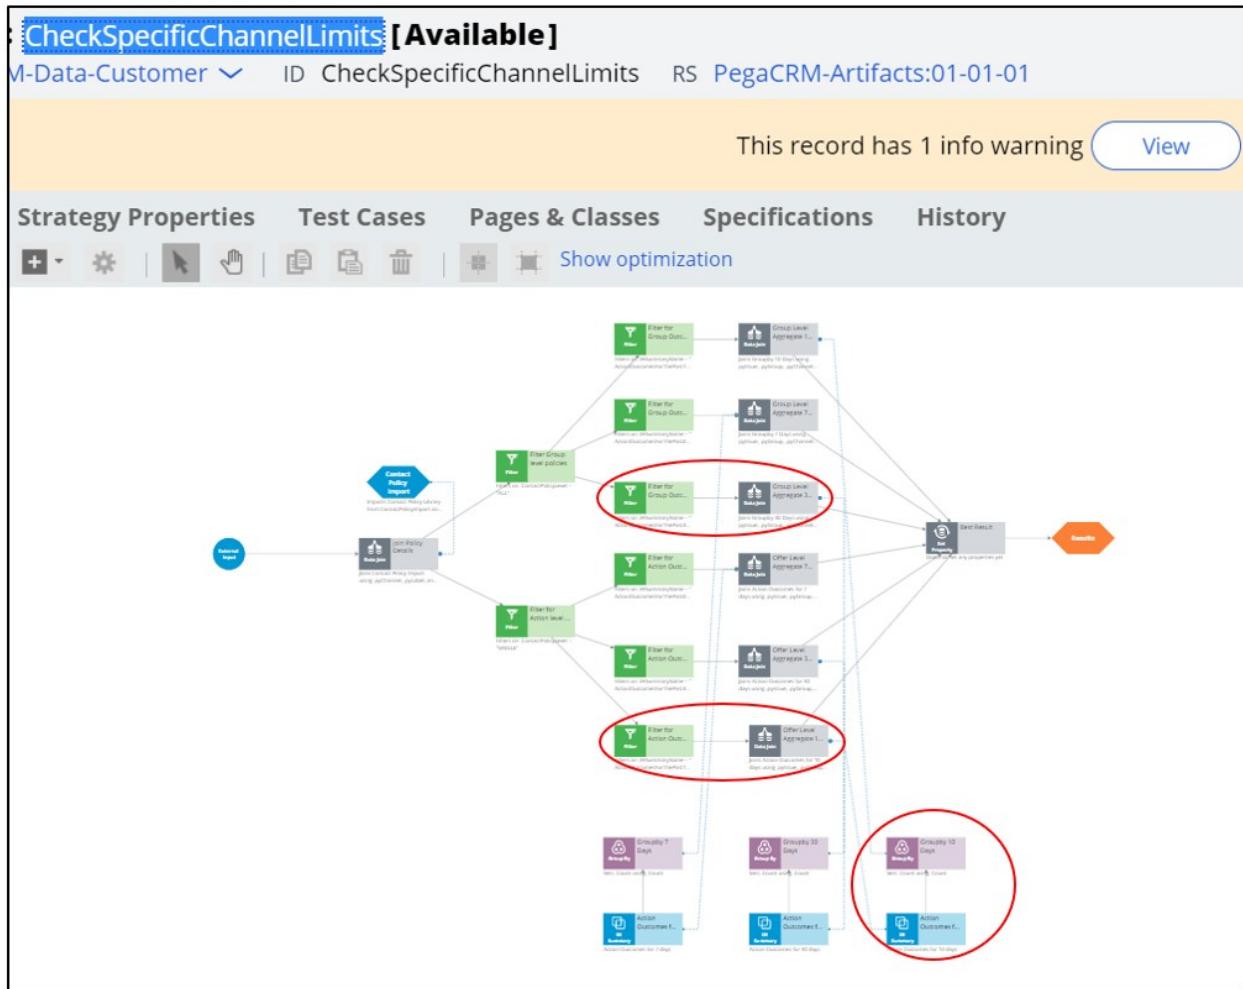
Also, the generated framework has some extension points where you can create strategies.

For example, while configuring values for Arbitration, you can specify a business value for an Action, or you can use a strategy to calculate the value. This can be done by adding a strategy to the existing framework.

Similarly, in defining the engagement rules, you can use a new strategy as a definition instead of an existing condition. Strategy designers can create such strategies from scratch using the decision strategy canvas.

Or, while defining the suppression rules, you can add a strategy to define new suppression rule limits instead of the existing 7 or 30 days.

For example, in the screenshot below, the `CheckSpecificChannelLimits` rule has been extended to have a 15-day limit:



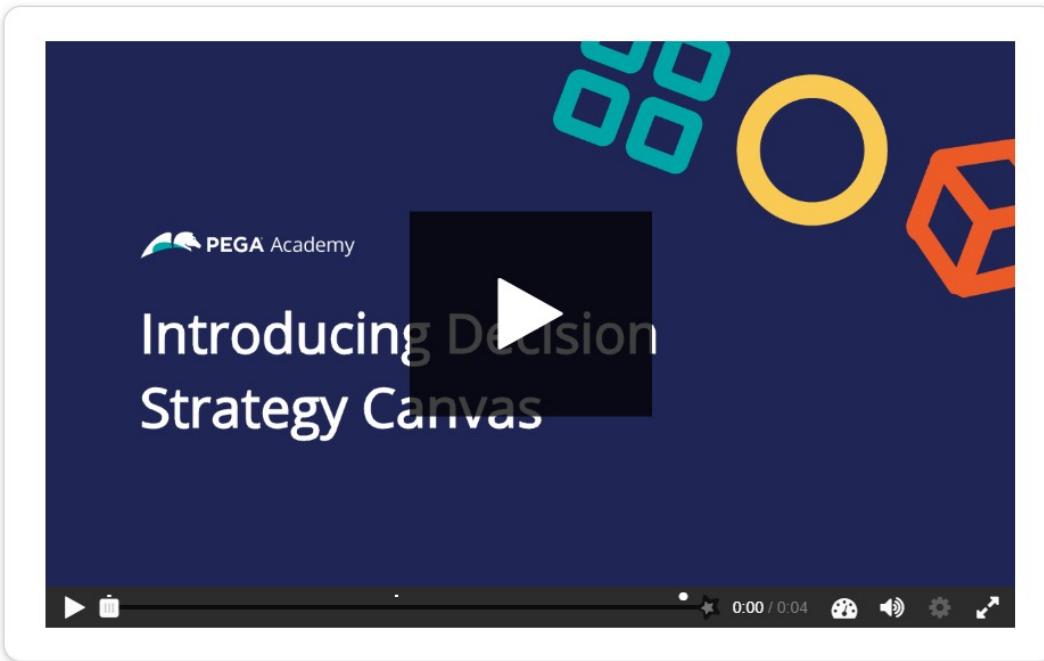
In conclusion, the NBA Designer provides a guided and intuitive UI to bootstrap your application development with proven best practices. NBA designer generates the underlying strategies for you, which can be extended using existing values in the designated extension points or by building decision strategies from scratch, depending on the business requirement.

Decision strategy canvas

Introduction

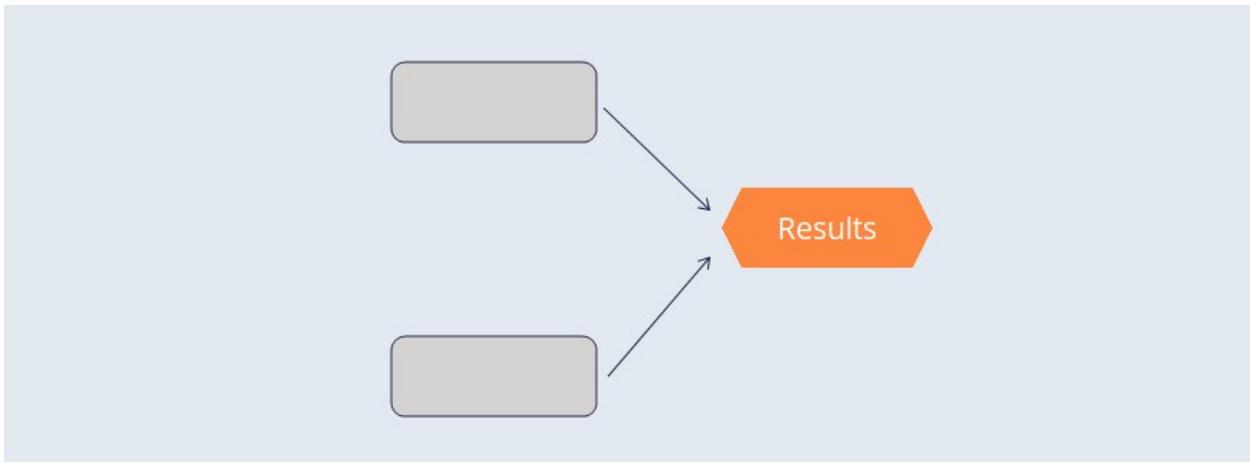
Decision strategies drive the next best action and comprise a unit of reasoning represented by decision components. You use the Proposition Data component to import actions into a strategy canvas. The sequence of the components in the canvas determines which action is selected for a customer.

Click the Play button to learn more about decision strategies.



Screen1: U+ business scenario

U+, a telecom organization, wants to promote two new phones in the contact center: iPhone and Galaxy. Click the + icons to learn more about the elements of a decision strategy that is created for this requirement.



Decision Components: A decision strategy is comprised of building blocks called decision components. You can add and connect components to implement the business requirements.

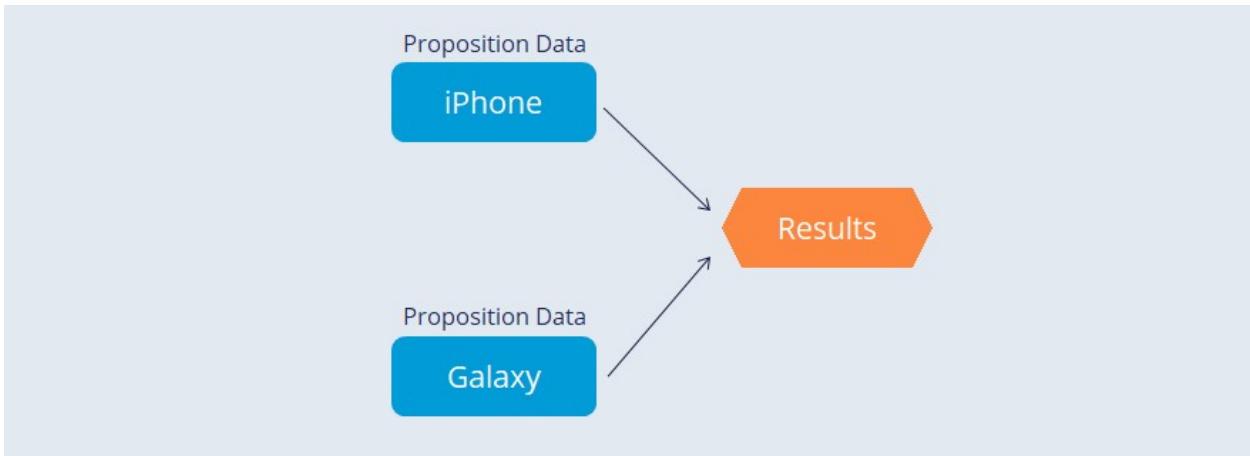
Arrows: An important element of the strategy canvas is the arrow. An arrow connects two decision components. A solid line means the data is copied from one component to another.

Strategy Canvas: In Pega, business users visually design decision strategies on what is known as a strategy canvas.

Screen2: Proposition Data component

The Proposition Data decision component imports the properties of an action. The result of this component is a flat list of all the properties.

Click the + icons on the proposition components to examine the components' results.



iPhone: This Proposition Data component outputs the Price and the Cost properties of the iPhone action.

Name: iPhone

Price: 150

Cost: 100

Galaxy: This Proposition Data component outputs the Price and the Cost properties of the Galaxy action.

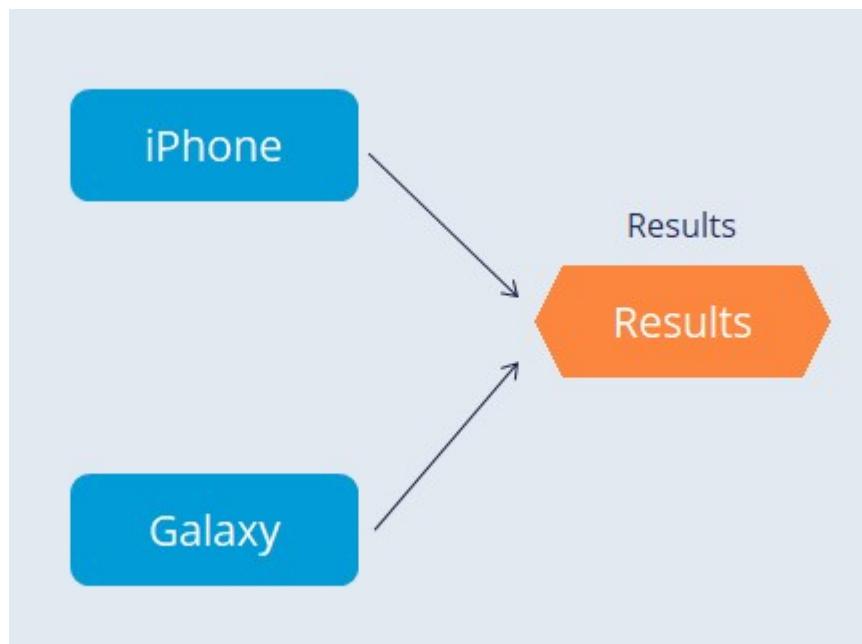
Name: Galaxy

Price: 250

Cost: 150

Screen3: Results component

Another decision component is the Results component. Each strategy always contains one Results component, which defines the output of the decision strategy.



How many actions do you think this strategy outputs?

- 0
- 1
- 2

Feedback: Because both iPhone and Galaxy are connected to the Results component with a solid arrow line, this decision strategy outputs two actions.

Dynamic pricing

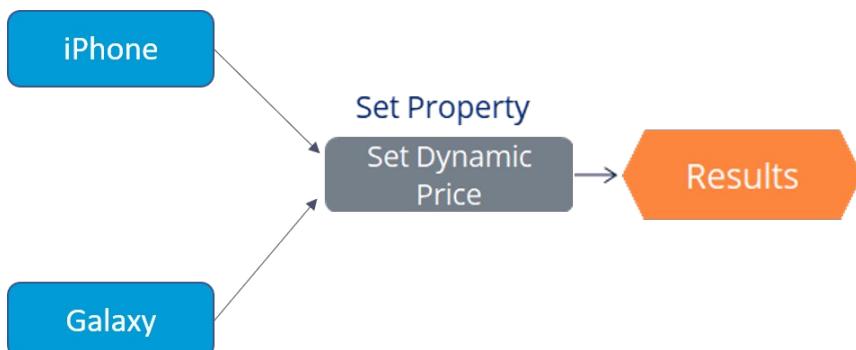
U+ Bank wants a dynamic Price for all offered actions. If the Customer value of a customer is higher than 60, the bank wants to offer a 10% discount to the customer.

To meet the new requirement, you must enhance the existing strategy to set the value of the Price based on Customer value. Changing the Price dynamically based on the Customer value makes the pricing customer-centric.

Set Property component

The Set Property component is used to dynamically alter the value of an action property based on a customer property. You use this component to set values to properties that are output by the strategy.

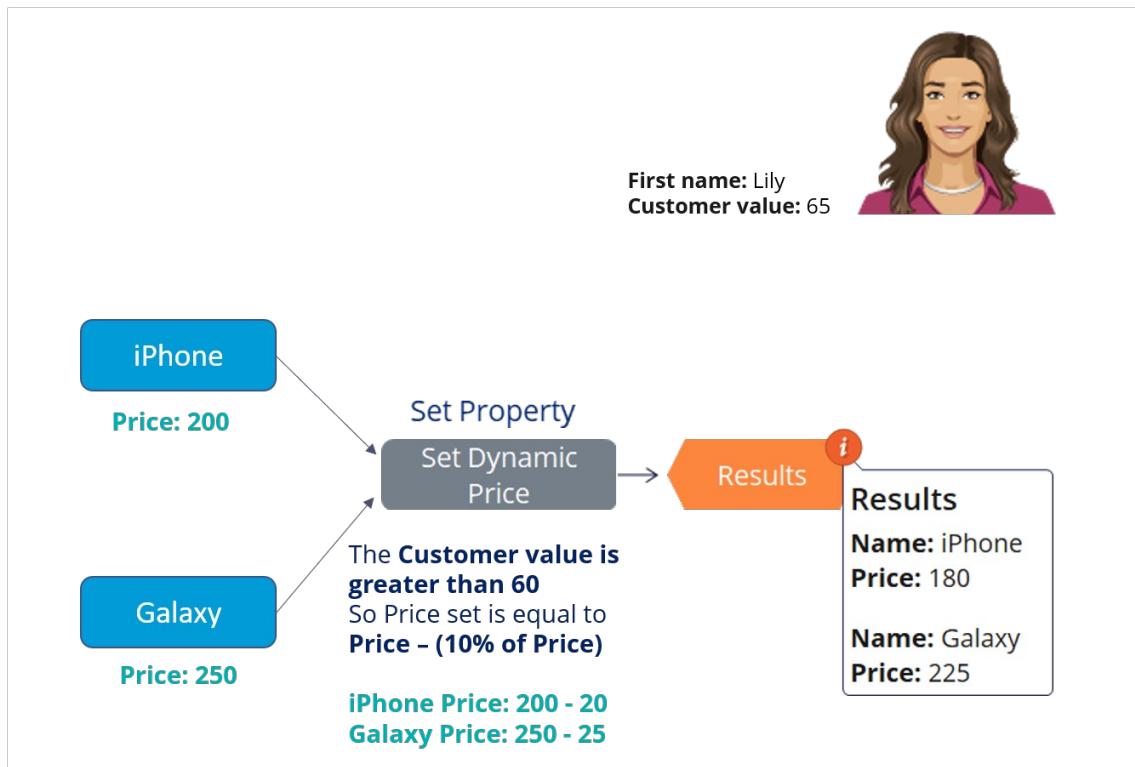
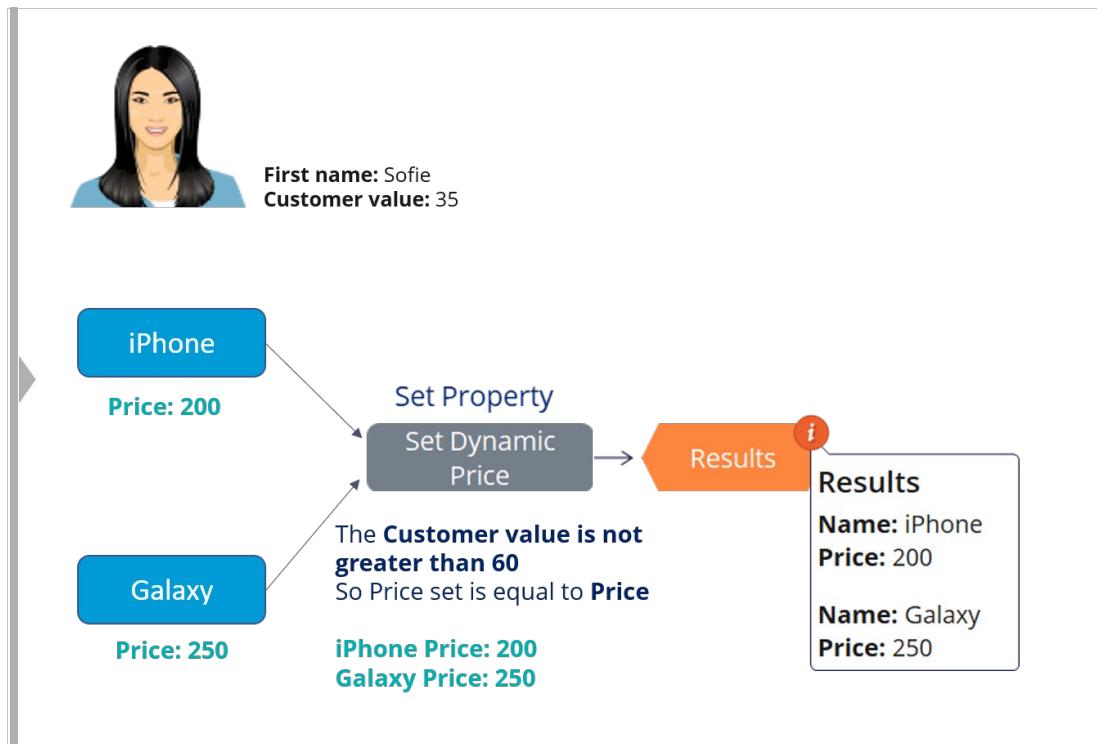
You can set properties to a constant or calculated value.



Example

Consider two customers: Sofie and Lily with customers value 35 and 65 respectively.

In the center of the following image, slide the vertical line to see how Sofie and Lily's Customer value affects the Price of the action offered to them.



Action ranking

U+ wants to offer the most profitable action to its customers.

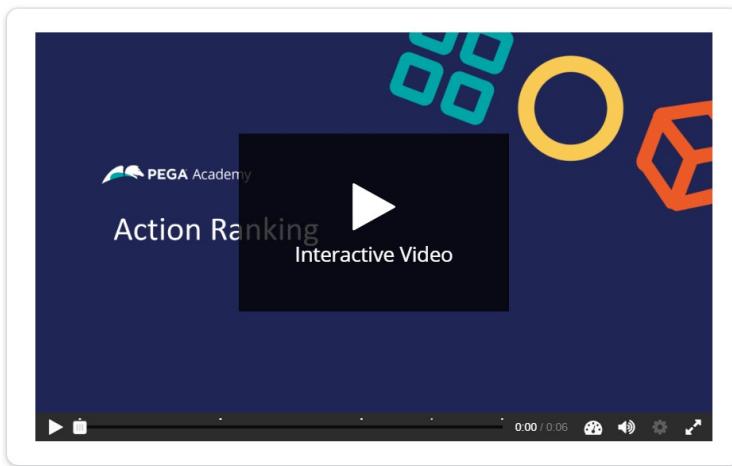
To enhance this strategy based on the new requirement, you need a new decision component that can rank the actions based on Profit and select the highest ranked action.

Profit is calculated based on Price and Cost action properties.

Prioritize component

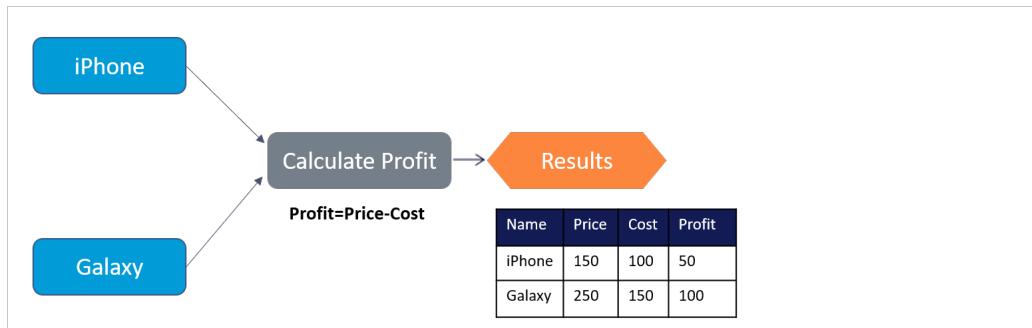
The Prioritize component is a decision strategy component used to rank actions. The Prioritize component is also used to select the top 1, top 2, or arbitrary top-n actions.

Click the Play icon to learn more about action ranking in detail with the help of a sample strategy.



Screen 1:

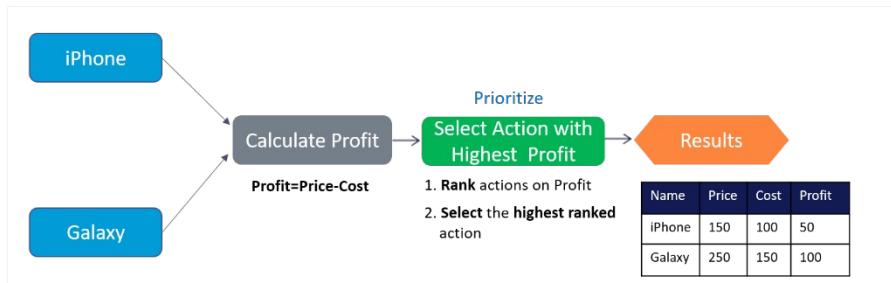
The initial strategy outputs price, cost and the profit calculated by the Set Property component. Click Add Prioritize to add the Prioritize component to the strategy.



Add Prioritize

Screen 2:

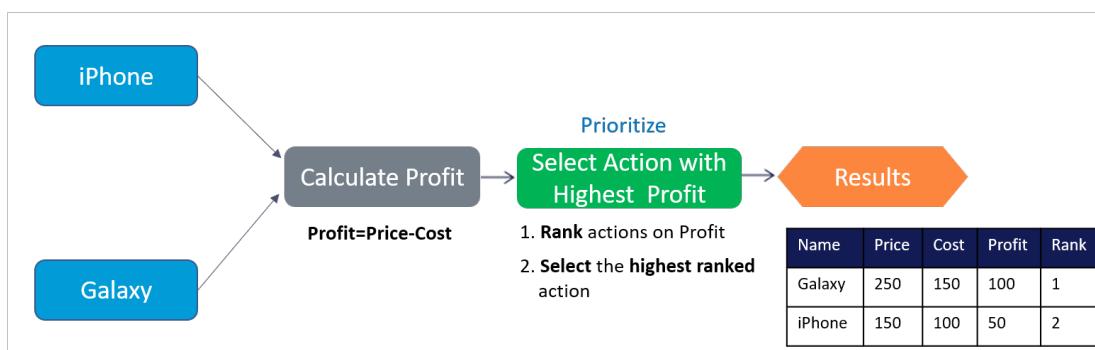
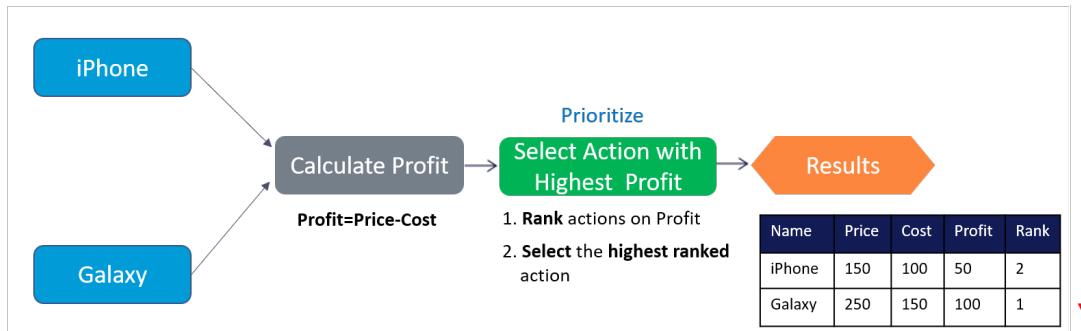
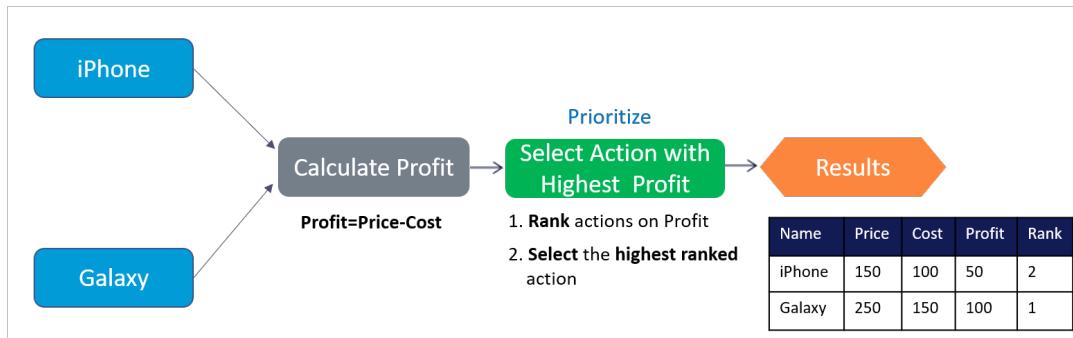
The Prioritize decision component can perform two operations: rank actions based on an expression, and select the highest ranked action. Click Rank actions to see how the component rank the actions.



Rank actions

Screen 3:

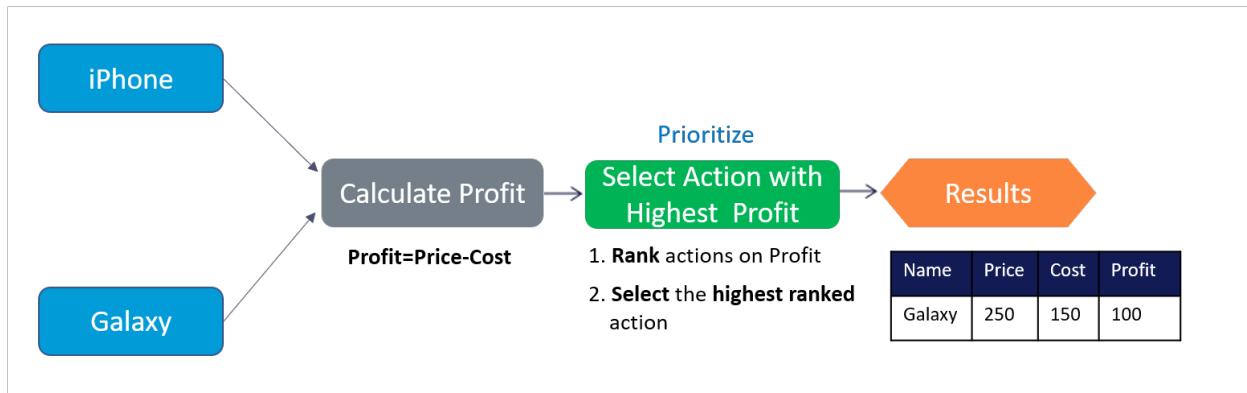
Once the actions are ranked, the prioritize component selects the highest ranked action. Click Select highest to see the action selected.



Select highest

Screen 4:

You can see that the output of the result component is the highest selected action: Galaxy with a profit of 100.



Quiz

Examine this strategy and then answer the following question to check your knowledge on action ranking.



What does the Results component of the strategy contain?

Sony with profit 150

LG with profit 100

Panasonic with profit 50

Feedback: The Prioritize decision component ranks the actions and selects the highest ranked action. Hence, the Results component of the strategy contains Sony with a profit of 150.

Creating a decision strategy

Introduction

Decision Strategies drive Next-Best-Action. They comprise a unit of reasoning represented by decision components. How these components combine determines which action will be selected for a customer: the Next-Best-Action. Learn the type of decision components and how they are used to create decision strategies. Gain hands-on experience designing and executing your own Next-Best-Action decision strategy.

Transcript

This demo will show you how to create a new decision strategy.

It will also describe three important decision components and the types of properties available for use in expressions during strategy building.

In this demo you will build a Next-Best-Label strategy. The Next-Best-Label strategy is a sample strategy, used to illustrate the mechanics of a decision strategy.

Start by creating a new strategy from scratch.

Decision strategies output actions, utilizing the so-called Strategy-Results class.

The Strategy-Results class limits the output of the strategy to the actions contained in the Business issue and Group.

The strategy you build will select a Label action from a set of predefined actions. The Label action selected will be the one with the lowest printing cost.

Notice that the complete definition of the Next-Best-Label strategy needs to include a reference to the PegaCRM-Data-Customer class.

This is the 'Apply to' class and it indicates the context of the strategy.

It ensures that from within the strategy, you have access to customer-related properties such as Age, Income, Address, Name, etc.

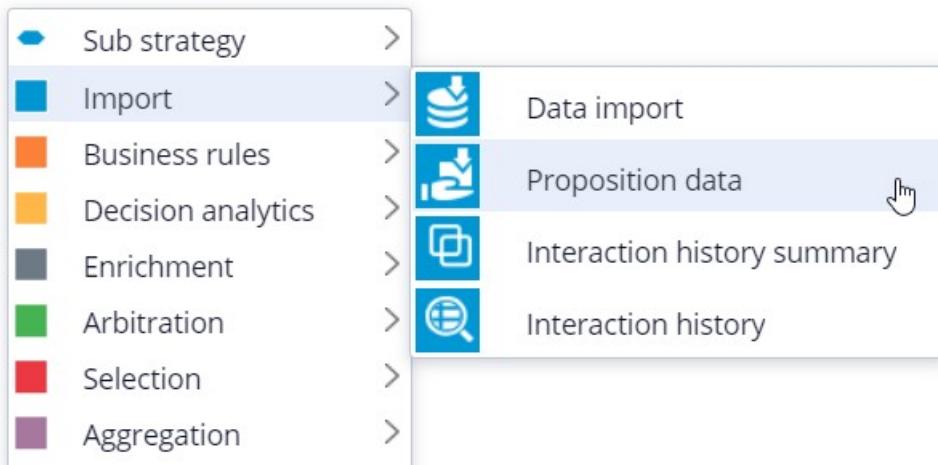


You can now start building the strategy. Right-click on the canvas to get the Context menu, which shows all component categories.

The first component to add is an Import component.

By expanding the Import category, you can see the Import component types available.

In this case you need a Proposition Data component to define the actions that will be considered by the strategy.



Now you need to configure the component. First, right-click to open the Proposition Data properties panel.

Notice that the Business issue and Group are grayed out.

This cannot be changed because the Enablement Business issue and Labels Group have already been selected for this decision strategy.

By default, the strategy will import all actions within that Group, unless you select a specific action.

For this component, you only want to import the Green Label, so let's select that.

Selecting the action from the drop-down menu automatically gives the component the appropriate name.

The description, which will appear under the component on the canvas, will also be generated automatically.

If you want to create your own description, you can do so by clicking the 'Use custom' radio button.

Now you want to import a second action into the strategy. You can use the Copy and Paste buttons to quickly add more Proposition Data components to the canvas.

You can use Alignment Snapping and Grid Snapping for easy placement of the components.

By turning these off, you can place a component anywhere on the canvas, but it makes it more difficult to align the shapes.



Now you need to add the next component in the strategy, which is an Enrichment component called Set Property.

You can add this component to the canvas by selecting it from the component menu.

Next, connect it to the Proposition Data components.

Ultimately, the result of this strategy should be the Label action with the lowest printing cost.

This printing cost is the sum of a base printing cost, which is specific to each label, and a variable cost, which depends on the number of letters.

The Set Property component is where you will calculate the printing cost for each of the actions.

The information in the 'Source components' tab is populated automatically by the Proposition Data components connected to this component.

Notice that the Black Label action is in the first row.

On the Target tab you can add properties for which values need to be calculated.

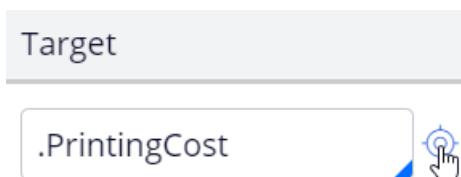
Click 'Add Item' to create the equation that will calculate the printing cost for each of the components.

Begin by setting the Target property to 'dot' PrintingCost.

In Pega, all inputs begin with a dot. This is called the dot-operator and it means that you are going to use a strategy property.

The PrintingCost property is a new strategy property that does not yet exist.

To create the new PrintingCost strategy property, click on the icon next to the Target field.



By default, the property type is Text. In Pega, there are various types supported. In this case, the PrintingCost is a numeric value, so change its type to Decimal.

Next, you need to make PrintingCost equal to the calculation you create. To create the calculation, click on the icon next to the Source field.

Using the Expression builder, you can create all sorts of complex calculations, but in this use case, the computation is very basic.

PrintingCost should equal BaseCost + 5 * LetterCount.

To access the BaseCost you type a dot. Notice that when you type the dot, a list of available and relevant strategy properties appears.

This not only makes it easy to quickly find the property names you're looking for; it also avoids spelling mistakes.

In a decision strategy, you have two categories of properties available to use in Expressions.

The first category contains the strategy properties, which can be one of two types.

An Action property is defined in the Action form. Examples are the BaseCost and LetterCount properties you are using here.

These properties have a value defined in the Action form and are available in the decision strategy via the Proposition Data component.

The property values can be overridden in the decision strategy but will often be used as read only.

The second type of strategy property is a calculation like the one you just created, PrintingCost. Such calculations are often created and set in the decision strategy.

These types of properties are either used as transient properties, for temporary calculations, or for additional information you want the strategy to output.

The second category contains properties from the strategy context, also called customer properties.

Suppose you want to use a customer property in your Expression, such as Age or Income.

In that case, you would have to type the prefix 'Customer dot', instead of just dot.

This is the list of available properties from the strategy context, also known as Customer properties.

For now, you calculate the printing cost for each action that does not use customer properties.

Finalize the Expression.

Expression builder

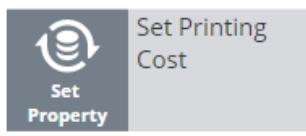
Browse

```
| 1 .BaseCost + 5 * .LetterCount
```

Even though you used the dot-operator to build your Expression, it's best practice to validate it, so click Test.

If the Expression isn't valid, you will receive an error message on screen.

On the canvas, you can see the automatically generated description for the component: Sets PrintingCost using BaseCost and LetterCount.



Set
Property

Sets .PrintingCost using .
BaseCost and .LetterCount

Now you want to ensure that the actions will be prioritized based on the lowest printing cost. So, you need to add the Prioritize component from the Arbitration category.

The prioritization can either be based on an existing property, or it can be based on an equation. Let's select an existing property using the dot construct.

Here you can select the order in which the top actions are presented. Since you are interested in the lowest printing costs, configure it accordingly.

You can also select the number of actions that will be returned by the strategy.

If you want to output only one label, select Top 1 here.

Expression*

Order by

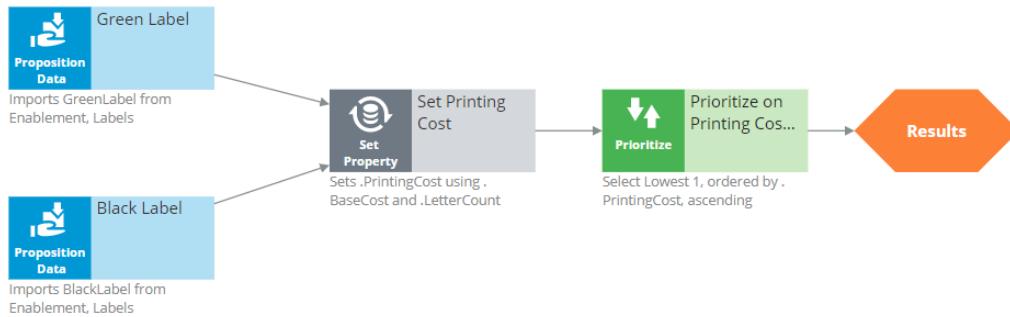
- Highest first (9 to 1)
- Lowest first (1 to 9)

Output

- Top
- All

1

Now you can connect the components and save the strategy.



To test the strategy, first check it out. Then, expand the right-hand side test panel and click 'Save & Run' to examine the results.

You can view results for any of the components by selecting that component.

If more than one action is present, each one is presented as a Page.

For the Set Property component, the Results contain a page for the Black Label and one for the Green Label.

For the Black Label the PrintingCost is 70.

For the Green Label the PrintingCost is 60.

On the canvas, you can show values for strategy properties such as Printing Cost.

For this exercise, you execute this strategy against a Data Transform called UseCase1.

If you open UseCase1, you can see the customer data the strategy uses when you run it.

To test the strategy on a different use case, you can create a Data Transform with different properties.

You can also select a Data Set that points to an actual live database table.

This demo has concluded. What did it show you?

- How to create a decision strategy from scratch.
- How to configure Proposition Data, Set Property and Prioritize decision components.
- How to build expressions in strategies.
- The two categories of properties available for expressions.
- How to test a decision strategy using a use case stored in a data transform.

Decision strategy execution

Introduction

Using Pega Decision Management, you do not need to be an expert in programming, math or data science to design and execute sophisticated decision strategies that engage your customers throughout the customer journey. With its highly intuitive graphical canvas, Pega Decision Management enables you to easily embed Pega or third-party predictive models into your decision strategies. The result is customer-centric interactions that improve the customer experience while increasing customer value, retention and response rates.

Transcript

This demo explains what's going on inside each component when a Decision Strategy is executed.

For example, what happens 'under the covers' when a Filter component is executed, and how does it interact with the components around it?

In the interest of keeping it simple, this example is limited to four actions. In reality, decision strategies will involve many more actions than that.



Here are our 4 actions: 'Green Label', 'Black Label', 'Red Label' and 'Blue Label'; they are represented by a Data Import or, more specifically, a Proposition Data component.

In this example, the Proposition Data components import three data properties for each action: Name, BaseCost and LetterCount.



Rank	Name	BaseCost	LetterCount
1	Green Label	10	10

The first action's Name is Green Label, its BaseCost is 10, and its LetterCount is 10.

Likewise, the other actions have a Name, BaseCost and LetterCount.



Rank	Name	BaseCost	LetterCount
1	Green Label	10	10

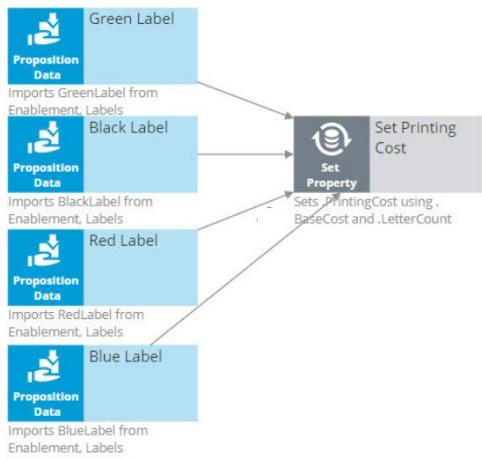
Rank	Name	BaseCost	LetterCount
1	Black Label	20	10

Rank	Name	BaseCost	LetterCount
1	Red Label	30	8

Rank	Name	BaseCost	LetterCount
1	Blue Label	40	9

One property is automatically populated for you; this is the Rank. We will come back to this later, but notice that, as separate components, each action has a Rank of 1.

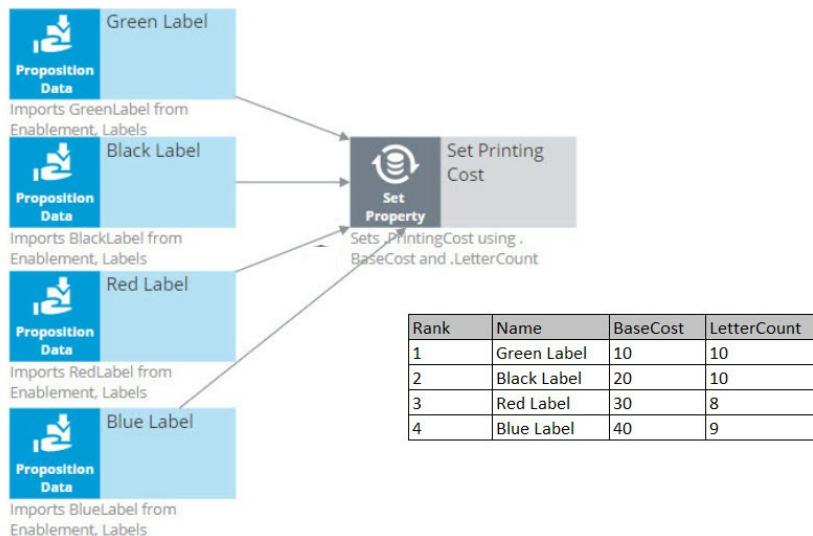
On the strategy canvas, components are connected by drawing arrows from component to component. So, what do these arrows mean exactly?



Well, when you draw an arrow, what happens is that, at runtime, all information in the component you're drawing the arrow from is available as a data source to the component you're drawing the arrow to.

So now, the Name, BaseCost and LetterCount for all of the actions are available in a single Set Property component.

The only data element that changes is the row number, or as we call it in the strategies, the Rank. In each decision component, the Rank value is automatically computed.



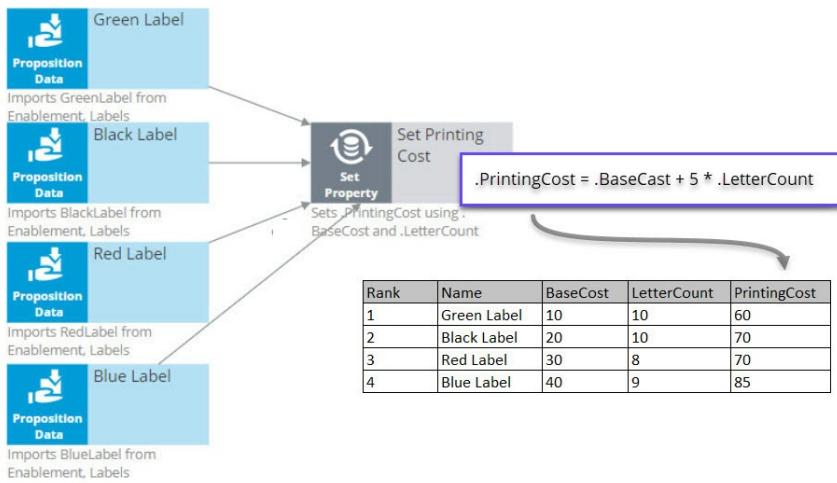
In the Set Property component, the Rank is determined by the order in which the actions are received by the component.

As a result, in this instance, the Green Label action has a Rank of 1, Black has a Rank of 2, Red has a Rank of 3, and Blue has a Rank of 4.

Ultimately, you want to select the best Label action. That is the Label with the lowest printing cost.

The printing cost of a Label is the sum of the BaseCost and a variable cost based on the LetterCount.

You configure the Set Property component to compute the printing cost of each Label action.



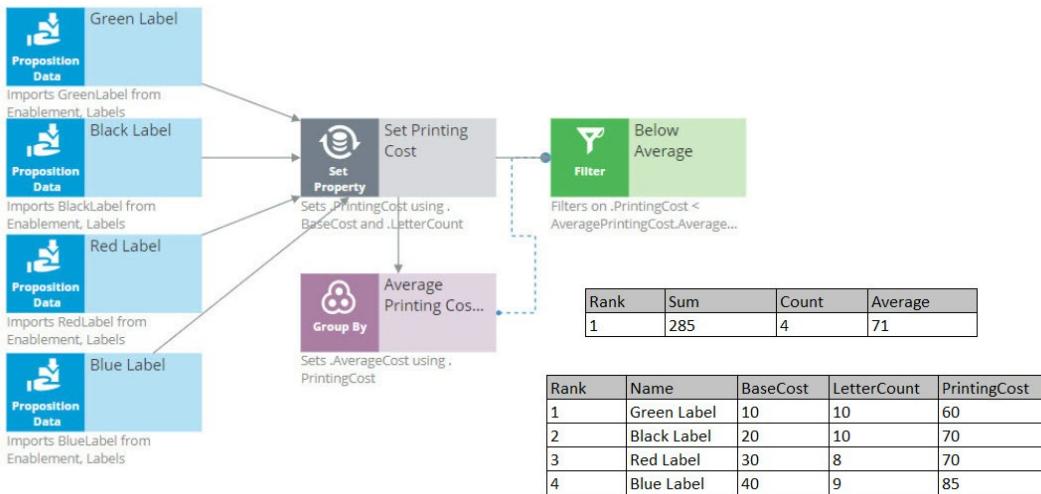
Because we are combining the data in our four Proposition Data components into one Set Property component, we only need to add one PrintingCost property to the new component, and it automatically computes the printing cost for all four actions.

For the Green Label action, PrintingCost equals a BaseCost of 10 plus 5 times the LetterCount of 10 which equals 60.

Similarly, the PrintingCost for the Black and Red Label actions is 70, and for the Blue Label action is 85.

Now, let's say the business rule is to select only Label actions with a printing cost lower than the average printing cost of all labels. For this requirement we use a 'Group by'/Filter component combination.

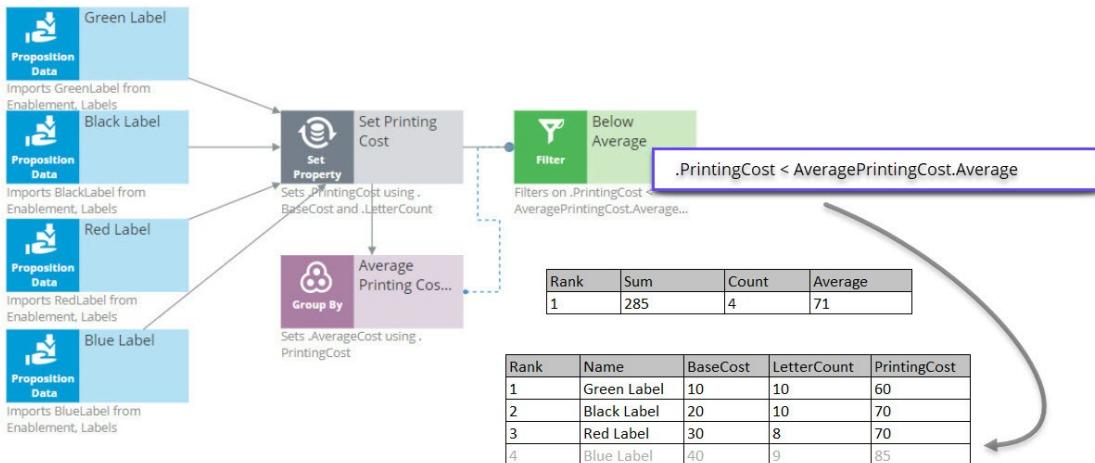
A 'Group by' component offers essential aggregation capabilities, like Sum and Count, that are used in many decision strategies. We will use it to calculate the average printing cost.



Again, we have our set of actions, each with their own specific PrintingCost value. The 'Group by' component combines all actions into one row. How does that work?

Well, it sums the PrintingCost values for all the actions, it counts the actions, and it calculates the average printing cost by dividing the summed printing cost by the count.

In this example, the sum of the PrintingCost values is 285, and the count of the actions is 4, so the average printing cost is 71.



Now that you have calculated the average printing price using a 'Group by' component, configure the Filter component to filter out actions that have a printing cost equal to or higher than this average.

So far in this strategy, we've seen only the solid line arrows, which copy information from one component to another. But now we also see a dotted line arrow.

This tells us that a component refers to information in another component.

Here, the Filter component is referencing the average printing cost that exists inside the Aggregation component. This is an important capability to understand.

The Filter component filters out actions when the printing cost for that action is equal to or above the average printing cost and propagates the other actions.

First, via the solid arrow, the filter looks at the actions sourced from the Set Property component.

Then, it applies the filter condition, which references the average printing cost in the 'Group by' component via the dotted arrow.

The Filter Condition in the Filter component is the Expression: 'dot PrintingCost is smaller than AveragePrintingCost dot Average'.

By using this ComponentName dot Property construct, any decision component can be referenced by any other component by name.

Important to note that the Filter component lets actions through when the condition Expression evaluates to **true** and filters out actions when the condition Expression is not met.

When you refer to a component, you always refer to the first element in the component, the one with Rank 1.

In this case, you are referring to the one and only row in the 'Group by' component, which naturally has Rank 1.

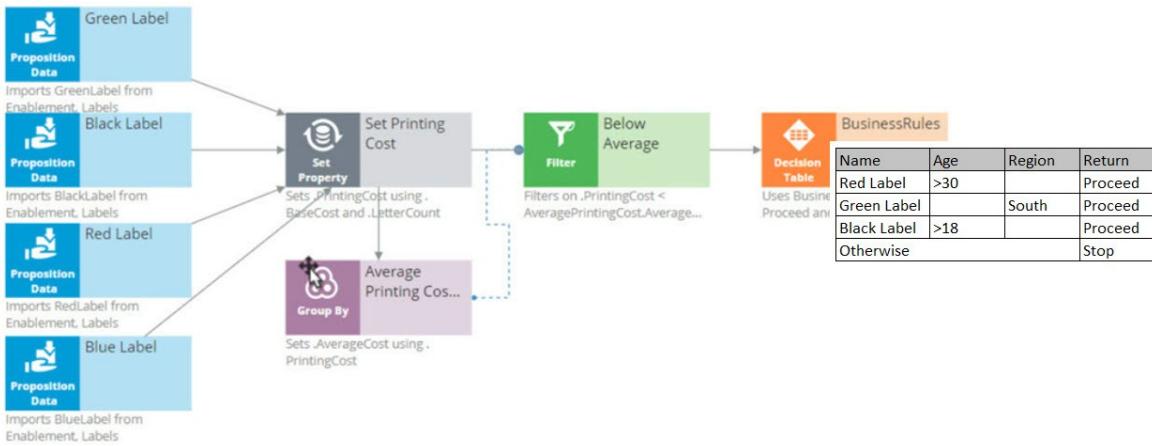
The Rank 1 average equals 71 in the 'Group by' component. This means that the filter will allow Label actions through that have a printing cost lower than 71.

By this standard, the printing cost of the Blue Label action is too high, so it is filtered out. The printing cost of the other Label actions are below 71, so they survive.

The result is that the table contains three surviving actions: Green Label with Rank 1, Black Label with Rank 2, and Red Label with Rank 3.

The next component is a Decision Table. A Decision Table in Pega is an artifact that can be used to implement business requirements in table format.

In a Decision Table, the business rules are represented by a set of conditions and a set of Return values.



The Decision Table receives information about the remaining actions via the solid arrow from the Filter component.

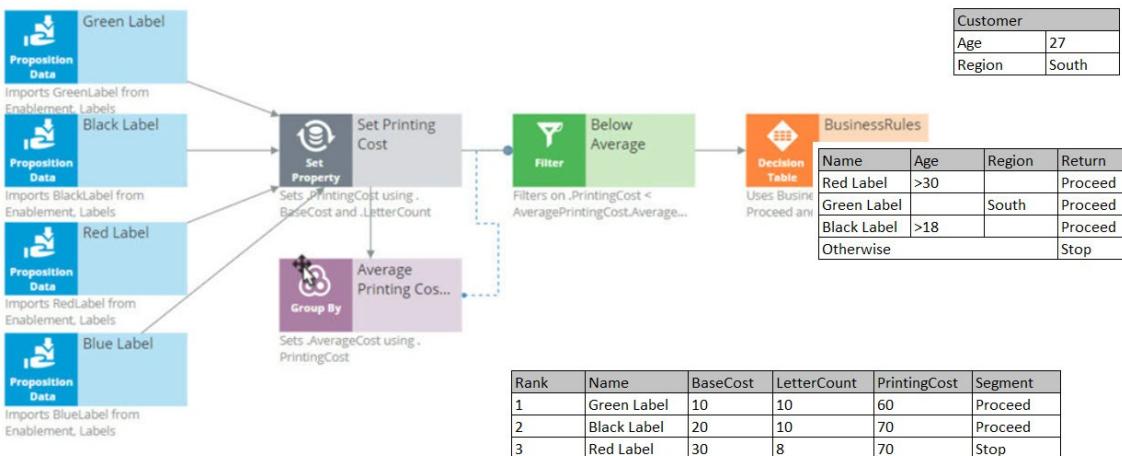
The business criteria say that the Red Label action can be offered if the customer's age is over 30 and they are from any region. If these criteria are met, the Return value is 'Proceed'.

The Decision Table also says that the Green Label action can be offered to anyone in the Southern region. So, if the Region value is South, the Return value for Green is 'Proceed'.

The Black Label action can be offered to anyone over the age of 18.

But in all other cases, or, Otherwise, no Label action meets the criteria, and the Return value is 'Stop'.

As an example, consider a customer with Age 27 and Region South.



Now, the Decision Table applies the business criteria for each action against the customer information and returns a value. The value returned by a Decision Table is also called a Segment.

The Decision Table checks the Green Label action with Rank 1 first, and in this case, it can proceed because the customer's Region is South.

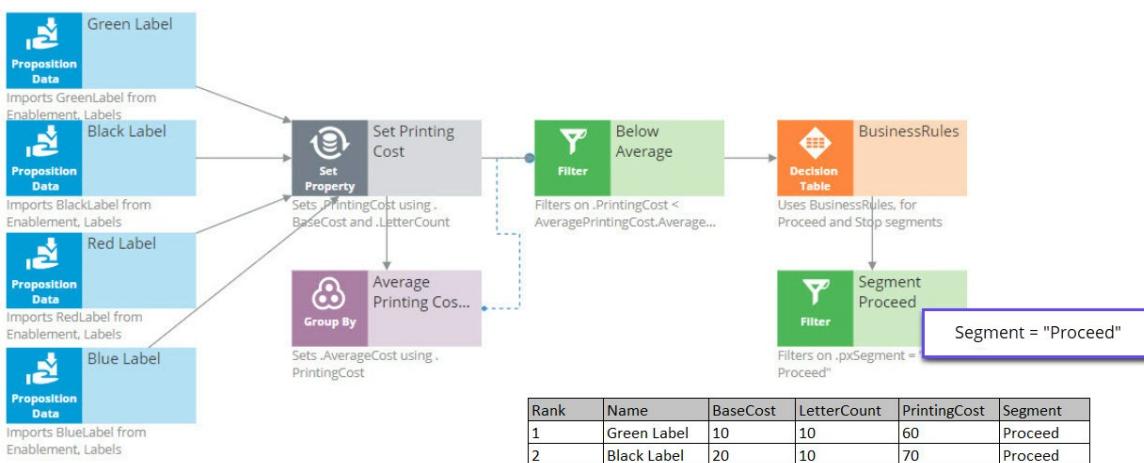
Next, it looks at the Black action and sees that the criteria for Black is that the customer's age is greater than 18. This customer is 27.

Black doesn't care about the Region, so the Segment value for the Black action is 'Proceed'.

Finally, it looks at the Red action, and the Age criteria don't match up, so the Segment value for Red is 'Stop'.

The result of the component is that you get a new segmentation column that flags which of the actions comply with the business rules.

You're now going to filter out the actions that do not match the business rules. This happens in the 'Segment Proceed' Filter component.



Again, via the solid arrow, the strategy copies the data over from the Decision Table component into the Filter component.

Now each action has a Rank, Name, BaseCost, LetterCount, PrintingCost and Segment. The filter condition is applied to this data.

The filter condition says: allow this action through if the Segment value equals 'Proceed'.

What this Filter component now does is go through the list of actions to find the actions with value 'Proceed' in their Segment property.

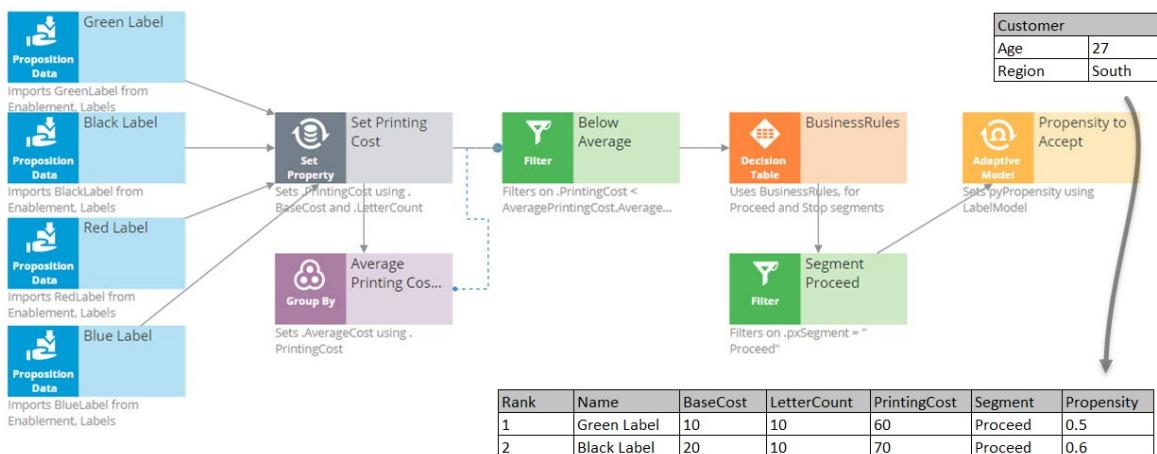
First is the Green Label. Green is allowed through, which means its properties will be available in the new component.

Then the Black Label. It is also allowed through because it also has 'Proceed' in its Segment property.

But the Red Label action is not allowed through, because Red has 'Stop' in its Segment property. Therefore, Red is not part of the output.

The strategy so far has selected two of our original actions, Green and Black.

Now, in the Adaptive Model component, you will use predictive analytics to determine the propensity of each of the remaining actions.



Propensity is the probability that a customer will accept an action, or, their likelihood of interest in it.

In order to calculate the propensity, we use an Adaptive Model component. The referenced model is configured to monitor customer characteristics such as Age and Region.

In this case our test customer has an Age of 27 and is from the South Region.

Again, just to keep it simple, we are using a model that makes predictions based on only this information. In reality, models will take into account many more properties.

The Adaptive Model determines the propensity.

First, we supply the action and the customer profile to the Adaptive Model, and the model says: 'Oh, it's the Green Label action; we have some evidence that young people like the Green Label action, but people from the South don't like it.'

Combining both factors, we get an overall propensity of 0.5 for the Green Label action.

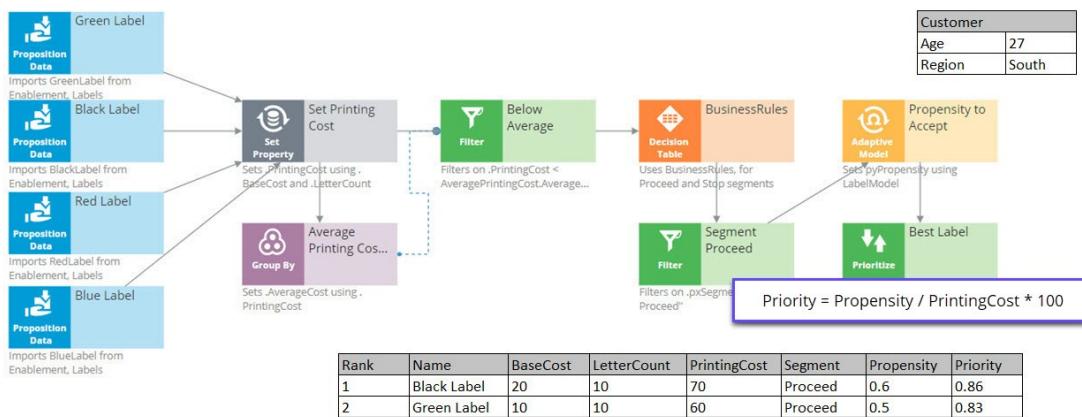
For the Black Label action, the likelihood turns out to be 0.6.

After consulting the Adaptive Model, the Propensity to Accept component sets the Propensity property value for each action.

Remember, the propensity is always a number between zero and 1.

It shows something along the lines of, half of the customers that are like this customer accepted the Green Label action in the past, and 3 out of 5 customers like this customer accepted the Black action last month.

The next component in our chain, called Best Label, is the Prioritize component. This component determines the priority of each action and ranks them. Let's see how this works.



A key element of this component is the priority Expression, which calculates a priority value for each action. According to this Expression, the higher the value, the higher the priority and rank.

In this case, the priority calculation weighs likelihood of acceptance in its equation: 'Propensity divided by PrintingCost times 100'.

When performing this calculation on the Black Label action, we can see that it has a PrintingCost of 70 and a Propensity of 0.6, therefore its Priority is 0.86.

The Green Label action has a lower PrintingCost and a lower Propensity, resulting in a Priority of 0.83.

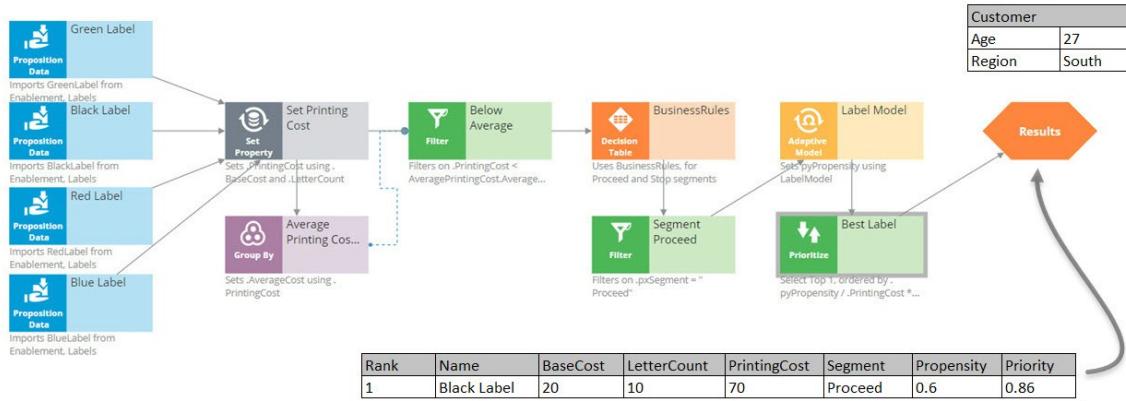
Because 0.86 is higher than 0.83, the Black Label action is now ranked number one.

So, even though the printing cost of the Black Label action is higher than that of the Green Label action, the Black Label action still comes out on top.

In this case, the Priority component reversed the Ranks of the two actions. Black is now the primary action and Green is the secondary action.

The same Prioritization component is also configured to output only the top action.

Therefore, it filters out the Green action altogether, and at the end of our strategy chain, the Black Label is left as our best action.



Creating predictions

Description

Predicting customer churn is one among many business use cases involving predictive models. Pega Customer Decision Hub™ uses predictions that use predictive models to improve one-to-one customer interactions. Learn how to create a new prediction in Prediction Studio and use the new prediction in an engagement strategy.

Learning objectives

- Create a new prediction that uses a predictive model
- Use a prediction in an engagement strategy

Predictive models drive predictions

Introduction

Learn how to better address your customers' needs by predicting customer behavior and business events with predictions. For example, you can determine the likelihood of customer churn or chances of successful case completion.

Predictions combine predictive analytics and best practices in data science.

Transcript

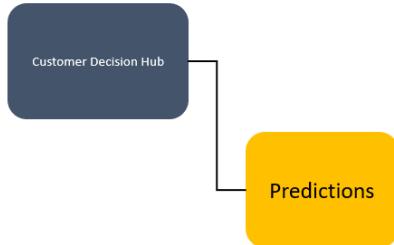
This video shows you how predictions can predict events in a business activity.

You can create and use predictions for several use cases.

Pega Customer Decision Hub™ predictions are used in decision strategies to optimize engagement with customers.

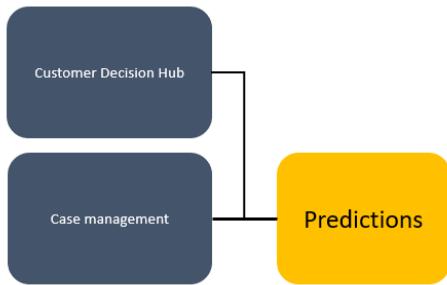
The most common predictions are provided out of the box. You can use the default predictions or create your own.

For example, you can use predictions to predict whether customers are likely to accept your offer or cancel a subscription.

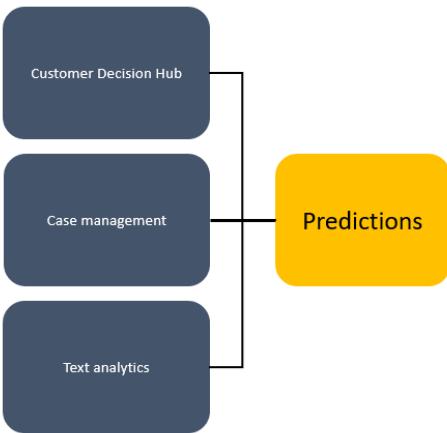


Case management predictions are used in case automation to route cases or prioritize work.

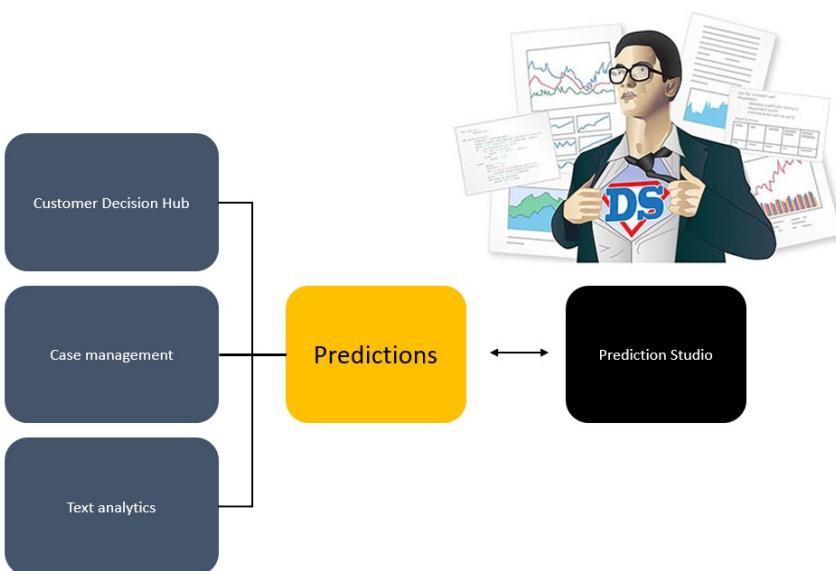
For example, you can create a prediction that predicts whether a case is likely to be completed successfully, or the probability that a claim is fraudulent, which might be useful in straight-through processing of claims with a low score.



Text analytics predictions analyze the text that comes through your channels, such as email or chat, to predict the topic or sentiment of the text.



Predictions are created and managed in Prediction Studio, the dedicated workspace for data scientists.



A next-best-action designer can include the predictions in decision strategies in Customer Decision Hub, to better adjust to customer needs and achieve business goals at the same time.

This handover of predictions from the data scientist to the next-best-action designer contributes to a separation of concerns.

Predictions combine predictive analytics and best practices in data science.

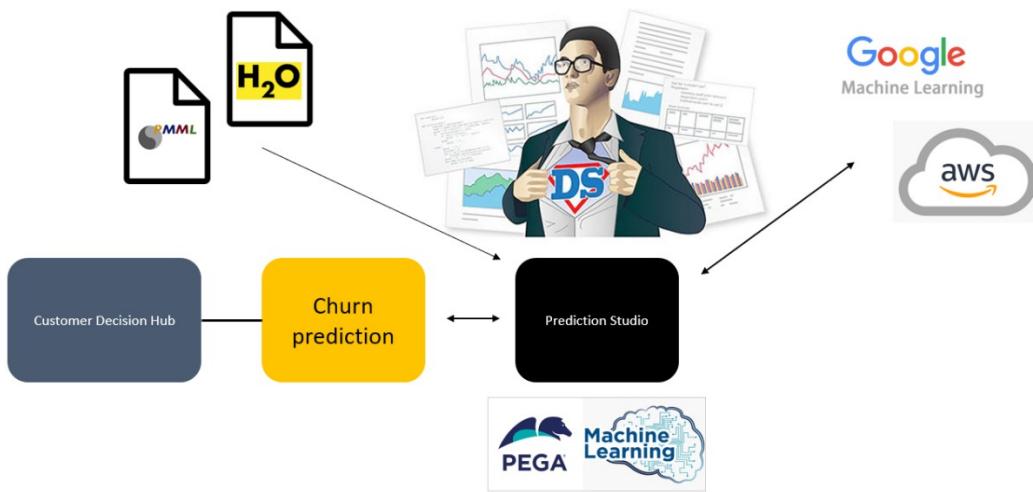
A predictive model drives a prediction. A data scientist can replace the predictive model at any time. However, the prediction always predicts the same outcome.

For example, you can create a new Customer Decision Hub prediction that calculates the likelihood that a customer might end a subscription soon.

To drive the prediction, a data scientist creates a predictive churn model on a platform of choice.

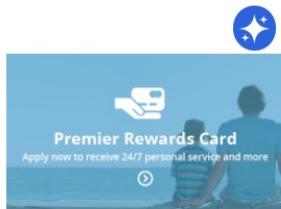
Prediction Studio supports the PMML and H2O model file formats. It can connect to machine learning services like Google ML and Amazon SageMaker.

You also have the option to use Pega machine learning.



One of the predictions that is shipped with Customer Decision Hub, **Predict Web Propensity**, calculates the likelihood that a customer clicks on a web banner.

Customer Decision Hub decides which banner to show based on the calculated propensity and weighting in business requirements and context.



Premier Rewards card

Receive 24/7 personal service and more

[Learn more](#)

In this case, the Premier Rewards card has a propensity of 0.8, and it is ranked as the next best action.

The **Predict Web Propensity** prediction is driven by an adaptive model that learns from each customer interaction.

Best practices in data science include the use of a control group. Customers in the control group are shown a random banner instead of one based on the calculated propensity.

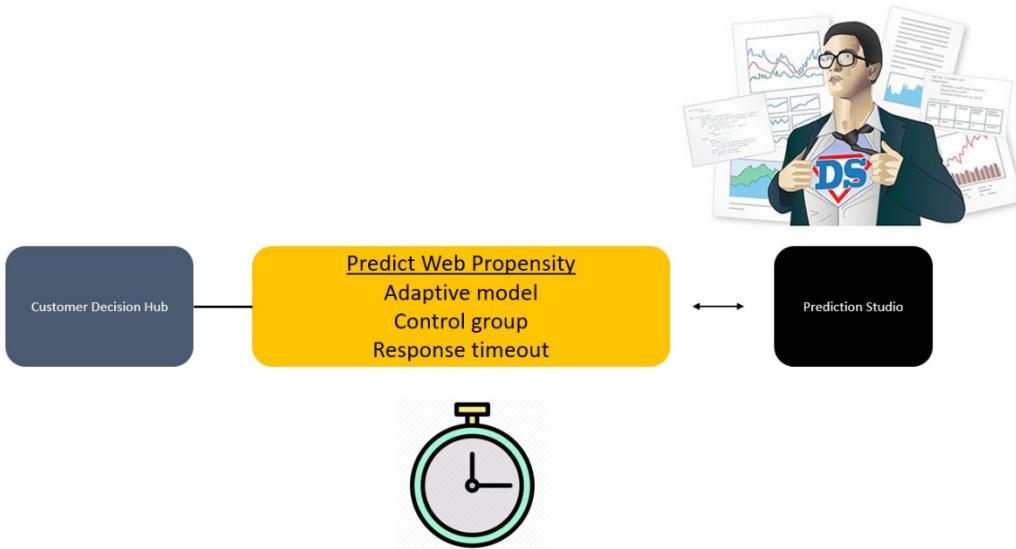
Comparing the success rate in the control group and the test group measures lift, which is the boost in success rate that the prediction generates. Lift is an important business metric.

Also, the use of a control groups allows continuous exploration by the predictive models.

When the lift drops significantly over time, Prediction Studio notifies the data scientist.



The response timeout is built into predictions. When the timeout expires, NoResponse is automatically recorded as the outcome of the interaction.



You have reached the end of this video. What did it show you?

- How predictions can predict events in the business activity
- How a prediction uses a control group to measure the boost in success rate that is generated by the prediction
- How NoResponse is automatically recorded as the outcome of the interaction when the response timeout expires

Creating a prediction

Introduction

Acquiring new customers can be more costly than retaining current customers. U+ Bank uses Pega Customer Decision Hub™ for their customer engagement and wants to reduce the churn rate. Learn how to create a new prediction in Prediction Studio that calculates the likelihood that a customer might churn in the near future.

Transcript

This demo shows you how to create a new prediction in Prediction Studio.

Predictions combine predictive models and best practices in data science.

U+ Bank uses Pega Customer Decision Hub™ to personalize the credit card offered to customers on their website.

If a customer is eligible for multiple offers, artificial intelligence (AI) decides which offer to show.

For customers that are likely to leave the bank soon, the bank wants to make a proactive retention offer instead of a credit card offer.

The bank has recorded historical churn data for its customer base, and a data scientist used this data to create a predictive churn model.

With this model, you create a prediction to use in Customer Decision Hub to display a retention offer to customers with a high churn risk on the website.

Predictions are managed in Prediction Studio.



You can create three types of predictions.

To improve customer engagement with retention offers, choose Customer Decision Hub.

Predictions for case automation and text analytics are also available.

To create a prediction that aims to calculate the likelihood that a customer might churn, set the outcome to **Churn** and the subject of the prediction to **Customer**.

Notice that initially, a placeholder scorecard is generated and used to drive this prediction.

This placeholder is useful in case you do not have a predictive model yet, as it allows the Next-Best-Action specialist to continue work while a predictive model is built.

Churn			
Name	Type	Performance	Status
Test2	Scorecard	—	ACTIVE

As you already have a predictive model, the next step is to replace the scorecard that drives this prediction with the predictive churn model.

When the replacement is ready for review, approve the candidate model, and save the configuration.

Churn			
Name	Type	Performance	Status
Churn	Predictive model	—	ACTIVE

Once the prediction is created, test your work. Select a persona as the data source and run the prediction.

Troy is predicted to leave the bank in the near future; therefore, the outcome is churn.

Barbara has a low propensity to churn; therefore, the outcome is loyal.

The prediction is now ready for use in Customer Decision Hub.

You have reached the end of this demo. What did it show you?

- How to create a new prediction
- How to replace the generated scorecard with a predictive model in a prediction

Using predictions in engagement strategies

Introduction

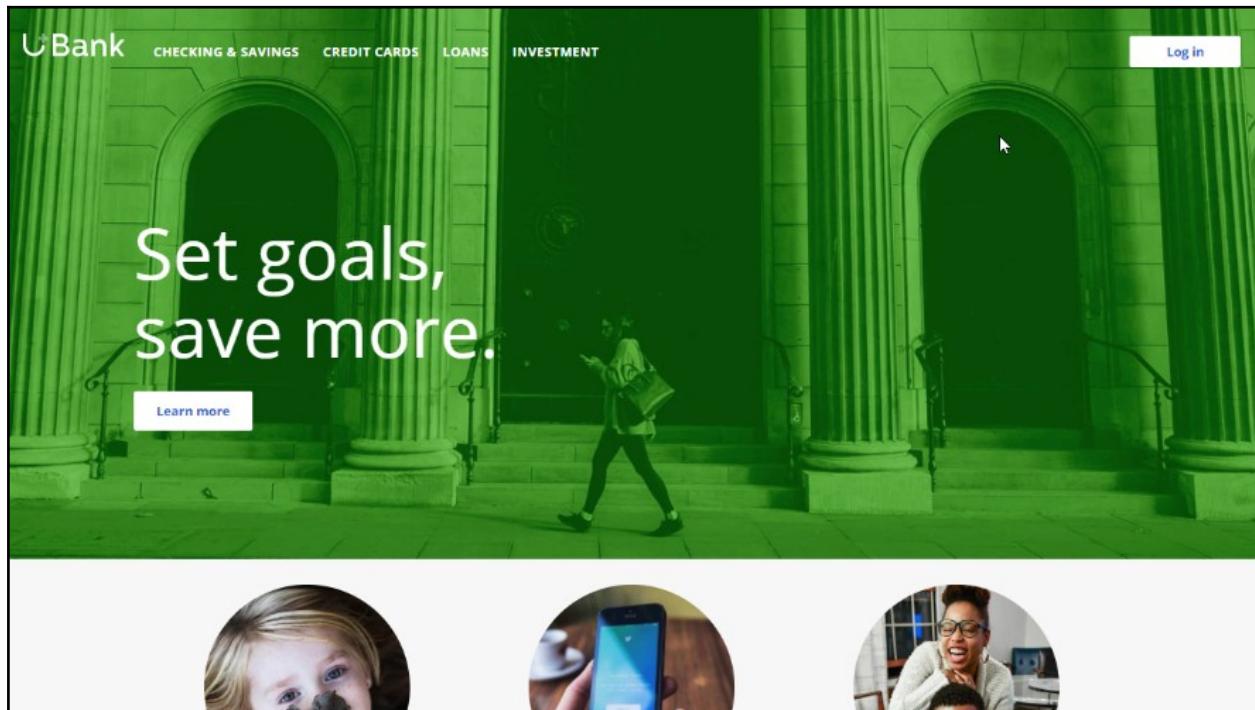
A prediction is used to predict customer behavior such as offer acceptance and churn based on characteristics such as credit risk, income, and product subscriptions. Learn how to arbitrate between different groups of actions to display more relevant offers to customers. Gain experience using a prediction in a decision strategy and learn how applicability rules can be defined to reflect the bank's requirements in a decision strategy.

Transcript

This demo shows you how to use a prediction in an engagement strategy to determine customer applicability for a retention offer.

Currently, U+ Bank is cross-selling on the web by showing various credit cards to eligible customers who log in to its website.

The bank now wants to show a retention offer, instead of a credit card offer, to customers who are likely to churn in the near future. The credit card offers are shown only to loyal customers.



To meet this business requirement, a decisioning administrator has already set up the taxonomy by defining a new business issue called Retention, and an offer group.

Taxonomy

[Business structure](#)

 **Business structure**

Issues / Groups

Retention

 ExtraMiles

Sales

 CreditCards

This ExtraMiles group contains a retention offer, Extra Miles 5K.

Actions

Search Issue / Group

Showing 1 of 1 results

[Extra miles 5K](#)
ExtraMiles5K

The next step is to create an applicability condition that makes a customer qualify for a retention offer when there is a high likelihood that the customer might churn.

A data scientist has created a prediction that identifies these high-risk customers.

When you open the prediction in Prediction Studio, notice that the possible response labels are **Churn** and **Loyal** to predict customer behavior.

The result of the prediction is stored in the pxSegment property.

Response labels

Labels for the possible values of the responses.

Churn

Target label Alternative label

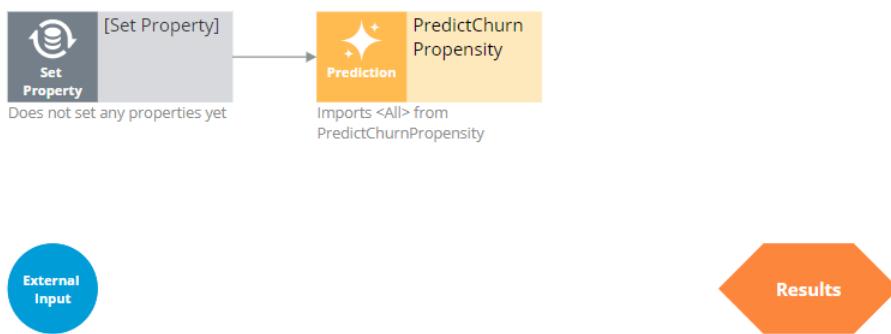
Churn Loyal

To define the applicability condition, you create a decision strategy to output a retention offer only if the response label of the prediction is **Churn**.

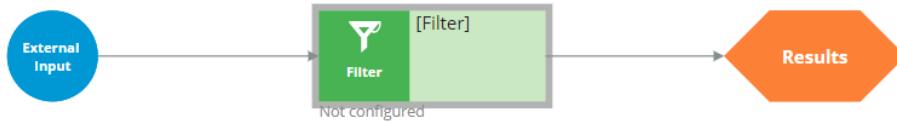
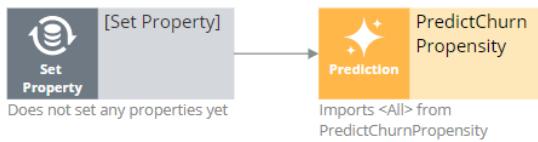
Add a **Prediction** component to the canvas and configure it to reference the churn prediction.

Add a **Set Property** component and connect it to the Prediction component.

You can configure the Set Property component at a later point to accommodate parameterized fields.



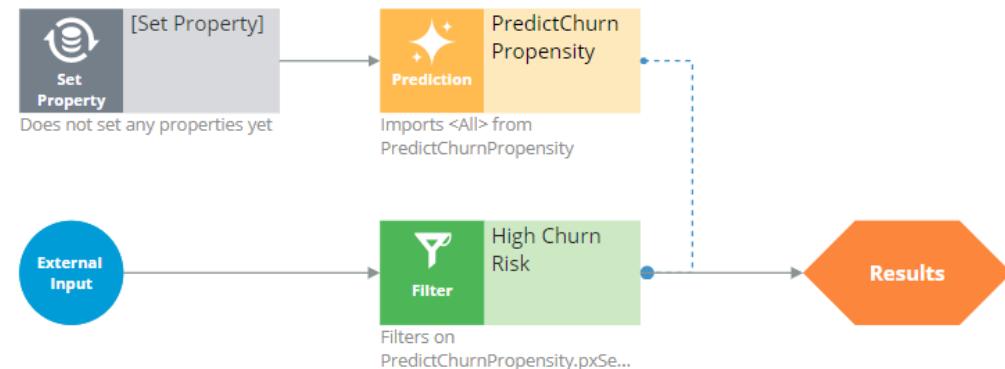
Next, add a filter component to filter out the loyal customers and pass retention offers to high churn risk customers only.



The filter condition is defined to output a retention offer when the pxSegment property of the prediction is equal to **Churn**.

Expression builder **Browse** **Test**

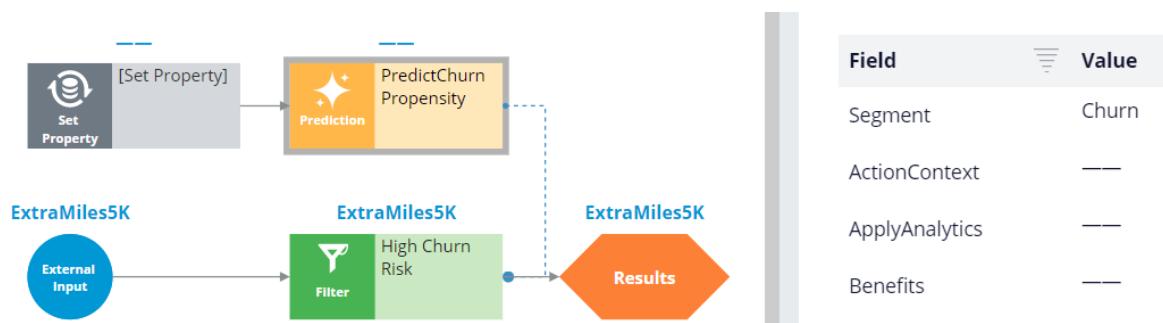
`| 1 PredictChurnPropensity.pxSegment = "Churn"`



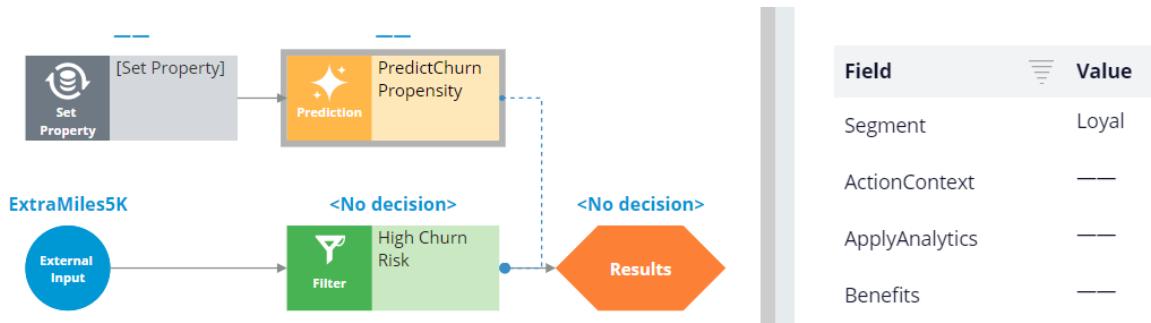
Next, test the strategy using two customer profiles, **Troy** and **Barbara**.

For external inputs, consider all available retention offers.

The strategy outputs a result for **Troy** because the result of the prediction is **Churn**.



The strategy does not have a result for **Barbara**, because the Segment value is **Loyal**.



By checking in the strategy, you commit your changes so that they go into effect.

You can now use this strategy in the Next-Best-Action Designer engagement policy as an applicability condition.

The first business rule you need to implement is that the **ExtraMiles** group is applicable only to high churn risk customers.

To implement this rule, in the **Applicability** section, define a condition for the customer field.

Select the **RetentionStrategy**. The condition is: the RetentionStrategy has results for the High Churn Risk component.

The second business rule you need to implement is: U+ Bank wants to show credit card offers to low-risk customers only; meaning the **CreditCards** group is not applicable for high-risk customers.

To implement this rule, modify the **Applicability** section of the **CreditCards** group.

The condition is: the RetentionStrategy doesn't have results for the High Churn Risk component.

Once the applicability conditions are defined, you need to amend the **Channels** configuration.

Because U+ Bank introduced a new group, **ExtraMiles**, which belongs to a new business issue, **Retention**, you need to select the results from the appropriate business structure level.

In this case, the bank wants to arbitrate between two different business issues: Sales and Retention. Therefore, select All Issues/All Groups from the business structure level.

Saving the configuration implements the business requirement.

On the U+ Bank website, when you log in as **Troy**, notice that the retention offer is displayed because **Troy** is predicted to churn in the near future.

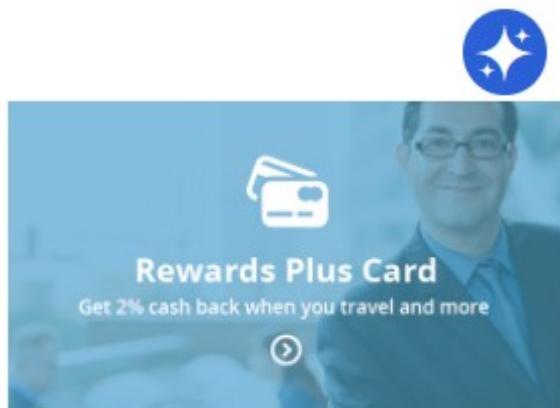


Extra miles 5K

5,000 extra miles

[Learn more](#)

Now, when you log in as **Barbara**, notice that the credit card offer is displayed because she is predicted to remain loyal for now.



Rewards Plus card

Get 2% cash back when you travel and more

[Learn more](#)

You have reached the end of this demo. What did it show you?

- How to use a prediction in a decision strategy
- How to arbitrate between different groups of actions to display more relevant offers to customers
- How to define applicability rules using a decision strategy in Next-Best-Action Designer

Creating predictive models

Description

In Prediction Studio, three options to leverage historical data are available: creating models using Pega machine learning, importing models created in a third party tool and referencing external models. Learn how to create, import and reference predictive models that can be used in driving the next best action.

Learning objectives

- Describe the role and usage of predictive models in the Pega landscape
- Use Pega machine learning to build predictive models
- Import third party predictive models
- Use machine learning services
- Explain the model transparency settings in Prediction Studio

Predictive models

Introduction

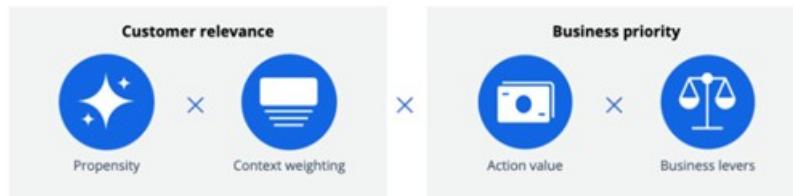
Enhance decision strategies with predictive models built on customer interaction data and let Pega Customer Decision Hub™ bring even more relevance to every customer engagement. Build models using Pega's machine learning capabilities, import models built with third-party tools and incorporate the latest AI algorithms into the Pega AI engine by connecting to the Google AI Platform and Amazon SageMaker machine learning services.

Transcript

This video will describe the use of predictive models to enhance the next best actions that Customer Decision Hub generates.

Next best actions balance customer relevance and business priorities by selecting the actions with the highest priority.

The priority is calculated by multiplying the values for propensity, context weighting, action value and business levers.



Propensity is the likelihood of a customer responding positively to an action by, for example, clicking on a web banner or accepting an offer.

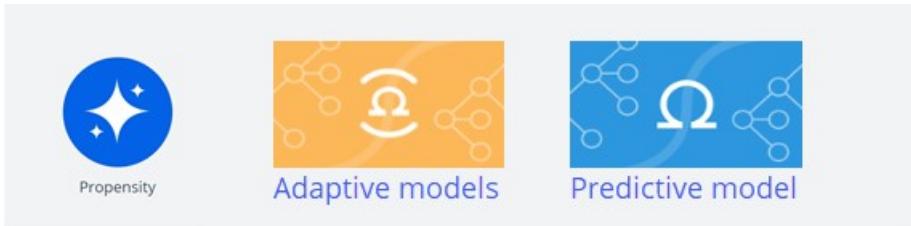
This is calculated by predictive models. In Pega, self-learning Naive Bayes models, which are generated for each action, are a key feature.



These adaptive models are automatically updated after new responses have been received and can start without any historical information because they learn on the fly.

When the use case requires a more advanced modeling technique, for example to predict customer churn or to estimate credit risk ...

... Prediction Studio offers several methods to create the artifacts that represent an actual predictive model or that reference a predictive model.



The first method is to use Pega machine learning. You can import a file containing the historical customer interaction data set and build a model in Prediction Studio.

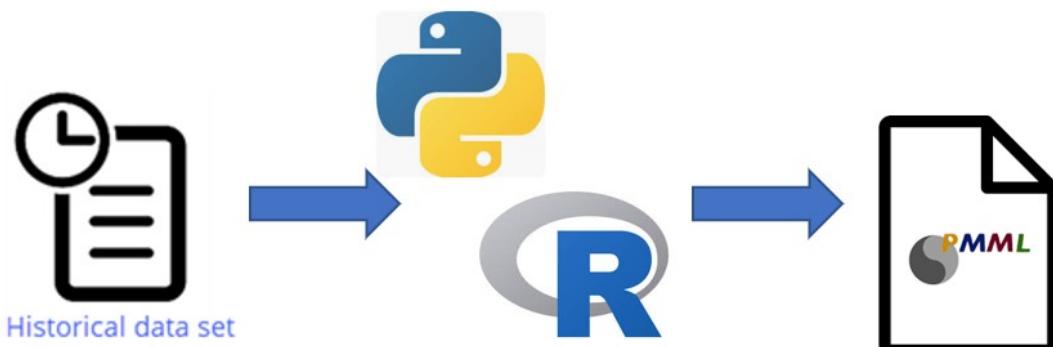
This model can then be used in decision strategies. When the decision strategies execute, the models are executed inside the Pega platform.



The second option is to import an existing model. You can build a model using a third-party tool like R or Python and export it as a PMML file.

PMML is an XML-based standard that is designed to facilitate the exchange of models between applications.

Import the PMML file into Prediction Studio and map its predictors to the fields in the customer data model.



Similarly, you can import model files that have been generated in H2O.ai. H2O is a modelling platform, and the procedure for using the generated model file is identical to that for a PMML file.



Just like with Pega machine learning models, the imported model can then be used in decision strategies.

When decision strategies using the imported models execute, the models are executed inside the Pega platform.

The third option is to reference a model on an external platform like the Google AI Platform.



Just like with Pega machine learning models, the referenced model can then be used in decision strategies.

In this case, when the decision strategy requires a prediction, a request is sent to the external model, which calculates the outcome and sends it back to Pega.

Like with the Google AI Platform, you can connect to AWS SageMaker and run your model remotely.



To summarize, you have three options for leveraging predictive models built on customer data.

You can build models using Pega machine learning, you can import models built with third-party tools, and you can use machine learning services to reference predictive models.

When the decision strategies using predictive models execute, the models are executed inside Pega or externally by Google ML and the Amazon SageMaker platform.

Building models with Pega machine learning

Introduction

Prediction Studio offers several options for leveraging customer data to create predictive models. Learn how to develop powerful and reliable models that can predict customer behavior, such as offer acceptance, churn rate, credit risk, or other types of behavior by using Pega machine learning.

Transcript

This demo will show you how to build a predictive model using Pega machine learning in Prediction Studio.

In an effort to proactively prevent churn, U+ Bank wants to predict the likelihood that a customer will leave the bank in the near future.

When starting to build a new model, you will be presented with the option to create a model on a template that is used for streamlining model development. One of these is churn modeling.

New predictive model X

Name *
ChurnPegaML

Create model [?](#)

Use Pega machine learning Import model Select external model

Category Retention Template Churn Modeling

Churn Modeling

Aims at ordering cases in terms of their propensity to churn within a defined length of time. Score bands are created to enable cases with different levels of propensity to be selected or deselected. Behavior: Churn can be defined as closure of a relationship in a following period (e.g. within three months after the potentially predictive data was captured). Cases can be restricted to those who suffered some adverse experience or those who would be targeted by some competitive offer. Predictions: In addition to the probability of churn, a model may analyze and forecast the probability of each possible reason for dormancy and retention.

The model build itself consists of 5 steps: Data preparation, data analysis, model development, model analysis, and model selection.

In the data preparation step, the data source containing the historical data is selected, the sample is constructed, and the outcome of the model is defined.

The data source can be a csv-file, a database table, a data flow or a data set.

Source selection

Select the data source for the creation of predictive models and preview the first 100 records.

CSV Database Data flow Data set

Upload flat file No file chosen

Separator character: , Quote character: " First line contains field names:

The preview of the first ten records in the data set allows you to verify that all fields will be correctly imported.

Preview for first 10 records of historical_data.csv

Field	Record 1	Record 2	Record 3
CustomerID	14	15	16
ACCOUNT_ID	—	—	—
Title	—	—	—
pyFullName	Troy Murphy	Barbara Stockton	Joanna Williams
Gender	M	F	F
Age	26	32	25

Next, construct the sample.

Using a weight field is optional; it is only used when the data source contains such a field. If you do not specify the field, each case counts as one.

The type of field to be sampled can be set to either numeric or categorical.

Select the fields to sample

Field	Type
CustomerID	Categorical
ACCOUNT_ID	Categorical
Title	Categorical

By default, all fields are considered potential predictors. When setting predictors, it's important to use some common sense.

For example, the customer ID is a random number and has no impact on the behavior to be predicted.

Likewise, the name of the customer has no predictive value. For such fields, change the type to 'Not used'.

If the data contains a relatively small number of cases, you will want to use 100% of the records. If the data source is large, a sample will be sufficient.

Select sampling method

- Uniform sampling
 Stratified sampling

Set sample size using

% or Cases

Next, you define the hold-out sets for validation and testing during model development. Your models will be trained with the remainder of the data.

Once trained, the validation set is used to check for robustness of candidate models and to compare their performance.

Finally, the test set is used to analyze the performance characteristics of candidate models, and to select the best model.

Hold-out sets

Split the sample into a development, validation and test set. [?](#)

Create hold-out sets by

- Setting percentages for each set
 User defined field

Retain % of the sample for validation (201 cases)

and % of the sample for testing (201 cases)

60.0 % of the sample for development (604 cases)

Finalize the data preparation step by defining the outcome to be predicted.

You can predict a binary outcome, as in this example, or predict a continuous outcome.

For a binary outcome type, the outcome field must be categorical. For a continuous outcome type, the outcome field must be numerical.

Here you also map the values of the outcome field to the outcome category. With that, you specify how to differentiate between good and bad behavior.

Outcome definition

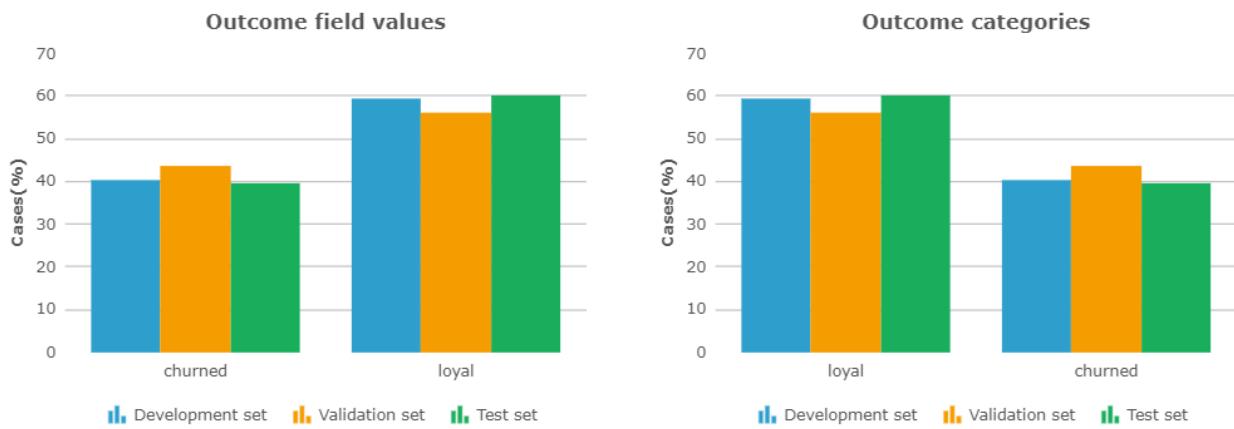
Define the outcome to be predicted. Predict a binary outcome, only categorical fields can be selected for this. Or predict a continuous outcome, only numeric fields can be selected for this.

Outcome type	Outcome field to predict
Binary	Segment

Map possible values of outcome field to outcome category

	Value	to	Outcome category
Map	loyal	to	loyal
Map	churned	to	churned

It is worthwhile to verify that the customer distribution across the development data set is similar to the whole sample.



In the data analysis step, you analyze the individual predictors. By default, only predictors with a performance higher than 52 are included.

For fields that have a very high performance, the Role is set to *value* to protect models from accidentally using predictors that might be directly correlated to the outcome.

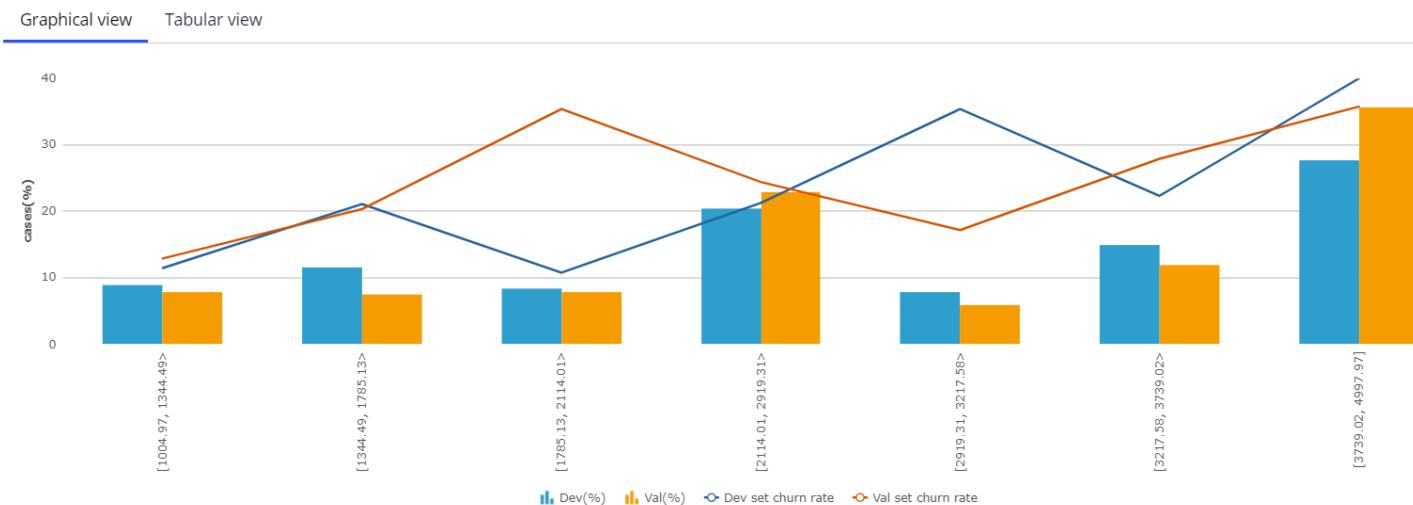
Exclude predictors with a performance below	52.0	Apply			
<input type="button" value="Change role"/> <input type="button" value="Reports"/> <input type="button" value="New virtual field"/> <input type="button" value="Reset"/> <input type="button" value="Show groups"/>					
<input type="checkbox"/> Predictor	Type	Role	Binned intervals	Grouped intervals	Grouped performance
<input type="checkbox"/> RiskScore	Numeric	VALUE	200	11	90.27
<input type="checkbox"/> AverageSpent	Numeric	PREDICTOR	200	7	68.28
<input type="checkbox"/> MonthlyPremium	Numeric	PREDICTOR	200	10	64.84
<input type="checkbox"/> Age	Numeric	PREDICTOR	70	8	63.87

You can also manipulate features to create a better predictor by creating a 'New virtual field'. This is a fundamental step towards having good models.

Income*CLV is such a virtual field. The performance of this new predictor is higher than that of the individual fields.

Predictor	Type	Role	Binned intervals	Grouped intervals	Grouped performance
RiskScore	Numeric	PREDICTOR	200	11	90.27
AverageSpent	Numeric	PREDICTOR	200	7	68.28
Income*CLV	Numeric	PREDICTOR	200	10	65.08
MonthlyPremium	Numeric	PREDICTOR	200	10	64.84
Age	Numeric	PREDICTOR	70	8	63.87
Income	Numeric	PREDICTOR	200	10	63.74
CLV_VALUE	Numeric	PREDICTOR	200	9	63.47

Data analysis creates a binned, ordinal view of individual predictors. Both Binning and Granularity are automatically set but can be manually adjusted.



As part of model development, the grouping and predictor selection process is automated.

When multiple predictors are correlated, considering them all for the machine learning process will lead to unnecessary model complexity.

It is best practice to select the best performing predictor in each group.

Predictor grouping

In predictor grouping, correlated predictors are grouped all predictors (default) or alternatively, continue with t

Grouping level

Prediction Studio provides a rich model factory that supports industry standard models.

You can create 4 types of models: Regression models, Decision tree models, Bivariate models and Genetic algorithm models.

By default, a Regression and a Decision tree model are automatically created. These models are highly transparent.

Bivariate models and Genetic algorithm models have a lower transparency score.

Bivariate model Pega Compliance All business issues	3	Genetic algorithm Pega Compliance All business issues	2	Regression Compliance All business issues	4	Tree model Compliance All business issues	5
--	---	--	---	---	---	---	---

The purpose of the next step, Model Analysis, is to select the best model for your use case.

In the 'Score comparison' step, you can compare the scores generated by the models in terms of behavior, lift, gains and discrimination.

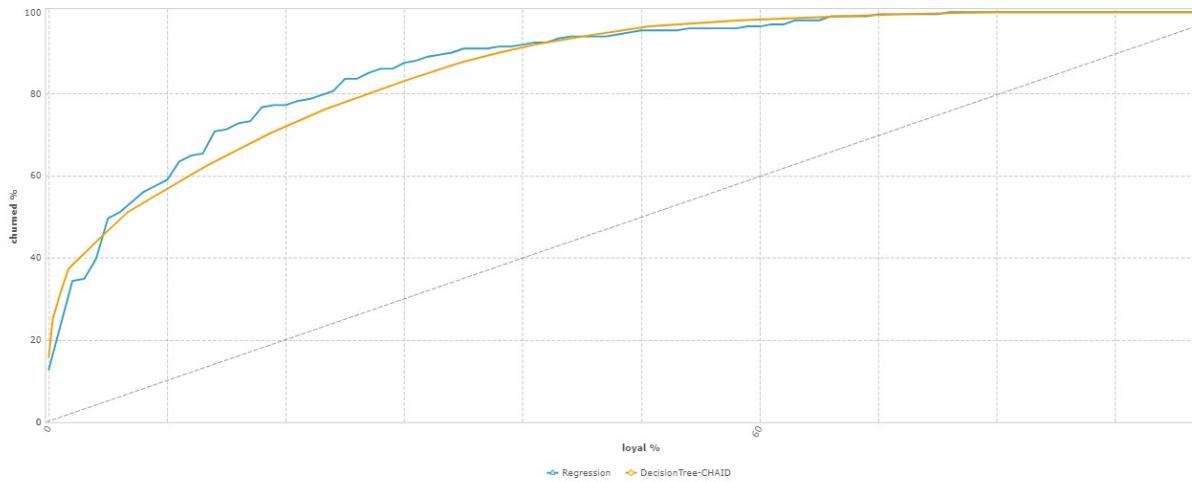
Prediction Studio uses Area Under the Curve (or AUC) to measure the performance of predictors and models.

You can describe AUC as the measure of how well the model is able to discriminate between good and bad cases.

The value of AUC ranges from 50%: random distribution, to 100%: perfect discrimination.

Behavior Lift Gains Value Discrimination

Export as CSV



In the 'Score distribution' step, the model scores are segmented based on a method you select. A typical example divides the scores into deciles: 10 classes with an equal number of cases.

Graphical view Tabular view



The 'Score distribution' settings give several methods for defining these segments.

Score distribution settings

Segmentation method

- Create bands with equal number of cases
- Create bands with equal number of cases**
- Create statistically significant bands
- Create monotonically increasing bands
- Create user defined bands

Max. # of bands *

10

Number *

100

or

Percentage *

9.94

Only count cases where the outcome equals

In the 'Class comparison' step, you can analyze and compare models after the score distribution has been adjusted.

Finally, you select the model that best fits your needs and specify the context in which to save it. The default context where the models are saved is the customer class.

Before you can save the model, check the mapping of the predictors to the properties in the customer class.

If the properties exist and have a name similar to a predictor field name, they will be mapped automatically.

You also have the option to create missing properties, but this should be discussed with the system architect beforehand.

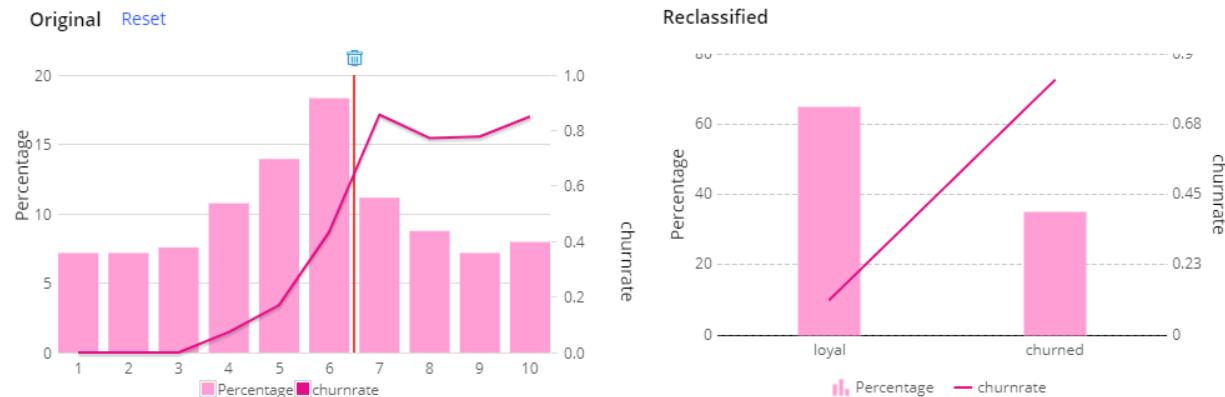
Monitor Model Mapping

The model predictors are automatically mapped to fields in the data model.

[Create missing fields](#) [Refresh mapping](#)

Predictor	Predictor Type	Field
Gender	Symbolic	.Gender
Age	Numeric	.Age
MaritalStatus	Symbolic	.MaritalStatus

If needed, you can adjust the score distribution segments by clicking on the original score distribution chart. In this example, two segments are appropriate: loyal and churned.



The model can now be saved and is ready for use in a decisioning strategy.

You have reached the end of this demo. What did it show you?

- How to create a predictive model in Prediction Studio using Pega machine learning.

Importing predictive models

Introduction

During a Pega Decision Management implementation project, you may discover that the company already uses predictive models. These assets can be reused in Pega Decision Management to help make customer predictions.

Transcript

This demo will show you how to import third-party predictive models into Prediction Studio and use them natively in Next-Best-Action strategies.

Prediction Studio supports two external model formats.

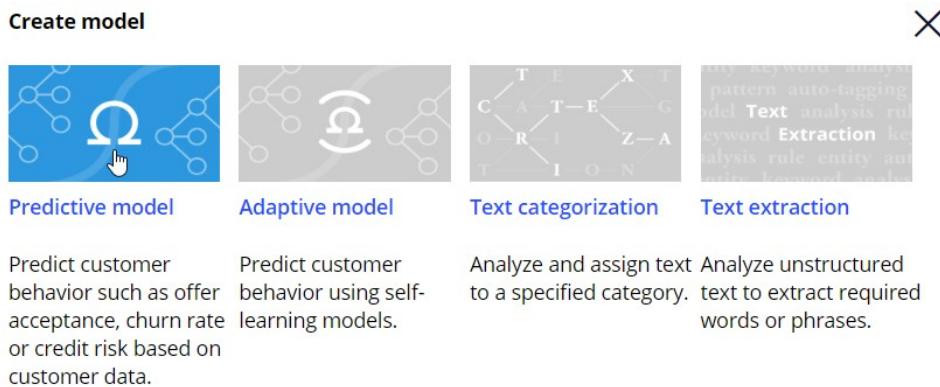
First, you can import models in the Predictive Model Markup Language (PMML) format.

PMML is an XML-based language aimed at easily sharing predictive models between applications.

It is the de facto standard for representing not only predictive models, but also data, pre- and post-processing.

Additionally, you can import models built with H2O.ai, an open source machine learning and predictive analytics platform that allows you to build machine learning models on big data.

The processes for importing PMML and H2O models are identical and start with creating a new predictive model strategy component.



Prediction Studio offers three options for creating a predictive model: using Pega machine learning, importing a previously built model, or using an external model.

To leverage an existing model file, select the **Import model** option.

Upload the PMML or H2O model file.

The default context of the model is the **Customer class**, where the customer data model properties are stored. You can change this class if required.

New predictive model

Name *
ChurnPMML

Create model [?](#)

[Use Pega machine learning](#) [Import model](#) [Select external model](#)

Import model file * [?](#)

[Choose File](#)

File name
ChurnPMML.pmml

Context
Customer [Change](#)

In the **Outcome definition** dialog box, you define which probability you would like to predict ...

... and the expected performance of the model, which is used as a benchmark when monitoring the model.

New predictive model

Outcome definition [Set labels](#)

The objective of the model is to predict Segment

Predict the probability of churned loyal

Modeling technique [X](#)
Tree model

Expected performance (AUC) [?](#)
80

Import the model and, on the **Mapping** tab, make sure that all predictors are mapped to fields in the data model.

Missing fields can be created, but this should be discussed with the system architect beforehand.

The model predictors are automatically mapped to fields in the data model.

[Create missing fields](#)

[Refresh mapping](#)

After the model is saved, you can test it for a single customer or run it for a batch of customers.

Run predictive model

Troy

Data Transform

Field name	Type	Input
Age	Double	26
Gender	string	M
NetPromoterScore	Double	9
MaritalStatus	string	Married
AverageBalance	Double	1500.67
AverageSpent	Double	3200.53
EmailOptIn	string	Y
DebtToIncomeRatio	Double	45
SMSOptIn	string	Y
MonthlyPremium	Double	0.0
HasMortgage	string	Y
DMOptIn	string	Y

Run

When you test the model for a single customer, you can use a data transform as input data.

When customer Troy is used as the data source, the model predicts that he is likely to churn.

The model also outputs his propensity to churn, which is, in this case, 93.42%.

Run predictive model

Single run

Inputs

Outputs

Results

Result	Monitoring performance
churned	0
Propensity	Monitoring evidence
0.9342621091861922	0.0

Output

Segment	Value
	churned

In contrast, the model predicts that customer Barbara is likely to remain loyal, with a low propensity to churn of 35.83%.

Run predictive model

Single run Batch run

Inputs

Outputs

Results

Result	Monitoring performance
loyal	0
Propensity	Monitoring evidence
0.3583554398897344	0.0

Output

Segment	Value
	loyal

You can also run the model on a batch of customers.

When the model is run for a larger input data set, the output shows the number of customers that are classified as either likely to remain loyal or likely to churn in the near term.

Run predictive model

Single run Batch run

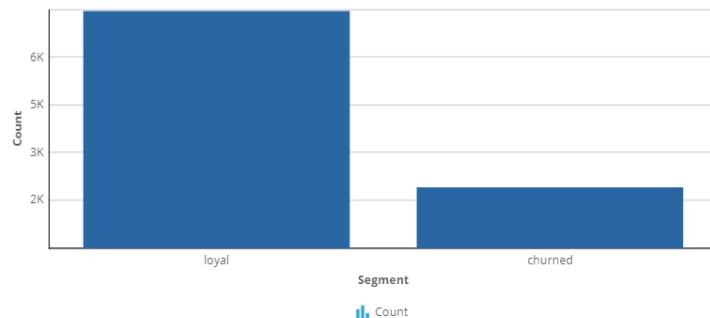
Data source*
CustomerBatch
This data source contains approximately 10,000 records.

Source type
Data set

Total records executed: 10000 Total failed: 0

Output

Segment ▾



You have reached the end of this demo. What did it show you?

- How to import third-party predictive models into Prediction Studio.
- How to test the model for a single customer.
- How to run the model for a batch of customers.

Using machine learning services

Introduction

Enhance the Pega AI engine with the latest AI algorithms by connecting to models in Amazon SageMaker and Google AI Platform machine learning services. Learn how to leverage a model, created in and running on Amazon SageMaker, in Pega's Prediction Studio.

Transcript

This demo will show you how to leverage a machine learning service by running a churn model created externally and using its outputs in Pega Prediction Studio.

We will showcase this using Amazon SageMaker. The steps are similar to using other machine learning services such as Google AI Platform.

Using a machine learning service instead of a model that runs locally may involve costs and possible down time of the service.

However, for certain use cases such as churn or credit risk models, machine learning services can be the optimal choice.

To showcase how to use a churn model created in Amazon SageMaker, let's first consider the high-level steps involved in creating a machine learning model.

Amazon SageMaker allows you to build, train and deploy machine learning models in a fully managed service.

The Autopilot feature automates this process and trains and tunes the best machine learning models for classification or regression, based on your data.

After setting up your AWS environment, you can open Amazon SageMaker Studio to create a new Autopilot experiment.

The screenshot shows the Amazon SageMaker Studio interface. The left sidebar has a red box around the 'Amazon SageMaker Studio' link. The main content area has a box around the 'AWS Marketplace' section. The 'Overview' section is detailed below:

Category	Description	Sub-options
Ground Truth	Set up and manage labeling jobs for highly accurate training datasets using active learning and human labeling.	Labeling jobs
Notebook	Availability of AWS and SageMaker SDKs and sample notebooks to create training jobs and deploy models.	Notebook instances
Training	Train and tune models at any scale. Leverage high performance AWS algorithms or bring your own.	Training jobs Hyperparameter tuning jobs
Inference	Create models from training jobs or import external models for hosting to run inferences on new data.	Models Endpoints Batch transform jobs
Processing Run	Pre- or post-processing and model evaluation workloads with a fully managed experience.	Processing jobs

In the Job settings, select the data file you want to build the model on, specify the outcome field, choose the location where the output should be stored and create the experiment.

Create Amazon SageMaker Autopilot Experiment

JOB SETTINGS

Experiment Name
ChurnAWS
Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Input data location (S3 bucket)
Enter the location in S3 where your training data is stored. You can point to a single data file, an S3 object key prefix that contains only data files, or a manifest file that contains the location of your input data. See more in the [AWS Docs](#)

Find S3 bucket Enter S3 bucket location

Note: The S3 bucket must be in the same AWS Region where you're running SageMaker Studio because SageMaker doesn't allow cross-region requests.

S3 bucket name
sagemaker-studio- us-east-1

S3 object key prefix
historical_churn_data.csv

Is your S3 input a manifest file?
For more information on the format of a manifest file, please see the [AWS Docs](#)

Target attribute name
The target attribute is the attribute in your dataset that you want Amazon SageMaker Autopilot to make predictions for.
Segment

The attribute name is case-sensitive and must match exactly the name in your input dataset

Output data location (S3 bucket)
Enter the location in S3 where you want to store the output.

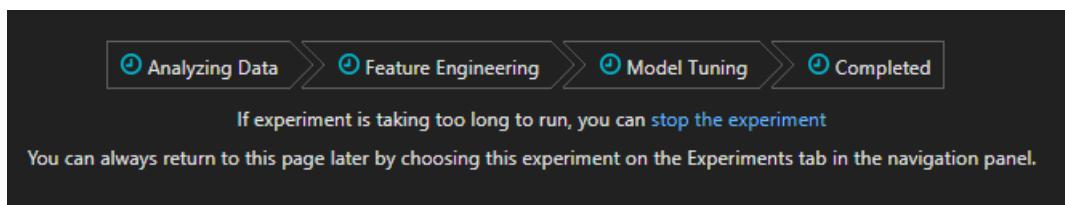
Find S3 bucket Enter S3 bucket location

Note: The S3 bucket must be in the same AWS Region where you're running SageMaker Studio because SageMaker doesn't allow cross-region requests.

S3 bucket name
sagemaker-studio- east-1

S3 object key prefix
historical_churn_data.csv

The Autopilot process analyzes the data, performs a feature engineering step, and tunes the candidate models.



To deploy the best candidate model, select the tuning job with the highest Objective value. This value indicates the predictive power of the model.

EXPERIMENT: CHURNAWS									
Trials		Job profile							
TRIALS									
1 row selected									
Trial name	Status	Start time	Objective		Deploy				
★ Best: tuning-job-1-a89f0f3cd5343889f-205-335f90e8	Completed	2 hours ago	0.9393600225448608						
tuning-job-1-a89f0f3cd5343889f-184-1056fb3	Completed	2 hours ago	0.934220016002655						
tuning-job-1-a89f0f3cd5343889f-171-ce9ec6a4	Completed	3 hours ago	0.934220016002655						
tuning-job-1-a89f0f3cd5343889f-238-e6521619	Completed	2 hours ago	0.934220016002655						
tuning-job-1-a89f0f3cd5343889f-211-e0032194	Completed	2 hours ago	0.934220016002655						
tuning-job-1-a89f0f3cd5343889f-144-27960f16	Completed	3 hours ago	0.934040009754333						
tuning-job-1-a89f0f3cd5343889f-148-5afdfad7	Completed	3 hours ago	0.934040009754333						
tuning-job-1-a89f0f3cd5343889f-248-caa2bc6	Completed	2 hours ago	0.934040009754333						
tuning-job-1-a89f0f3cd5343889f-166-c13901fa	Completed	3 hours ago	0.934040009754333						
tuning-job-1-a89f0f3cd5343889f-168-0ddff684	Completed	3 hours ago	0.934040009754333						

An endpoint that can be reached from Pega is automatically created.

A binary classification, as in this example, predicts if an event will happen or not, based on a cut-off value. By default, the response content for a binary model is set to this 'predicted_label'.

However, it is best practice to include a value for the probability that the event will happen in the response content ...

... as it contains the most information and allows the cutoff value to be adjusted in Pega. Also, it allows for monitoring of the probability with respect to observed outcomes over time.

Deploy model

REQUIRED SETTINGS

Endpoint name

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Instance type Instance count
 ▼

Data capture
 SageMaker Studio will save prediction requests and responses from the endpoint to an Amazon S3 location specified below

Save prediction requests
 Save prediction responses

Inference Response Content
 Select the response content the endpoint should return per input data point. The inference response will be in the order in which the keys are selected.
▼

ADVANCED SETTINGS - Optional

In Prediction Studio, you can define a machine learning service to connect to your cloud service instance.

To move messages securely to and from Pega, the system architect has set up an authentication profile.

AmazonML X

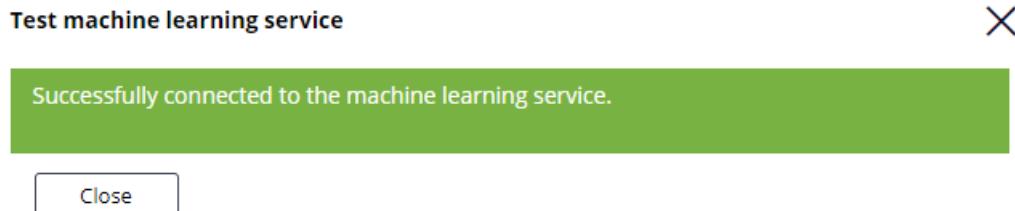
Service type
Amazon Sagemaker

Name
Amazon Machine Learning

Authentication profile * Type AWS
AmazonML

Region
US East (N. Virginia)

Cancel Save



Once the connection to the machine learning is established ...

... start by creating a new predictive model to leverage the service.

Select the machine learning service and the model that you want to reference.

New predictive model

Name *

Create model [?](#)

[Use Pega machine learning](#) [Import model](#) [Select external model](#)

Machine learning service * [Amazon Machine Learning](#)

Model [SageMakerChurn-model](#)

Next, upload the required model metadata file. A template for this JSON file, containing example values, is available for download.

The screenshot shows a Windows Notepad window titled "ModelMetadataTemplate.json - Notepad". The window contains a JSON object with the following structure:

```
{  
    "predictMethodUsesNameValuePair": false,  
    "predictorList": [  
        {"name": "GENDER",  
         "type": "CATEGORICAL"},  
        {"name": "AGE",  
         "type": "NUMERIC"}  
    ],  
    "model": {  
        "objective": "Churn",  
        "outcomeType": "BINARY",  
        "expectedPerformance": 78.5,  
        "framework": "SCIKIT_LEARN",  
        "modelingTechnique": "Naive Bayes Classifier",  
        "outcomes": {  
            "values": ["Yes", "No"]  
        }  
    }  
}
```

The status bar at the bottom of the Notepad window shows "Ln 1, Col 1" and "100%".

The JSON file must contain the list of predictors in the data set and their property type.

It must also contain the objective of the model and the outcome type.

Available outcome types are binary, categorical and continuous.

Optionally, include the expected

performance. The metric for binary models is AUC, F-score for categorical models and RMSE for continuous models.

For SageMaker, the file must include the framework property. This property determines the input format and output format of the model. In Google AI Platform, this property is automatically fetched.

Finally, the metadata file must include the modeling technique and the outcome values.

For binary outcome models, enter the values for the outcome for which you want to predict the probability, and the alternative outcome.

For categorical outcome models, enter all values that represent the possible outcomes.

For continuous outcome models, enter minimum and maximum outcome values.

Best practice is to generate the file as part of the model-building process to avoid human errors.

Next, set the correct context of the model if required. The default context is the customer class.

You can review the model metadata, such as the objective of the model and the type of problem to solve, before proceeding.

New predictive model X

Outcome definition	
The objective of the model is to predict Churn	Modeling technique xgboost
Predicting Two categories	Framework scikit-learn
Predict the probability of churned	Expected performance (AUC) ? 78.5
With alternative outcome loyal	

Back Cancel Create

All predictors must be mapped to the corresponding fields in the data model.

After saving the model, you can run it through the new service connection.

Customer Troy has a high risk of churning; the model returned a high probability to churn for him.

▽ Outputs

Results

Result	Monitoring performance
0.9071381688117981	0
Propensity	Monitoring evidence
---	0.0

Customer Barbara will probably remain loyal; the model returned a low probability to churn for her.

▽ Outputs

Results

Result	Monitoring performance
0.0010575958294793963	0
Propensity	Monitoring evidence
---	0.0

By default, the results of the model are shown in the Results field.

Model results are unique for each framework type on which a model is built.

Pega offers full support for the xgboost, tensorflow, kmeanclustering, knn, linearlearner and randomcutforest frameworks.

Once the predictive model rule is created, it can be used in next-best-action strategies ...

... in a similar way as native Pega machine learning models and third-party models imported using PMML or H2O.ai.

But there is an important difference to keep in mind.

Native and imported models, using the required input data, execute inside Pega.

In the case of machine learning services, the input data required by the model is sent to the external platform, the model is executed externally, outside of Pega ...

...and the result is sent back to Pega using a secured connection.

You've reached the end of this demo. What did it show you?

- The high level steps involved in creating a model using Amazon SageMaker Autopilot.

- How to connect to external machine learning services and run a model externally.

Model transparency

Transparency score

Transparent artificial intelligence is becoming an important requirement for many businesses. In risk management, decisions need to be explainable, and opaque predictive models are not allowed. In marketing, the policy for the transparency of models might be less strict and allow for the use of opaque models.

Each model type that comes with Pega® Platform is assigned a transparency score ranging from 1 to 5, where 1 means that the model is opaque, and 5 means that the model is transparent. Highly transparent models are easy to explain, whereas opaque models might be more powerful but difficult or not possible to explain. For example, a decision tree has a high transparency score, whereas a neural network model has a low transparency score.

Transparency scores by model type ?				
The scores and thresholds determine in which business issues use of the model will be compliant.				
Adaptive model Pega Compliance All business issues	Bivariate model Pega Compliance All business issues	Genetic algorithm Pega Compliance All business issues	Regression Pega or PMML Compliance All business issues	Tree model Pega or PMML Compliance All business issues
Association rules PMML Compliance All business issues	Clustering model PMML Compliance All business Issues	Ensemble model PMML Compliance All business issues	General regression PMML Compliance All business issues	k-Nearest neighbors model PMML Compliance All business Issues
Naive Bayes model PMML Compliance All business issues	Neural network PMML Compliance All business Issues	Random forest PMML Compliance All business issues	Scorecard PMML Compliance All business issues	Support Vector Machine PMML Compliance All business issues

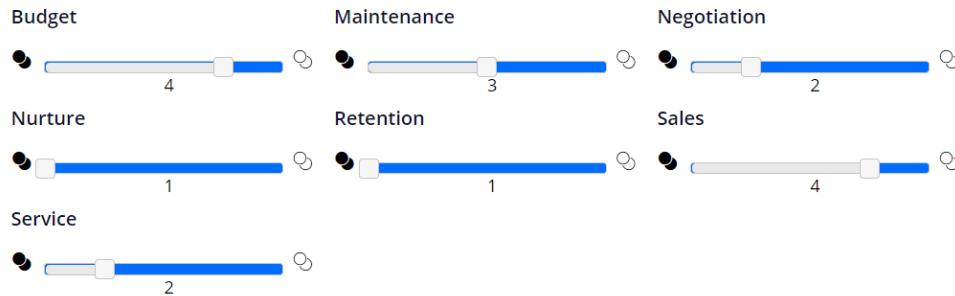
Configuring the model transparency policy

Depending on the company policy, models are marked as compliant or non-compliant for a specific business issue. By default, the transparency threshold is set to 1 and all model types are allowed in all business issues. Lead data scientists can modify transparency thresholds for different business issues. For example, they can increase the threshold for Budget to indicate that opaque models are non-compliant in that area. To set the thresholds in Predictive Studio, go to the Settings section and choose Model transparency policies.

Transparency threshold per business issue [\(i\)](#)

Models with a transparency score above or equal to the threshold of business issues, are compliant. A high transparency score of 5 indicates that models are fully auditable.

Range of scores  1 (Opaque)  5 (Transparent)



Upgrading predictions

Description

Learn how to replace the predictive model that drives a prediction with a new model. You can import a predictive model that is built on an external platform or connect to a machine learning service. Deploy the model in shadow mode. In shadow mode, the candidate model ingests production data but does not affect decisioning until it is promoted to the active model status.

Learning objectives

Import a model built on an external platform

Connect to a machine learning service

Deploy a new model in shadow mode

Promote the new model to the active model status

Model operations process

Introduction

Learn how you can improve the performance of your predictions by using a standardized model operations process (MLOps). MLOps lets you replace a low-performing predictive model that drives a prediction with a superior model created in a third-party platform.

If a candidate predictive model is deployed in shadow mode, it can be monitored with real production data without impacting the business outcomes. If the model proves effective, it is deployed as the active model.

Transcript

This video shows you how to update a predictive model in a prediction.

In the standardized model operations process, the active model is replaced with a better one in a production environment by using the shadow mode option.

Pega Customer Decision Hub™ uses predictions to determine the next best action for each customer.

A prediction is driven by an adaptive model, a predictive model, a scorecard model, a field model, or a combination of these models.

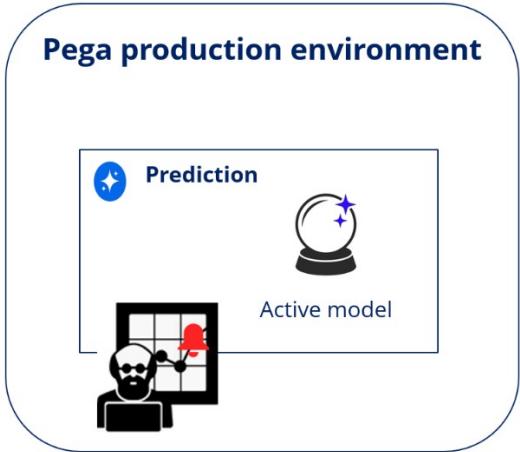


You can replace a model in the production environment at any time through the model operations process.

As a data scientist, you may respond to a Prediction Studio notification that an active model does not generate enough lift ...

... and decide to replace the low-performing model with a high-accuracy external model.

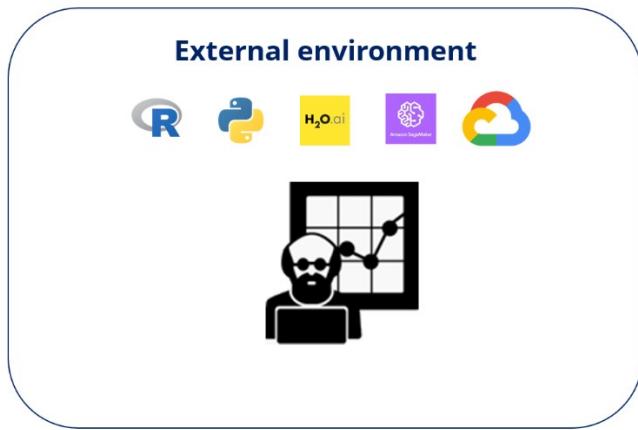
Or you can update a prediction regularly, for example, whenever you develop a new model.



To build a new model, you can use Pega machine learning or an external environment.

You can use data science tools that can export models in the PMML format, such as R or Python. The H2O format is another option.

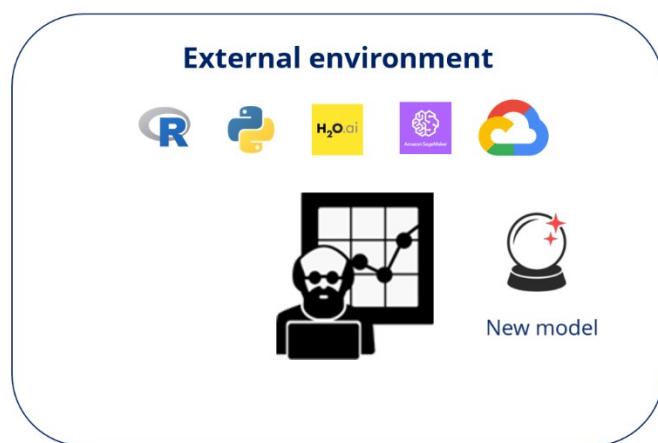
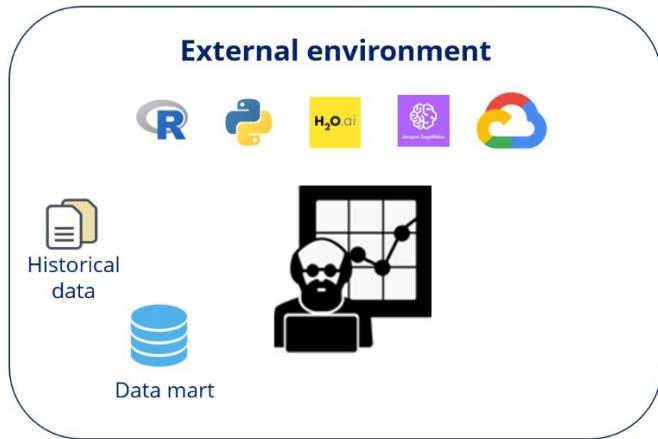
You can also connect to the Amazon SageMaker or Google Cloud machine learning services.



When an application runs in the production environment, you can capture predictor and outcome data for every customer interaction.

You can utilize the historical data of models captured by the system by importing these records into your external environment of choice.

The historical data can be combined with data from other sources to build a new model.



Once the new predictive model is developed, you validate the active model and the candidate model against the same data set to compare their metrics in Prediction Studio.



If the candidate model outperforms the active model, approve the model.



You can choose to replace the active model immediately or place the new model in shadow mode.



In the business operations environment, a change request is automatically generated.

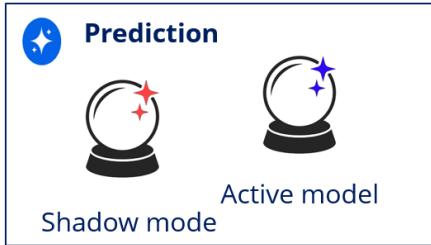


The revision manager approves the change request and packages it in a revision.



The revision is then promoted to the production environment.

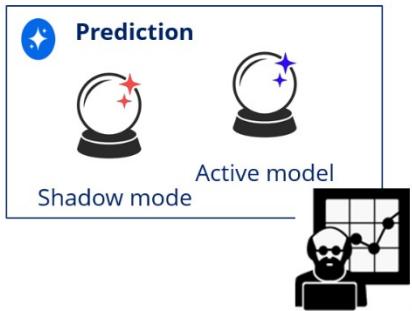
Pega production environment



If you deploy the new model in shadow mode, the new model is exposed to the production data but does not drive the prediction yet.

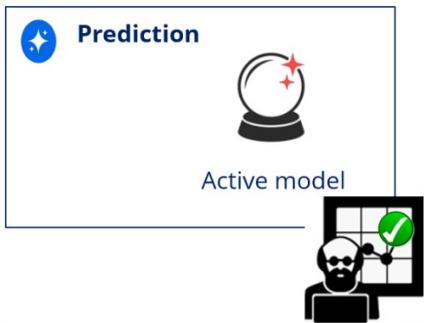
Shadow mode allows you to monitor the model performance in a production environment before deploying it as an active model.

Pega production environment



After monitoring the prediction for some time, you can promote the shadow model to active.

Pega production environment



You have reached the end of this video. What did it show you?

- How the model driving a prediction is updated with a new predictive model
- How the shadow mode allows monitoring of a new model in a production environment

Placing a predictive model in shadow mode

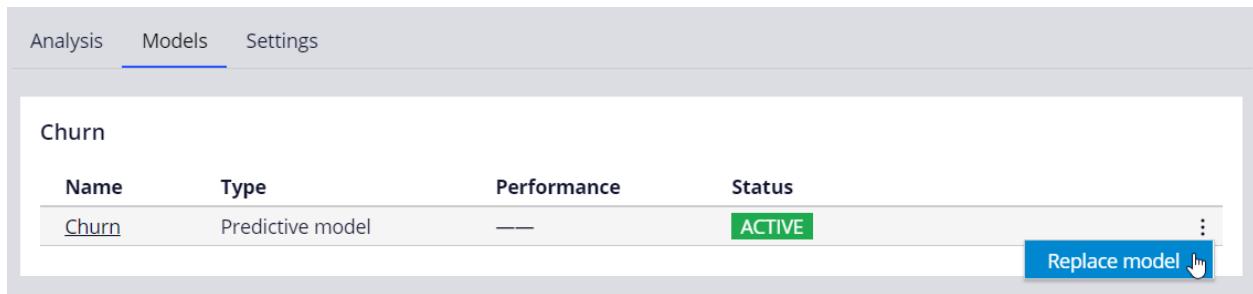
As a data scientist, you can approve changes to the models that drive predictions for deployment to the production environment. You can change models independently or respond to a Prediction Studio notification that a prediction does not generate enough lift.

To improve the performance of a prediction, you can replace a low-performing model with a high-accuracy external model that you upload to a Pega repository or directly to Prediction Studio. As a result, you start a standard approval and validation process to deploy the model update to production. Before you approve any changes, you can compare the candidate model with the existing model based on data science metrics, such as score distribution or lift.

Model deployment

In Pega Customer Decision Hub™ environments, changes to models that you approve in Prediction Studio are deployed to production through Pega 1:1 Operations Manager and the Business Change pipeline.

The process begins when you update a prediction in your non-production environment.



Name	Type	Performance	Status	⋮
Churn	Predictive model	—	ACTIVE	Replace model

You can replace the active model with a predictive model, scorecard, or field in the data model that contains a score.

Replace model

What do you want to replace it with?

- Model
A machine learning model to calculate a score in real-time
- Scorecard
A simplified method to calculate a score in real-time
- Field
An existing field in the data model that already contains a precalculated score

You can select a model from Prediction Studio, upload a PMML or H2O model to Prediction Studio, or connect to an external model that you developed on Amazon SageMaker or Google AI Platform.

Replace model

Compare the models [?](#)

[Upload](#) [Machine learning service](#) [Model list](#)

Select a PMML, H2O MOJO or Pega OXL file

[Choose File](#)

When creating a machine learning model in a third-party environment, you can use historical data from Pega Platform™ and other sources to train the model.

Recording historical data

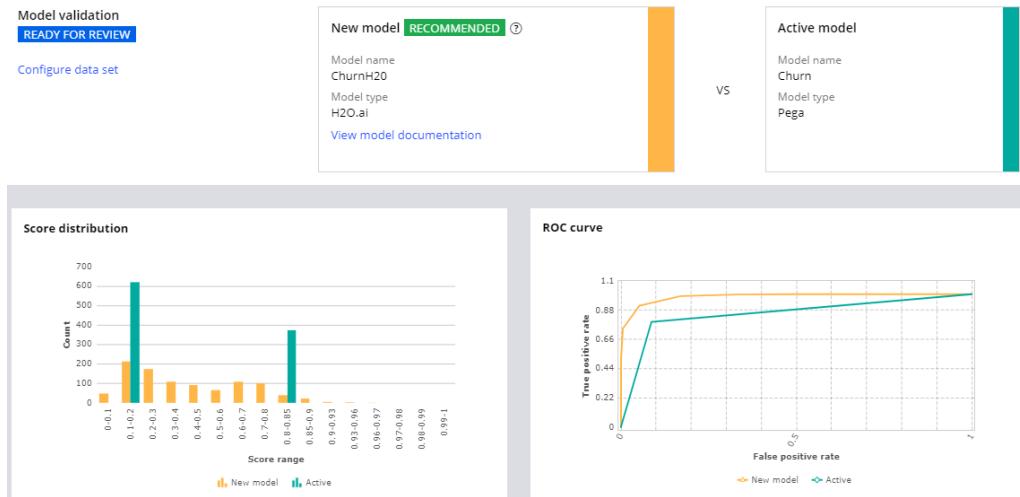
Save historical data in a repository to use for offline analysis. You can find an overview of the historical data in [Historical data overview](#).

Record historical data

churned	Sample percentage <input type="text" value="100.00%"/>
loyal	Sample percentage <input type="text" value="100.00%"/>
File location <input type="text" value="file:///defaultstore/Predictive/Rule-Decision-PredictiveModel/PegaCRM-Data-Customer/Churn/records*.json"/>	

You can validate the candidate model against a data set and compare the new model with the current model.

This analysis provides relevant metrics to help you decide which model has better performance with a static data set.



After you evaluate the models, you can approve or reject the candidate model for deployment to production. You can place the model in shadow mode or immediately replace the current model with the new model.

Evaluate ChurnH20

Evaluate the model and provide your feedback. (?)

Evaluation

- Approve new candidate model and start shadowing (recommended)
- Approve candidate model and replace current active model
- Reject candidate model

In a Pega Customer Decision Hub environment, the system creates and resolves a change request in Pega 1:1 Operations Manager. A team lead can verify the changes in the rules and the relevant documentation. The change request is packaged into a revision, and a deployment manager can promote the prediction with the candidate model to production.

If you deploy the candidate model to production in shadow mode, it runs alongside the original model, receives production data, and generates outcomes, but the outcomes do not impact business decisions.

Churn

Name	Type	Performance	Status	⋮
▼ Churn	Predictive model	---	ACTIVE	⋮
ChurnH20	Predictive model	---	SHADOW	⋮

Model promotion to active

If the model proves ineffective, you can reject it and add another model to the prediction. If the new model performs well in production, you can promote it to the active model position.

Name	Type	Performance	Status	⋮
Churn	Predictive model	—	ACTIVE	⋮
ChurnH20	Predictive model	—	SHADOW	⋮

Name	Type	Performance	Status	⋮
ChurnH20	Predictive model	—	ACTIVE	⋮

Text analytics for email routing

Description

Humans can effortlessly interpret a single tweet but are unable to parse a large volume of information efficiently. Businesses are exploring ways to use machine learning to extract meaningful information from a large quantity of text messages. These insights help improve business performance and customer experience.

Learning objectives

- Explain text analytics
- Name a few practical applications where text analytics can be used
- Describe the role of artificial intelligence (AI) in text analytics
- Explain how text models for text categorization are trained

Text analytics

Introduction

A company wants all incoming emails to be automatically routed to the correct department. The company must find a way to process the text and extract the information for which it is looking. Text analytics is the most efficient and effective way to obtain information.

Transcript

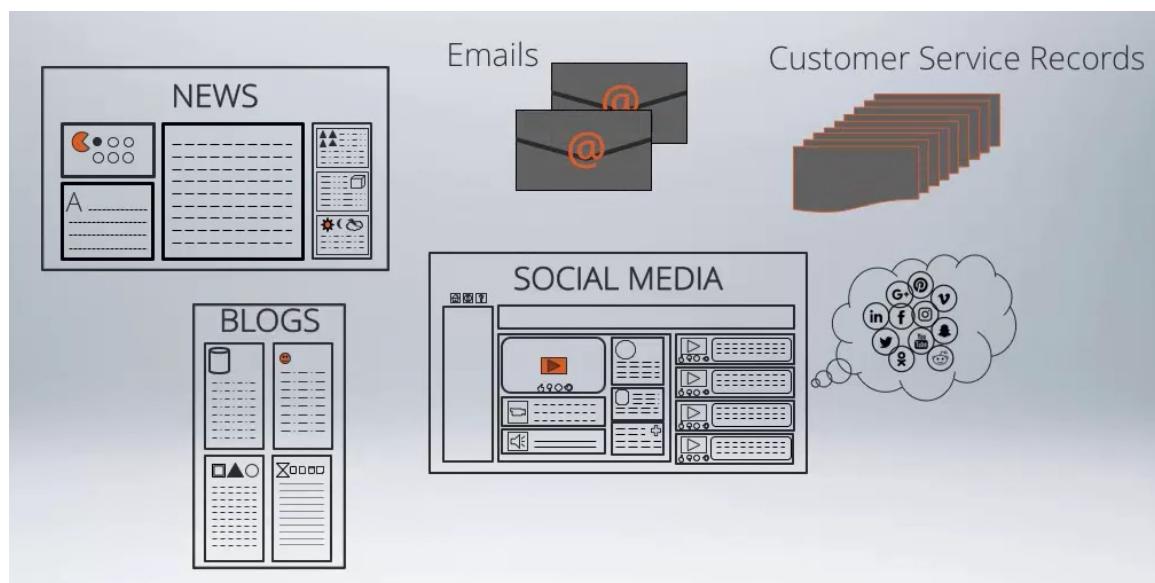
In this lesson you are going to learn what text analytics is and how it is used to help understand customer sentiments and needs.

Consider this message from a customer.

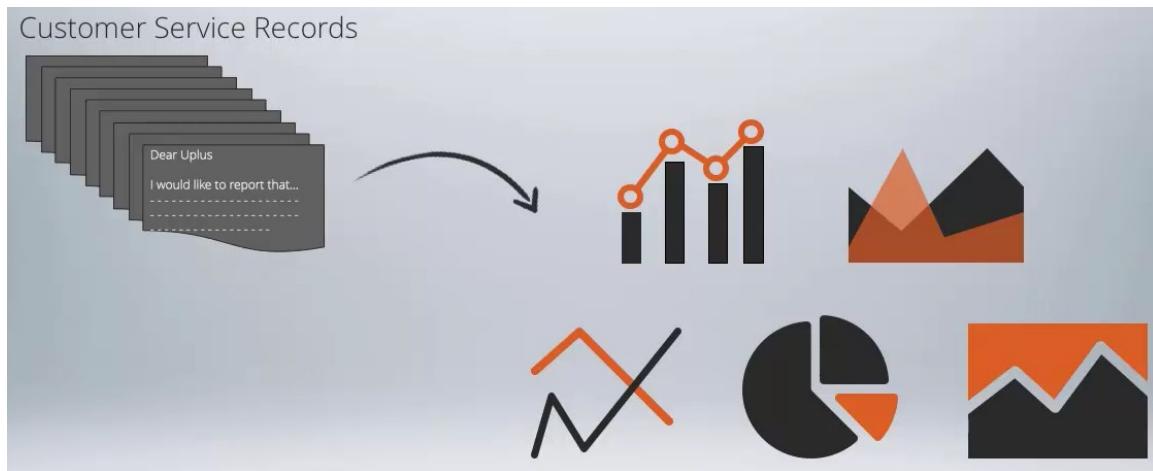
I've been a customer of UPlus for over 5 years.
I love the cell phone connectivity. But now I am being overcharged for the services in the August-2017 billing statement.

Humans have the ability to understand the implications of a piece of text. We can summarize it or elaborate on it or even re-phrase it and still preserve the original meaning. In fact, we do this very well. But can we do it on a large scale?

Today, digital sources such as news, blogs, social media and emails are generating more text than ever before.

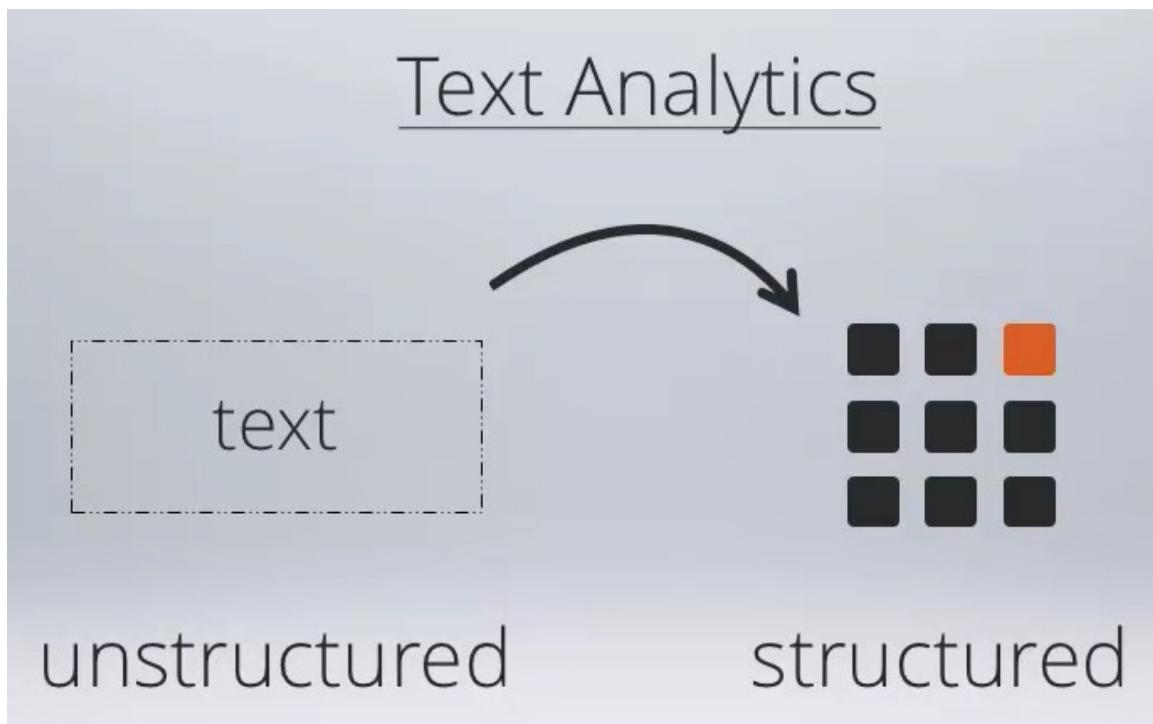


It's not possible for humans alone to analyze and extract valuable information from such a large volume of text. Consider a company that has tens of thousands of customer service records in text form.

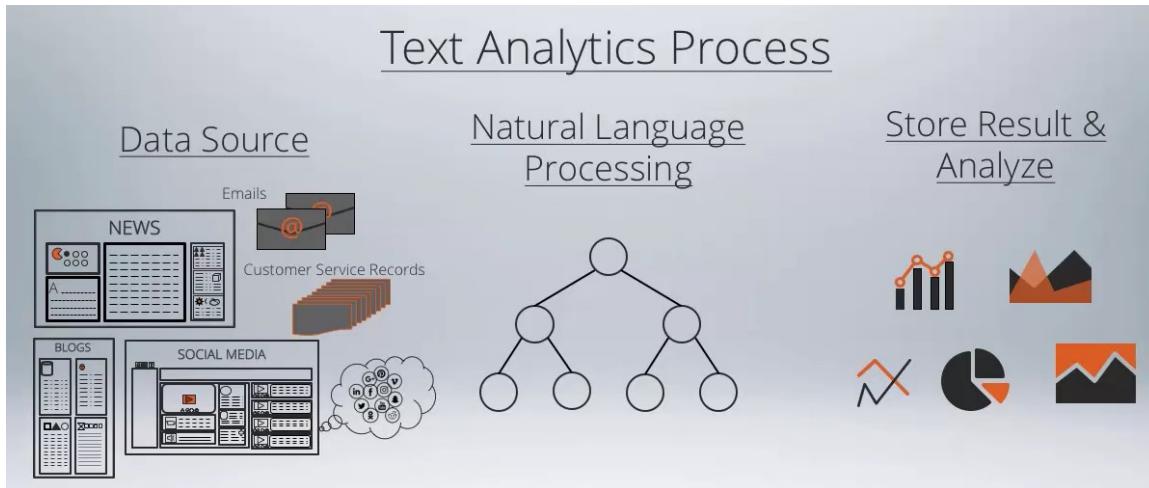


Let's say the company wants all of the complaints to be automatically routed to the right department. Then it wants a report that breaks down the complaints by department. The company also wants to use these records to understand how satisfied its customers are about various products and services. The company must find a way to process the text so that the information it's looking for can be extracted. Text analytics is the most efficient and effective way to do this.

Text Analytics is the process of deriving high-value information from text.

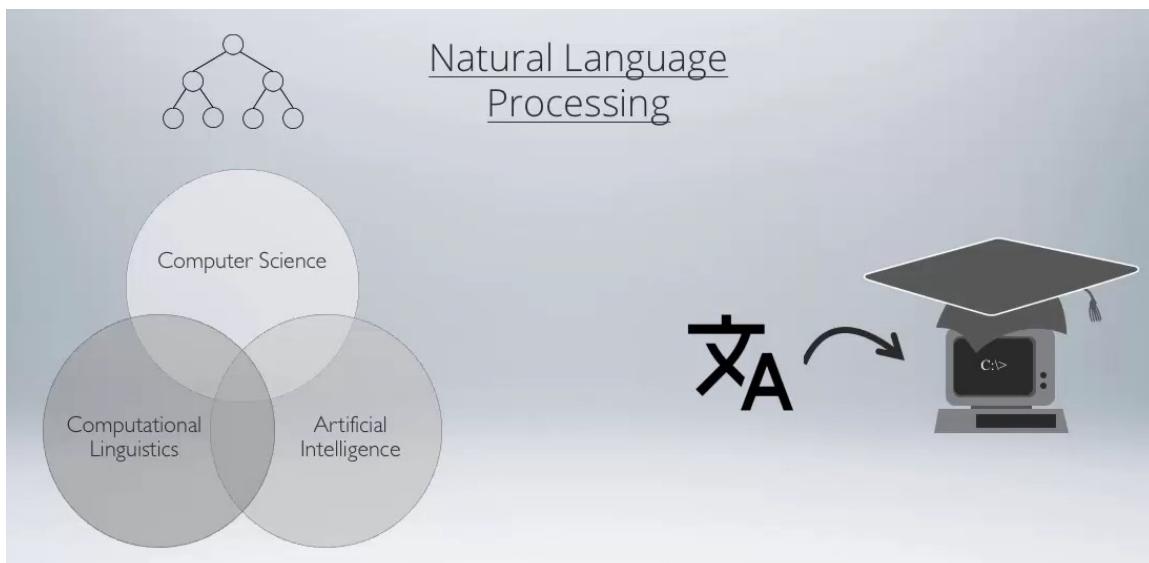


It involves converting unstructured text into structured data that can be analyzed using statistical methods. There are three main steps involved in text analytics.



First, data is fetched from a source, such as social media, emails or customer service records. Second, natural language processing techniques are applied to extract certain attributes from the text and present them as structured data. And third, the extracted structured data is stored and analyzed.

Natural language processing is a field of computer science, artificial intelligence and computational linguistics concerned with the interactions between computers and human language.



In particular, it is concerned with programming computers to fruitfully process large volumes of natural language text. At a high level, the following analyses occur during natural language processing.

Natural Language Processing

High-level Analyses

Syntactic Analysis

- Study of the structure
- Find words, word roots, phrases, sentences

Semantic Analysis

- Study of the meaning
- Establish relationship between words and sentences

Linguistic Analysis

- Study subtle nuances in a language.
- Ascertain the true meaning of a sentence.

Syntactic analysis analyzes the structure of text and recognizes the parts of speech, such as nouns, verbs, adjectives, etc. Semantic analysis establishes the meaning of a piece of text. So the relationship between words and sentences is analyzed. In the linguistic analysis phase, the nuances of a language are considered to ascertain the exact meaning of the text.

Here are the outcomes of natural language processing.

Natural Language Processing Outcomes

I've been a customer of UPlus for over 5 years. I love the cell phone connectivity. But now I am being overcharged for the services in August-2017 billing statement.

Language detection

- English

Sentiment detection

- Overall: Neutral
- 1st sentence: Neutral
- 2nd sentence: Positive
- 3rd sentence: Negative

Classification

- Billing/Payment > Accuracy
- Service Experience > Call Quality

Entity detection

- UPlus
- August-2017

The language of the text is detected. The sentiment is detected as positive, negative or neutral. Sentiment is the general attitude expressed towards a subject. The overall sentiment as well as the sentiment per sentence is reported.

The text is classified into one or more pre-defined categories, such as the business functions in a company. For each category detected, the sentiment of the corresponding sentence is also reported. Entities refer to the proper nouns found in the text such as names of people, places, dates and times, organizations, etc. This helps in establishing the subject of the text.

You have now reached the end of this lesson which provided an overview of text analytics and how it works.

Email routing

Introduction

A company has an email channel where its customers can send any email — from service requests to compliments to sales inquiries — to the product support team. This email variety makes it difficult to provide each customer with a prompt and personalized response. Pega Infinity™ uses AI-powered text analytics to perform intelligent email routing.

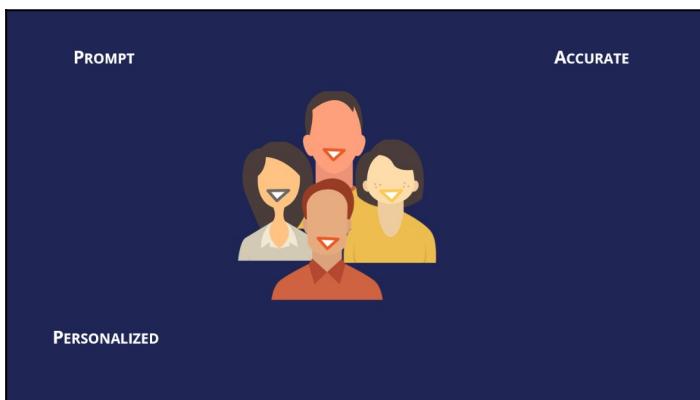
Transcript

In this video, you will learn how Pega Infinity™ uses AI-powered text analytics to do intelligent email routing.

Customer satisfaction is a reflection of what a customer expects from a company vs. what they experience from the company.

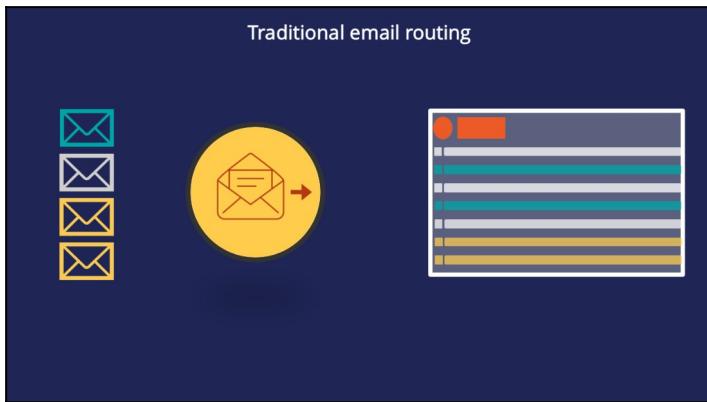
Meeting, or even better, exceeding customer expectations means addressing their service requests and complaints promptly, accurately, and with personalized solutions.

Doing this will ensure the customer has a great experience.



Assume a company has an email channel in which its customers can send any type of email —from service requests to compliments to sales inquiries—to the product support team.

These emails are often routed to the same container and are uncategorized. This makes it difficult to provide each customer with a prompt and personalized response.



In Pega Infinity, this issue can be addressed by routing the emails using artificial intelligence.

Pega Infinity uses AI-powered Natural Language Processing to detect the topic of an email and route the email to the appropriate container.

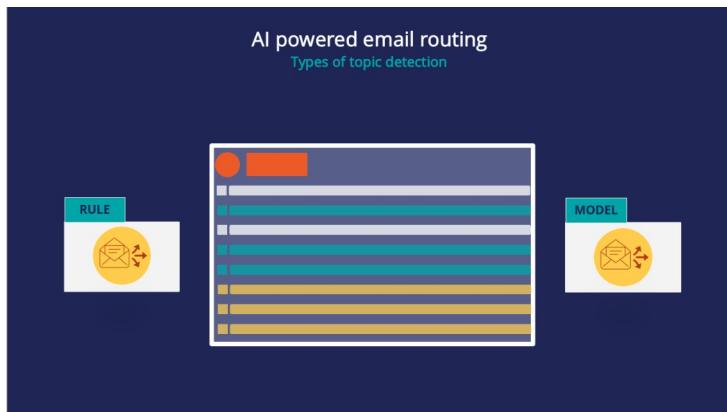
Consider emails from customers requesting an account address change, making a compliment, or requesting a new credit card.

With the help of AI-powered text analytics, Pega Infinity is able to read and understand the content of each email and route it appropriately.

This means customer service representatives can be alerted to any account-related service requests and resolve them quickly.



Email routing is done using the topic detection mechanism. The two types of topic detection are rule-based and model-based.



In rule-based topic detection the routing is based on the rules configured in the email channel.

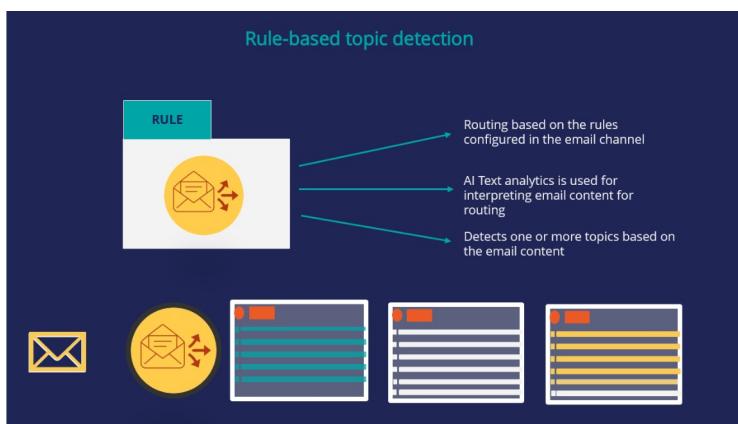
AI-powered text analytics is used to detect the topic of the email, and the channel rules route it to the right container.

This type of topic detection may detect one or more topics if the email contains words associated with more than one topic.

Let's consider a service request email from a customer. The email content is analyzed and routed to the right container.

If an email from another customer contains words that are associated with two topics, the rule-based topic detection detects both topics.

The email can then be routed to two different containers depending on how the channel rules are configured.



In model-based topic detection the routing is based on AI models built by a data scientist using machine learning.

AI powered text analytics is used to detect the topic of the email. Building these models requires a training data set and a test data set.

The data sets consist of a list of emails and the associated topic for each email.

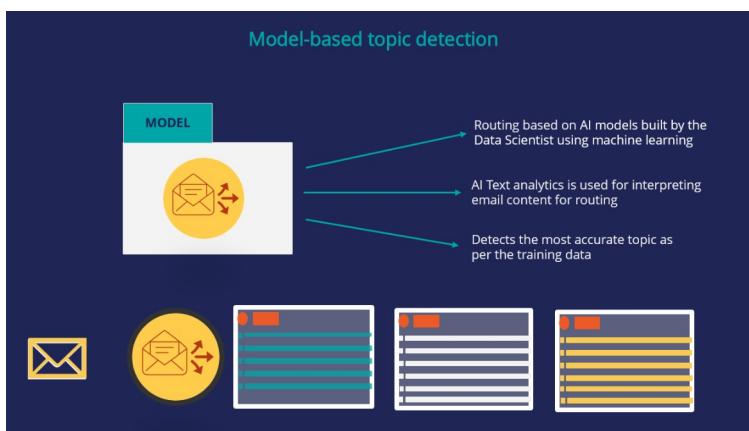
This type of topic detection identifies the most accurate topic based on the AI model and training set used by the data scientist.

Let's consider the same service request email from the customer. The email's content is analyzed and routed to the right container.

If an email from another customer contains words associated with two topics during the training of the models ...

...the model-based topic detection detects both topics but typically with a different accuracy factor.

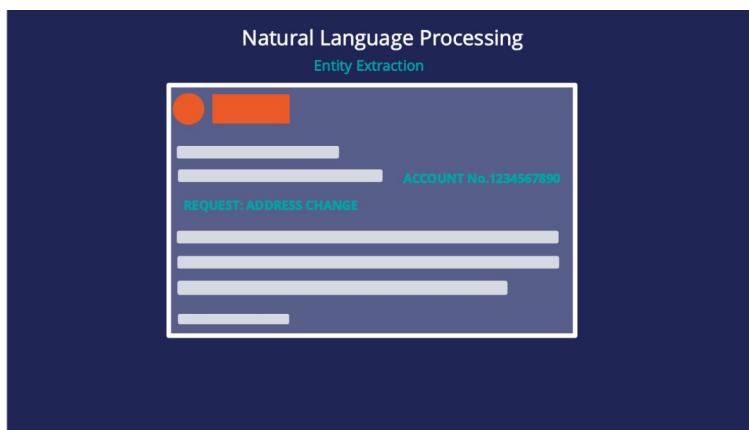
In this case, the topic with the highest accuracy factor is chosen.



Pega Infinity also enables you to extract entities from an email.

This means that when an email is sent, certain entities such as account number, email address, street address, etc. can be automatically detected and extracted.

This allows certain emails to be automatically processed or given priority.



Besides Topic detection and Entity Extraction, Pega Infinity uses its AI-powered text analytics to enable you to detect the sentiment of an email based on its content.

Suppose a customer sends the following email to customer service. The email's content is a combination of different types of sentences.

The first sentence has a neutral sentiment. In the second sentence, the customer expresses his satisfaction with the connectivity, so it has a positive sentiment.

The third sentence is negative, as the customer is complaining about an overcharge.

The overall sentiment of the email is determined based on the weight of the content sentiments. In this specific example the overall sentiment is neutral with a sentiment score of 0.05.

Detecting email sentiment

I've been a customer of UPlus for over 5 years.
I love the cell phone connectivity.
But now I am being overcharged for the services in August-2019 billing statement.



Neutral (0.05)

The sentiment score is a value between -1 and 1.

In the out-of-the-box configuration, a sentiment score <-0.25 results in a Negative sentiment, a sentiment score between -0.25 and 0.25 results in a Neutral sentiment ...

... and a sentiment score above 0.25 results in a positive sentiment.

Detecting email sentiment

Sentiment Score
-1 to 1

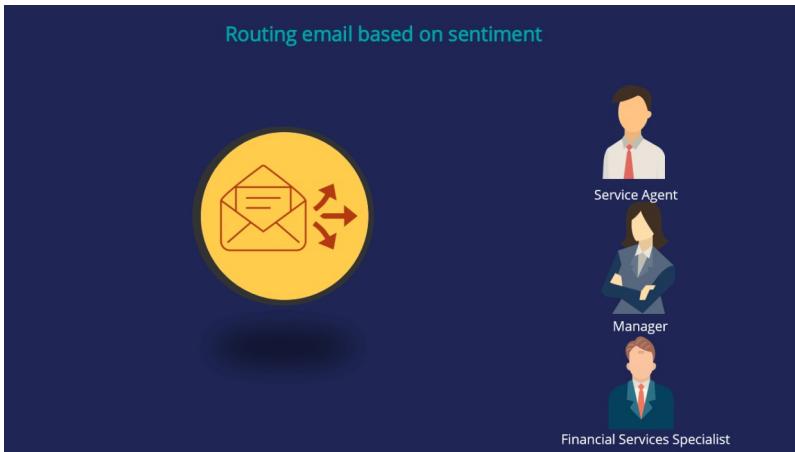


<-0.25 -0.25 to 0.25 > 0.25

Once the email sentiment is detected, you can configure the email channel to route a specific topic with a specific sentiment to a specialized agent for a quick and personalized response.

For example, you could route an address change with a neutral sentiment to a Service Agent, a complaint email with a negative sentiment to a Manager ...

... and a credit card inquiry with a negative sentiment to a Financial Services Specialist.



In summary, Pega Infinity's AI-based email routing capability enables customer service representatives to be more productive, reduces request processing time ...

... and improves the customer experience by providing prompt and personalized service.



Training a topic model to improve email routing

Introduction

U+ Bank uses Pega Customer Service™ to route incoming emails to the appropriate department based on the topic of the email. For several use cases (for example, an address change), emails are routed based on keywords that are detected in the message. To improve the email routing, learn how to train the text prediction with a data set that contains classified messages.

Transcript

This demo shows you how to train a text prediction to improve email routing.

U+ Bank uses machine learning to route inbound messages in the email channel to the appropriate department based on the topic of the email.

A text prediction that aims to detect the topic of the message drives the routing.

U+ Bank customer support



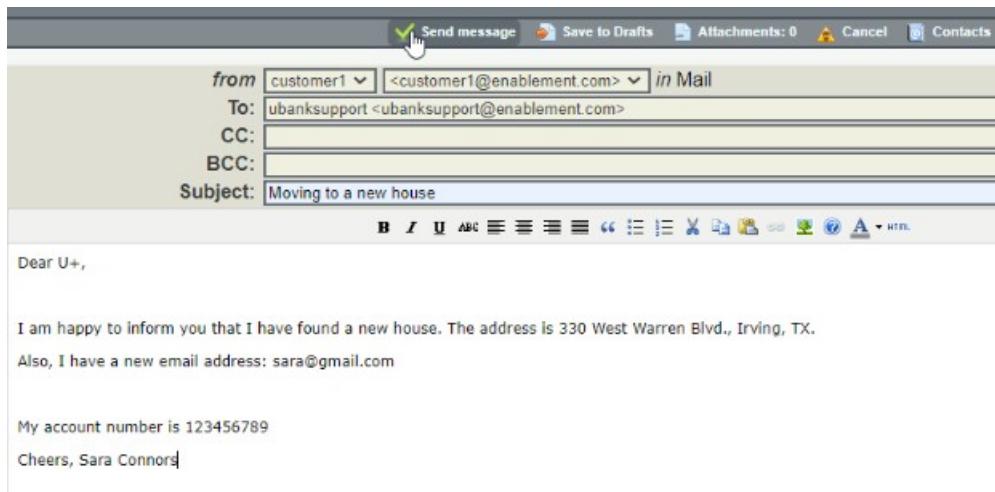
U+ Bank customer support

Mar 11, 2021

[Open prediction](#)

⋮

For example, when Sara writes an email to inform U+ Bank that she has moved to a new house, the text prediction detects an address change as the topic.



The **Account maintenance** department receives all emails where an address change is detected as the topic of the message.

My Work

Messages & coaching

Handling account address changes

Service case tip- Address Change: It's important to keep account addresses up to date. Check out this benefit.

Select... Account Maintenance Branch processing Coaching session Customer Circumstance Customer data rights Customer suggestions GDPR requests General service requests Inbound correspondence Material fulfillment Outbound Call Outbound campaigns Outbound correspondence Pending outbound call Statement copy requests Transaction disputes

My work View queue for Select...

The text prediction also detects entities such as a ZIP Code and an email address and the overall sentiment of the message.

Email OPEN

 customer1 Tue, Mar 09, 05:52 AM (42m ago)

Subject: Moving to a new house

Dear U+,

I am happy to inform you that I found a new house. The mailing address is
222 West Las Colines Blvd., Irving, TX 75390, USA

Also, I have a new email address: sara@gmail.com

Cheers, Sara Connors

Notice that the topic Account Address Change is detected with 100% confidence.

Text analysis

Collection of all information extracted from this email and attachments

All content	Topic	Confidence
Entities	Account Address Change	100%
All entities		
Email	Language	
State	English	
Postal code		
City		
Other		
Topics		
Email header info		

The intelligent routing is set up in the email channel, in App Studio.



U+ Bank Support

ubanksupport@enablement.com

Suggested cases

Close Account, Account Address Change, Order statement copy, Dispute Transaction

Suggested replies

Apology, Clarification, Thank You

An email is routed to the work queue of the Account maintenance department when the detected topic is an address change.

Complaints are routed to the **Transaction Disputes** department.

If the address change and complaint topics are not detected, the email is routed to a default work queue.

Intelligent routing

Process: First matching action ▾

Triggers only the first action that matches conditions.

Action	Route to work queue	Value	AccountMaintenance	
When	Topic	Is equal	Account Address Change	
Action	Route to work queue	Value	TransactionDisputes	
When	Topic	Is equal	Complaint	
+ Add condition				
Otherwise	Action	Route to work queue	Value	InboundCorrespondence

Note that in the **When** condition of the email routing, you use the outcome of the text prediction.

Optionally, you can also use the sentiment of the email in the routing conditions.

Topic Is equal Account Ac
Header NLP Entities
Condition Topic Decisioning result
In Sentiment Language

Route to work queue Value Inbou

You can set a threshold for the confidence with which a topic must be detected to trigger routing.

Confidence threshold

Is Greater Than 0.8 OK

Every channel is associated with a text prediction.

Text Analytics Use advanced configuration

[Edit topics](#)

Record training data

Enable subject analysis

Analyze email attachments

[Open Text Prediction](#) 

A text prediction is driven by predictive models that detect topics, entities, and sentiments.

Outcomes

Manage all topics, sentiment and entities that should be part of this prediction

[Topics](#) [Entities](#) [Sentiments](#)

Many entity extraction models, and a sentiment model, are available out of the box.

For topic detection, Prediction Studio supports keyword models as well as machine learning models based primarily on training data.

A keyword model uses **Should words**, **Must words**, **And words**, and **Not words**.

Action > Account Address Change

Machine learning [Keywords](#)

A single keyword can consist of multiple words. Use tab/enter to separate keywords.

Should words

[move](#) [change](#) [address](#)

Must words



And words

Not words

If any of the **Should words** appear in a piece of text, topic detection assigns that text to the corresponding topic. To achieve accurate results, create an exhaustive list of **Should words**.

Only if all **Must words** appear in a piece of text, will the topic detection assign that text to the corresponding topic.

Use **And words** to distinguish between similar topics while using identical **Must words**.

If a **Not word** appears in a piece of text, the text is not assigned to the corresponding topic.

You can test the output of a prediction on a sample message.

Two topics are correctly detected: an address change and a complaint.

Notice that the confidence score for both topics is **1**. The keyword model performs a Boolean match based on the presence or absence of words and detects the topic with absolute certainty.

▼ Topic detection

Granularity
Document

Analysis
Rule

Topic	Sentiment	Sentiment score	Model name	Model type	Confidence score
Action > Account Address Change Neutral	0		pyInteractionTaxonomy Pega NLP		1.00

It is highly recommended to use a machine learning model for topic detection. It determines the confidence score based on evidence.

A machine learning model is based on training data that consists of categorized texts.

You can get training data from two sources. The first is training data accumulated through the application that is running in production.

When a customer service representative corrects the topic of an incoming email, this change is added to the training data.

After the messages are reviewed and approved, they are used to rebuild the topic model.

Also, you can choose to import data, provided you have accumulated data in the past.

The file must contain the message and the associated topic.

A	B	C
1 content	result	type
20 Effective 11/28/17 please change mailing address to	Action > Account Address Change	Test
21 Please amend the mailing/billing address to the follow	Action > Account Address Change	Test
22 Confirm - mailing address: 1430 E 33rd St., Signal Hill	Action > Account Address Change	Test
23 Effective 8/28/17, please change the insured's mailin	Action > Account Address Change	Test
24 Per our insureds request, please issue the following c	Action > Account Address Change	
25 Effective 8/25/17, please correct the mailing address	Action > Account Address Change	
26 Effective 8-16-17 please amend the mailing address e	Action > Account Address Change	
27 Please amend the mailing address.	Action > Account Address Change	
28 Effective 8/28/17 please change the mailing and billir	Action > Account Address Change	
29 AMEND insured's MAILING ADDRESS to read:	Action > Account Address Change	
30 Effective: 09/01/2017 Please amend mailing address	Action > Account Address Change	
31 Please update mailing address per the below:	Action > Account Address Change	
32 We have received a request to change your mailing a	Action > Account Address Change	
33 Please change mailing address to read: 197 Warren /	Action > Account Address Change	
34 Please change the mailing address as per the attache	Action > Account Address Change	
35 Please amend mailing & billing address to 3253 Phillip	Action > Account Address Change	

Notice that training data for the address change and the complaint topics are in the pending state.

You can now rebuild the models.

You can select individual models or rebuild all models across model types and languages.

While training is in progress, you have the option to cancel.

When the process completes, you can view the latest model report on the **Models** tab.

The report contains the validation data, the confusion matrix, and the score sheet.

The topic detection now uses machine learning models based on the training data and the keywords provided for the address change and complaint topics.

Outcomes			
Manage all topics, sentiment and entities that should be part of this prediction			
Topics	Entities	Sentiments	
Language			
English ▾			
Topics	Total keywords	Total training data	F-score
Action > Account Address Change	10	82 (0 pending)	1.00
Action > Close Account	10	0 (0 pending)	
Action > Complaint	7	84 (0 pending)	1.00
Action > Report a lost or stolen card	13	0 (0 pending)	

The **Should words** and **Must words** act as positive features for matching text to a topic, while the **Not words** act as negative features.

But the training and testing data have the greatest impact on your machine learning model, while keywords have a smaller impact.

Once the models are rebuilt, you can test the prediction.

Notice that the topics are detected with a confidence score below one.

In this case, the topic message is recognized as a complaint with high confidence. The address change topic is detected, but with a lower confidence score.

When multiple topics are detected in a message, comparison of the confidence scores allows selection of the one with the highest priority.

Topic	Sentiment	Sentiment score	Model name	Model type	Confidence score
Action > Complaint	Neutral	-0.25	U+ Bank customer support Pega NLP	0.75	
Action > Account Address Change	Neutral	-0.25	U+ Bank customer support Pega NLP	0.25	
Action > Report a lost or stolen card	Neutral	-0.25	U+ Bank customer support Pega NLP	0.00	

The other outcomes of the text prediction are the entities detected, such as an email address and the sentiment of the message.

Output

I am happy to inform you that I found a new house. The address: 330 West Warren Blvd.. Irving.
Also, I have a new email: sara@gmail.com For the third time, I have noticed an invalid transaction.
The amount is small, but could you please check?

Sentiment indication

Positive Negative Neutral

You have reached the end of this demo. What did it show you?

- How Pega Customer Service routes incoming emails to the appropriate department based on the topic of the email
- How text predictions work to predict the topic and sentiment of a message and detect entities
- How to train a topic model and use machine learning to identify the topic correctly